

Quantum Fault Tolerance Meets Toom’s Contours

David Ponarovsky

August 30, 2026

Abstract

We suggest a new proof of quantum fault tolerance based on the alternative proofs of Toom’s theorem [BS88] (see also [Gac21]). The fault tolerance construction is relative to the circuit model, subjected to local stochastic noise, where at each time unit only a constant depth of non-noisy classical computation is allowed. Our main error correction code is the (d, d') -Toric code, namely qubits are set on the d' -dimensional faces of the d -dimensional Toric. As byproducts of the study, we provide and prove the following:

1. An explicit description of the sweep-rule—a generalization of Toom’s rule for topological codes—for any dimension greater than two. In doing so, it answers affirmatively to the conjecture made by Kubica and Preskill [KP19].
2. A fault tolerant construction in the circuit model that both uses the (d, d') -Toric code and has a constant-depth correction cycle, where $d', d - d' \geq 2$. Until now, the use of a topological code as the primary encoding and a constant-depth correction cycle were only achieved separately; this is the first work to achieve them simultaneously.
3. A novel fault tolerance that doesn’t introduce an asymptotic depth overhead per Clifford gate. In particular, there is a threshold $p_0 \in (0, 1)$ such that for a noise rate $p < p_0$, and any n large enough, the construction takes a circuit with total depth μ and a $\mathbf{T}$ -depth $\mu_{\mathbf{T}}$, and results a circuit that simulates the given with probability:

$$1 - \Theta(\text{Poly}(\text{circuit’s size}))p^{\Theta(n)}$$

and has a depth $\Theta(\mu) + \Theta(\mu_{\mathbf{T}}n^{d+1})$.

Specifically, all these results hold for the 4D-Toric code.

Contents

1	Introduction	2
1.1	Background	2
1.2	Stating Our Results	3
1.3	Discussions, Applications, and Further Directions	4

2	Outline of the Proof	5
3	Notations and Preliminaries	15
4	The (d, d')-Dimensional Toric Code	19
4.1	Sweep Rule	22
4.2	Infinite Sides Torics	26
4.3	Reduced Error	30
4.4	Shrinking	36
4.4.1	Ideal decoder	38
4.5	Logical Gates	38
5	Toom's Contours	44
5.1	The Barrier Argument	44
5.2	Toom's Contours	46
5.3	Space-Time Graph	50
5.4	Minimal Explantation Argument	58
5.5	The Replacement Argument	61
6	Proof of the Theorem	64
7	Appendix	68
7.1	Time Coarse-Graining	68
7.2	Checks Commutation	70
7.3	Infinite Compatibility	71
7.4	Counting Subtrees.	73
7.5	Tree Separator Lemma	73

1 Introduction

1.1 Background

Fault-tolerant constructions take ideal quantum circuits and establish their logical equivalence such that they endure fault occurrences. Their efficiency is typically measured by the amount of additional resources – qubits, computation time, range connectivity, etc. – required by the resulting circuit, beyond the scope of theoretical discussion; those parameters translate into the machine complexity and energy cost required to run the circuits. Minimizing overheads brings quantum computing applications closer to the near future.

Since the introduction of the quantum threshold theorem [AB99], which showed that quantum fault tolerance is in principle feasible, the main effort in subsequent work has focused on reducing the additional qubit overhead. A line of work has led to a constant-factor overhead in qubits, first in the instant-classical computation model [Got14], then in the standard circuit model [Gro19], [FGL18], [NP25]. Yet, although the best qubit overhead was reached, the depth overhead was only slightly improved from poly-logarithmic to 'tilde'-logarithmic [NP25]. In fact, it seems that the same means that are used to reduce the qubit overhead are the ones that challenge reducing the depth overhead.

A key concept for reducing qubit overhead is to encode qubits with a high-rate quantum LDPC code. For those codes, except for the control-not, there are no known gates that can be realized in constant depth; furthermore, since the structure of the codes is relatively complicated – based on the skeletons of expander graphs - it seems hard to look for realizations of their logical operators. Without a constant-depth realization of logical operators, any construction that doesn't interfere with the order of the logical operations would admit a non constant-depth overhead.

Contrary to high-rate LDPC codes, topological error codes have a geometrical image that is well understood and, as a result, a wide range of works suggest efficient realization of their logical operators. For example, consider the 3D-color code [Bom15], fold-transversal gates on the surface code [BB24], and realization of the Clifford group on the 2D Toric [GV24].

Another branch of work, whose primary goal was to show that the embedding dimension of the computation can be reduced to 2D [BDL26], and more recently [DST26], uses topological codes as the main error correction code. Those constructions simplify the hardware realization; they suggest that quantum computers can be implemented by nearest-neighbor interactions automata - what is believed to be easier to develop. Yet, as in the threshold theorem [AB99], and in contrast to the works that use high-rate LDPC codes, they rely on a hierarchical correction cycle whose execution time grows with system size. Hence, even though those directions have the potential to overcome the non-constant depth overhead per gate, it seems that the correction cycle imposes a barrier for them to come up with a constant factor in depth-overhead construction.

1.2 Stating Our Results

In this work, we present a fault-tolerance construction that uses the (d, d') -Toric code and has a constant-depth 'correction' cycle, for any $d', d - d' \geq 2$. In other words, the construction uses a relatively simple code - the 4D Toric code, yet doesn't incur the cost of a hierarchical correction cycle.

The computation model we consider is quantum circuits subject to *local-stochastic noise*, where only qubits are subjected to the noise, while classical bits are non-noisy. Classical computation is allowed, and the classical depth matters. Namely, at each time unit, only a constant depth of classical computation can be done. Local stochastic channel is defined to act such that the probability of any subset of locations S (space-time coordinates), which is picked prior to the channel's action, being affected decays exponentially with the subset size.

Definition 1.1 (Local stochastic noise channel). *Let $p \in (0, 1)$. We say that the noise channel $\mathcal{N}_{\mathcal{E}}$ is local stochastic with rate p if for every finite set of locations $S \subset \mathbb{N} \times \mathbb{N}$,*

$$\Pr_{\mathcal{E} \sim \mathcal{N}_{\mathcal{E}}} [S \subset \text{Loc}(\mathcal{E})] \leq p^{|S|}.$$

For circuits $\tilde{C}$ and C , we say that $\tilde{C}$ simulates C with probability p if, when one stops a noisy run of $\tilde{C}$ at a midpoint, one can, with probability p , recover by non-noisy operation an 'image' of a midpoint in the running of C , and if stopping at the last time point of $\tilde{C}$ enables recovery of the 'image' of C at the end of its run. We state our main theorem as follows:

Theorem (Informal version of Theorem 6.3). *There exists a constant $p_0 \in (0, 1)$ such that for any local stochastic noise channel $\mathcal{N}_{\mathcal{E}}$ with rate $p < p_0$, a quantum circuit C with depth μ and width w , and a big enough integer n , there exists a 'logical equivalence' circuit $\tilde{C}$ with depth $\Theta(\mu n^5)$ and width $\Theta((w + \mu)n^8 \text{Polylog}(n))$ that simulates C with probability:*

$$1 - O(\text{poly}(w, \mu, n)) p^{\Omega(n)}$$

Furthermore, $\tilde{C}$ has the following two properties:

1. *After a short initialization stage (at depth $\Theta(\text{Poly}(\log n))$), each of the qubits is encoded in the n -length 4D Toric code.*
2. *An error correction cycle takes $O(1)$ depth, and the gates in each correction cycle don't depend on time. In particular, if C is the identity circuit, then after the initialization stage, the layers in $\tilde{C}$ are the same constant-depth error reduction procedure.*

As mentioned earlier, both code simplicity and constant depth have already been achieved: the former as part of the quantum automation line of work, and the latter through the use of high-rate error correction codes. However, this is the first construction to achieve both simultaneously.

Clifford-free Depth Overhead. An application that results from a slight modification of the main construction and its proof is a fault-tolerant construction that introduces a constant depth overhead per Clifford gate. We name that construction *Clifford-free Depth Overhead*. ?? describes the construction, proves external claims that are needed, and eventually elaborates on how to modify the proof of the main theorem to achieve the desired result.

Theorem (Informal version of ??). *There exists a constant $p_0 \in (0, 1)$ such that for any local stochastic noise channel $\mathcal{N}_{\mathcal{E}}$ with rate $p < p_0$, and a Clifford+**T** quantum circuit C with depth μ , **T**-depth $\mu_{\mathbf{T}}$ and width w , there exists a 'logical equivalence' circuit $\tilde{C}$, with depth $\Theta(\mu) + \Theta(\mu_{\mathbf{T}} n^5)$ and width $\Theta((w + \mu)n^8 \text{Polylog}(n))$ that simulates C with probability:*

$$1 - O(\text{poly}(w, \mu, n)) p^{\Omega(n)}$$

Sweep-rule in High Dimensions. The correction cycle in our construction consists of local operations done in parallel. Concurrently, any vertex of the Toric is thought of as a computation unit that has to decide on a correction to be applied over its local neighborhood. The specific procedure it runs is the sweep-rule, a generalization of Toom's rule invented by Kubica and Preskill [KP19]. The sweep-rule can be sketched informally as follows: The vertex looks for a subset of bits to flip on its 'first quarter'¹, such that the flipping erases the syndrome that is seen on the 'first quarter' and all the bits to be flipped are supported on vertices that support the local syndrome. If there is such a subset of bits, then the vertex flips the bits. In [KP19], it is argued that for small dimensions, any vertex that sees only syndrome on its 'first quarter' has a subset of bits to be flipped that satisfy the sweep-rule requirement. Moreover, the authors conjectured that this is also true for higher dimensions. We prove their conjecture; furthermore, we give an explicit algorithm for finding the subset. We refer to the outline section for an elaborated discussion.

1.3 Discussions, Applications, and Further Directions

Our result can be understood as follow: when the qubit overhead is not a consideration, it means that the 4D Toric code is might be as good as the good quantum LDPC / Hyperproduct code. That translates to the two following applications:

1. New approach for quantum fault tolerance: Our proof is inherently different from prior works, describes in higher resolution the expected error, and might be used as a springboard for a constant depth overhead construction.
2. Memory/Storage components: Protecting a 4D Toric encoded state for an exponentially long time is achieved by performing the same operation at each time unit, suggesting that the 4D Toric code is a promising candidate for memory in the circuit model. Combined with the fact that the 4D Toric code

¹or the positive direction

also functions as memory in thermal noise models [Ali+08], and given its relative simplicity, it appears to be a natural choice for real-world memory applications. We emphasize that an experimental usage of codes in four dimensions has already been demonstrated [Aas+25], [Ber+24], [Rei+24].

2 Outline of the Proof

In studying fault tolerance, the typical method is to first distinguish between configurations that can be tolerated and those that cannot. For example, in von Neumann-style proofs [Neu56], it is common to track the number of flipped bits; if fewer than half are flipped, the logical states can be recovered. In the work of Aharonov and Ben-Or [AB99], a logical bit at the n level admits a logical error if at least a constant fraction of the $n - 1$ level logical bits that compose it admit an error. In the work of [Gro19], a tolerated configuration is one in which all encoded registers can be considered as distributed according to sufficiently weak local stochastic noise. The next step is to show that from a tolerated configuration, we move with high probability to another tolerated configuration. For our goals, each approach presents a challenge: The concatenation method, by definition, requires non-trivial depth computation at each step. Tracking the number of flipped bits fails almost immediately, as topological error correction codes have sublinear distance when a linear number of bits is expected to be flipped. And preserving the local stochastic noise distribution seems to heavily rely on the expansion properties [Gro19] of the underlying structure, where for topological codes, we expect to forge a 'spatial' dependence.

In contrast, Berman and Simon provided a proof that, instead of tracking the dynamics, considers the entire run as a possible outcome, ruling out improbable runs only at the end [BS88]. We outline their idea informally as follows: a logical bit is represented by $n \times n$ bits set on the vertices of the grid. In each correction cycle, the bits determine their value by taking the majority of their current value, and their north and east neighbors values, a rule known as Toom's rule. For simplicity, let's ignore the logical operations. Now, given a run, and in particular the set of space-time locations of faults, consider the following directed graph $\mathcal{H} := (V_{\mathcal{H}}, E_{\mathcal{H}})$ induced by it:

1. If at time t , the bit at vertex v diverges relative to the ideal run, then add (v, t) to the vertices.
2. If at time t , vertex v 'leads' one of its neighbors u to diverge, then connect (v, t) and $(u, t + 1)$ by an 'arrow'.
3. Connect any two vertices in $V_{\mathcal{H}}$, which are connected to a third by 'arrows', by a 'fork'.

Notice that the vertices of $\mathcal{H}$ with no incoming arrows correspond to faults; let's call them leaves. Thus, the probability of each running can be bounded by the number of leaves in the induced graph. Their key lemma states that a cluster of flipped bits with a 'perimeter' s , connected by arrows, induces a subgraph with at least $\Theta(s)$ leaves².

²And number of edges $\leq \Theta(\text{leaves})$.

In other words, the probability of a connected component (by arrows) being flipped decays exponentially with its perimeter.

The graph $\mathcal{H}$ was later named Toom’s contours, a term that, to the best of our knowledge, was coined in [SST26]. We chose to follow their approach. The rest of the section reviews the challenges that we encountered on the way to generalize it and the ideas to overcome them.

The Toric Code, the Sweep-Rule, and the Infinite Plane. The first issue was finding an analog for Toom’s rule. We decided to adopt the sweep rule invented by Preskill and Kubica in [KP19], which they introduced as a local rule for any CSS code induced by a simplicial Toric, equipped with a notion of ‘positive’ and ‘negative’ directions (in their paper ‘future’ and ‘past’). They also provided a rigorous proof that the iterative non-noisy applications of the sweep rule succeed in decoding faults that were sampled independently. For our purposes, we had to extend their proof for cubical Toric³. The main reason for this is that the alignment of the positive/negative directions with the axes direction of $\mathbb{R}^d$ space allows quantifying the ‘perimeter’ of error clusters. Yet, we are sure that our proof can be adjusted to support simplicial Toric complexes as well, and we leave it for further work.

In complete contrast to typical works in topology, where global properties are studied and local characterization is left as general as possible, here we need notations to describe the local view of each vertex of the complex. We needed it for defining the sweep-rule action and stating claims about how it shrinks errors. That led us to the $\theta_{\mathbf{v},S'}$ notation defined in detail in Definition 4.1. In general, we denote by S the generator set of the Toric, and for a subset of generators $S' \subset S$, we define $\theta_{\mathbf{v},S'}$ to be the face whose ‘bottom-left’ corner is the vertex v , and sides of the face correspond to edges on paths obtained by stepping through subsets of S' (for example, if $S' = \{g_1, g_2, g_3\}$ then $\{g_2g_1v, g_3g_2g_1v\}$ is an edge of $\theta_{\mathbf{v},S'}$). Figure 1 shows how the 2D-Toric complex is characterized via $\theta_{\mathbf{v},S'}$ notation.

³The standard Toric codes.



Figure 1: Here we consider a 2D Toric complex. We name the set of generators $S = \{g_1, g_2\}$, where g_1 is the generator that sends us one step in the north direction, and g_2 sends us in the east direction. For vertices v, u, w , we mark in blue three faces that are rooted at them.

The 4D-Toric code is defined as the binary assignment over the 2D-faces (plaquettes), where the Z -stabilizers (parity checks) are set on the edges, namely restrict the Z product of the faces adjacent to each edge, and the X -stabilizers (phase checks) are set on the 4D-faces (hypercubes). Similarly, the (d, d') -Toric code is defined by associating the d' -dimensional faces of the d -dimensional Toric code with the qubits, while associating the $d' \pm 1$ -dimensional faces with the X/Z -stabilizers.

To present the sweep-rule, we need to have a definition of directions. For this section, we will be satisfied with the following hand-waving: For a vertex v , the bits and checks on the first 'quarter' of it (high-dimensional quarter) will be named the positive bits and checks of v .

Definition 2.1 (Sweep Rule (Inforaml)). *Let z be a binary assignment on the Toric faces representing 'error'. The sweep rule decoding procedure at a vertex v is defined as follows:*

Measure the positive checks of v and denote the outcome by σ . Next, look for an assignment $\tilde{z}$ on the positive bits of v such that the vertex support of $\tilde{z}$ is contained in

the vertex support of σ and that the positive checks of v assign to $\tilde{z}$ the same syndrome as σ . If there is such, flip $\tilde{z}$.

See Definition 4.6 for the formal version. A full error correction cycle, or a sweep-cycle, consists of applying the sweep-rule to each vertex and then doing the same in the dual space. Figure 2 provides an example of the result of the sweep rule given a syndrome.

To see the effectiveness of the sweep rule, consider an error that does not close a non-trivial loop on the Toric. Such errors have 'poles', and in particular, one can think of the most 'negative pole'. In the 2D case, it is the most 'bottom-left' vertex that supports the error. For simplicity, let's call that vertex the corner of the error. In [KP19], it was argued that, at the end of a sweep-rule execution, the corner always suggests a non-trivial correction. Thereby advancing the boundary of the error toward shrinking. That argument is easy to verify for low dimensions; for higher dimensions, the authors left it as a conjecture. In this work, we prove that conjecture and, from the proof, derive an algorithm for finding the correction. Nevertheless, the proof doesn't distinguish between simplicial or cubical complexes.



Figure 2: An example of applying the sweep rule within the $(3, 2)$ -Toric code. In the $(3, 2)$ -Toric code, bits are placed on the squares and Z -stabilizers (parity checks) on the edges. The view is restricted to two vertices, u and v . On the left, the unsatisfied positive checks of u and v are highlighted in blue. On the right, the bits to be flipped are marked in green. Given the syndrome, vertex v flips one of its positive faces, whereas for u , there is no subset of positive faces that, when flipped, decreases the number of local unsatisfied checks. Therefore, u flips no bit.

The way we quantify the 'perimeter' of an error cluster is by considering the (hyper)triangle that bounds the vertices supporting the cluster. Now, the poles of an error cluster are the vertices that lie on the sides of the triangle. We followed [BS88], and by embedding the Toric in the infinite plane, we were able to identify the poles as the max/min points of linear functions which we call *potential walls*. Then, we show that the sweep rule advances only one pole - matching the corner mentioned earlier. And hence, advances one side, forward the others. As a consequence, a sweep-cycle decreases the sides of the bounding triangle, and hence the 'perimeter'.

The embedding restricts periodicity conditions on the assignments. We encounter this twice: first, an assignment on the Toric's faces that closes a non-trivial cycle

translates to an infinite assignment in the infinite plane. In that case, no triangle bounds the error, which agrees with the fact that a logical error—in that context, an error that changes the logical states—should be ‘decoded’ to a different codeword. From a technical perspective, it comes into effect as the claims about shrinking assume that a given assignment is a connected finite assignment (or a collection of connected components such that each of them is finite). The second time we have to deal with the mapping onto the infinite plane is in the calculation of the probability of a given Toom’s contour to occur. Since different leaves (faults) of Toom’s contour might be associated with the same qubit⁴ and hence dependent, to bypass this, we took from [Gac21] the claim that states any such graph has a subgraph whose leaves are independent, and their number is at least a constant factor of the original leaves.

⁴Actually, the leaves are vertices that support the face associated with a faulty qubit.



Figure 3: A graphical example of the maximal values of the potential walls. Here we consider the $(2, 1)$ -Toric code, namely the bits are set on the edges where the Z -stabilizers (parity checks) are set on the vertices. The grid is a finite 6×6 patch of the infinite 2D plane. The edges marked in blue correspond to a bits assignment z . Each of the dashed lines marks the line of vertices whose L_i value converges with the maximal value reached by a vertex in z . Moving into the infinite plane enables us to determine those lines and draw the triangle that bounds z .

See Section 4.2 for embedding into the infinite plane, Section 4.4 for the proof of shrinking, and Section 7.5 for the claim about the existence of a sufficient subgraph.

The Intermediate Argument. Now that we have the code and decoder, the next difference we have to address is how to classify an error. Unlike the classical regime, flipping a stabilizer's qubits doesn't count as a fault. Moreover, the sweep-rule discussed above doesn't shrink stabilizers (as it acts trivially on them). Therefore, we need to redefine what is considered a divergence from the ideal computation. To overcome it, we define the reduced form of an error to be its translation by stabilizer that minimizes the 'perimeter' of the bounding triangle. Then we prove that the poles

of the error in it's reduced coalesce with the poles of the syndrome it induces. In particular, the proof that the sweep-rule shrinks an error (Section 4.4) assumes that the given error is in its reduced form. As a result, instead of defining the vertices of Toom's contour graph to be the flipped vertices as in [BS88], we define them to be qubits⁵ in the reduced form of the error.

We then face another challenge: if we analyze Toom's contour, which is induced by the single time unit running consisting of flipping all the qubits at once, then in the layer following the fault cycle, even though the error contains many qubits, its reduced form is empty. In the language of Toom's contour, as will be defined later, the Toom's contour is empty, and no bounds on the event's probability to occur can be inferred using it. More generally, there is a threshold, such that for errors exceeding that threshold, one cannot restore the history of the error. In fact, the threshold is far lower than flipping all the qubits and matching errors that close a non-trivial cycle.

We address this by introducing an analysis framework where all computation, including faults, is imagined to occur sequentially — in steps. Within this framework, we can prove an *intermediate claim* stating that on the path to an error configuration closing a non-trivial cycle, one must pass through a step with an error where the 'perimeter' its reduced form is large. Thus, a sufficient condition for a logical error to occur in some time unit is the existence of a step with a reduced error that is large. Moreover, having such a step implies the existence of a step for which it happens for the first time, meaning that prior to that, no error has closed a non-trivial cycle, the sweep-rule worked as it should, and a Toom's contour describing the history of the error can be constructed. Thus, an upper bound on the probability of the condition can be derived. In the proof of the theorem, we show that this condition happens with negligible probability. To elaborate on the issue, we first need to present the computation model.

We model a noisy running of a given circuit C as follows: first, we sample space-time locations for the faults and denote the outcome by $\mathcal{E}$. Then, we impose layers of computation alternately, such that the even and odd layers are taken from $\mathcal{E}$ and C according to their order. The combination of the circuit C with a subset of faults $\mathcal{E}$ defines a possible running, and we call it a subsection of C to $\mathcal{E}$ and denote it by $C[\mathcal{E}]$. Furthermore, we showed in Claim 3.8 that for a local-stochastic noise model, unbalancing partition to noisy and logical layers, for example placing a noisy layer after every three logical layers of C , is equivalent to odd-even alternations up to a rescaling of the noise rate (by a constant). We use that since the error correction cycle of the sweep rule takes more than a single layer.

As mentioned, we are going to analyze the likelihood of getting logical errors by looking back in time and asking what a run that has led to them is supposed to look like. The sweep rule ensures that an error cluster below a certain size, sufficiently low to not closing a non-trivial cycle, shrinks. The problem is that a single fault layer (i.e., a layer whose source is in $\mathcal{E}$) might increase the error cluster above that size. At first glance, one might be tempted to show that the expansion of error clusters beyond a fraction of the code distance happens with negligible probability, but that calculation

⁵In fact: vertices on the qubits' support, see Section 5.2.




Figure 4: Example of a subjected quantum circuit: the even layers (0, 2, 4) contain the original gates, while the faults are imposed on the odd layers and marked in red.

soon becomes complicated when considering the merging of disjoint clusters. Here comes the splitting of the computation into steps to advantage. In each step, only one gate or one fault is applied, and steps are ordered such that they agree with the computation. Figure 4 illustrates a quantum circuit and a graph of steps corresponding to it.

Since the gates and the faults have a bounded fan-in/out⁶ In each step, the maximal size of an error cluster can increase by at most a constant factor, with the maximal increase occurring when multiple clusters are merged. Equipped with that insight, we are able to infer the intermediate argument.

Section 3 formalize the computation model and the computation graph presentation (splitting to steps). In addition, see ?? for the formal proof of the intermediate argument.

Description of the Construction. Similarly to other fault tolerance approaches, we use the concatenation construction [AB99] in order to initialize zeros and resource states encoded in the 4D-Toric code, as referred to in [Gro19], [Got14]. From there, we assume that all the quantum registers are encoded in the 4D-Toric code, and we name each of the encoded block a quantum register. The resources states are spread across two registers.

Then, after the initialization step, the logical cycles are divided into error correction, which is a single sweep-cycle, and 'logical operation' cycle. For our universal logical gate set, we choose **X**, **Z**, **CX** (controlled NOT), **H**, **S**, and **T** gates. Except for the Paulis and the 'Controlled NOT' gate, which are realized by transversal Paulis and transversal wire-wise physical 'Controlled NOT' gates, all the others are implemented via gate teleportation [ZLC00].

The general realization by gate teleportation is done as follows: first, we XOR the target state into a resource state via **CX**, measure the input state and the first register of the resource state with a Bell measurement, the input register in the *X* basis, and the first register of the resource state in the *Z* basis. Then, on the classical outcomes,

⁶The faults are bounded-fan because we can assume that in the time units representation they are Pauli products and therefore in the steps presentation they are decomposed into single qubit Pauli matrices.

run a classical decoding procedure to extract the logical codeword. Meanwhile, on the second register of the resource, which is intended to contain the result of the action, perform a sweep cycle. At the end of the decoding, depending on the result, we apply either a transversal Pauli operator or the **S** gate (again by teleportation) depending on the original gate that was asked to simulate. The gate teleportation lasts for multiple logical cycles, depending on the depth required by the classical decoding procedure.

We emphasize that for measuring the input register in the X -basis, we apply the Hadamard through its fold-transversal realization [BB24], so the measurement is done in constant depth. However, the reason we don't implement the Hadamard by fold-transversal from the beginning is that it introduces an expansion in the bounded-triangle by the fold-transversal realization. Since it guarantees that after the measurement the register will not observe more faults, we can tolerate a one-time blowup in the size of the error and still be able to decode the logical state in the register correctly. Yet, for the Clifford-free depth overhead, we do use the fold-transversal implementation followed by sufficiently many followed sweep-cycles. As a result, we stretch the depth of the basic logical layer, which in turn comes into effect by a lower threshold for the noise rate.

In Section 4.5, we discuss in detail the structures of the logical codeword of the (d, d') -Toric code and the realization of its logical gates.

Toom's Contours. A sufficiently simple definition of Toom's contours in [BS88] would be a directed graph that describes a noisy run with the additional marking of diverged bits. In our case—the quantum case—the vertices of Toom's contour are neither bits nor checks; instead, they are the vertices on the Torices that 'diverged'. A proper conceptual definition would be:

*A graph that describes a noisy run from the Torices' perspective.
In particular, the vertices of the graph are the vertices of the
Torices, equipped with a time coordinate, that witness a faulty
qubit in a reduced form.*

We adjust their construction as follows: Denote the graph by $\mathcal{H} := (V_{\mathcal{H}}, E_{\mathcal{H}})$ and by $P_1, \dots, P_t$ a *Pauli-path*, such that P_i is the propagated error up to time i (including faults that occur at that time), and each 'cluster' of error in P_i is in its reduced form. Then,

1. If at time t , a vertex v on one of the Toric complexes⁷ corresponding to one of the quantum registers supports a face-qubit- that belongs to a reduced error, then add (v, t) to the vertices.
2. If at time t , vertex v belongs to a reduced error, and at time $t + 1$ its neighbor u belongs to a reduced error, then we think of the situation as ' v "leads" to u ' and connect (v, t) and $(u, t + 1)$ with an "arrow." The vertices v and u might belong to different registers (which matches the case that at time t we apply an operator on both registers). For convenience, we don't require u at time t to belong only to non-faulty qubits.

⁷Actually the infinite grid.

3. As in the classical case, connect any two vertices in $V_{\mathcal{H}}$ that are connected to a third by 'arrows' with a 'fork'.



Figure 5: Above: The application of a sweep-cycle, where the green faces represent the bits assignment z , and the blue edges correspond to unsatisfied checks $\partial^- z$. The vertices in the reduced form are labeled as $v_1, v_2, \dots, v_8$. Notice that the leftmost bottom-inner vertex is not in the reduced form even though it belongs to z . Below: The Toom's contour derived from the single computation layer that matches the cycle. Arrows are colored in blue, while forks are colored in green. Note that $(v_1, t+1)$ doesn't belong to the Toom's contour (since v_1 doesn't belong to a reduced error cluster at time $t+1$).

Later, we incorporate the 'splitting into turns' idea into the formation of Toom's counter $\mathcal{H}$ -graphs, by gluing the layers corresponding to the first $t-1$ layers with a final layer that matches a partial execution of the t th time unit. The layer encodes the situation of the running up to a certain step. Then, combining with the intermediate theorem, we have that any running that admits a logical error derives a Toom's contour, with sufficiently large vertices at final layer.

The rest of the analysis follows [BS88]. We define *slices* at time t to be connected (by arrows) components of vertices at time t . Then, in the merging of two slices into one, it is guaranteed that faults induced by each are disjoint; otherwise, those two slices will be connected by arrows, namely - the same slice. For a *slice*, we define its size to be the perimeter of the triangle that would bound the vertices of the slice if they were projected on the same Toric. Then, by induction, we prove that Toom's contour that induces at time t a slice with size s has a spanning tree with at most and at least a constant times s edges and leaves (and therefore faults). That allows us

to bound the probability of such a slice occurring by counting the number of possible spanning trees and assigning them the probability that the minimal number of faults occurs.

One technical change that we made and haven't discussed yet is that in the gluing of the last layer, which was defined to capture a partial layer execution up to a turn, we connect by arrows any two vertices that were added to $V_{\mathcal{H}}$ and share a check. By doing so, we ensure that connected components in the bits-to-checks graph would also be connected by arrows in Toom's contours. That is necessary since, without the last gluing, an error cluster that closes a non-trivial loop on the Toric might consist of two 'connected by arrows' error clusters which are tangent. We emphasize that such edges connection can be made only in the last stage, as it doesn't pass the induction step. However, the final slice that is formed passes the desired inequalities only by constants and is therefore sufficiently good to derive the probability bound.

Section 5.2 defines Toom's contours, Section 5.3 defines the slices and states a claim connecting a slice to its predecessors, and Section 5.4 provides the induction proofs.

3 Notations and Preliminaries

We follow the standard notations commonly used in the literature. For an integer n , we denote by $[n]$ the set of integers $\{1, 2, \dots, n\}$. To present a countable set, which might or might not be finite, where the number of elements does not matter, we use curly brackets with a subscript denoting the first index of the set. For example, $\{u_i\}_{i=1}$ represents the set $u_1, u_2, \dots$. For finite sets, we might add a superscript to present the size of the set. If the first index is clear from the context, then we might omit the assignment. We use a colon inside curly brackets to specify a set of elements satisfying a condition, i.e., for sets A, B , the notation $\{a : a \in A, a \in B\}$ equals $A \cap B$. Curly brackets without subscripts or conditions refer to singletons. For example, for a given integer i , the set $\{u_i\}$ is the set that contains only the element u_i . The size of a set A is denoted by $|A|$. For an integer i , we denote by $\mathbf{1}_i$ the vector consisting of one at the i th coordinate and zero elsewhere. The char $|$ symbolizes a restriction, for example $x|_y$, is the restriction of the 'entity x ' to ' y '. The particular behavior of the restriction would change depending on the target type.

When not otherwise mentioned, we consider quantum states over qubits, each lying in a two-dimensional Hilbert space $\mathcal{H}_2$, and denote by $\mathcal{H}_{2^n}$ the Hilbert space of an n -qubit system. The $\otimes$ operator stands for tensor product, for example: $\mathcal{H}_{2^n} = \mathcal{H}_2^{\otimes n}$. The space of linear operators on the Hilbert space is denoted by $L(\mathcal{H})$. Pure quantum states and their dual presentations are denoted by $|\psi\rangle, \langle\psi|$, and in the matrix representation, we use $|\psi\rangle\langle\psi|$ to denote its density matrix. We will use ρ to denote a mixed state. The action of an operator A on a mixed state is written as $A \circ \rho$. For example, if A is a unitary operator, then $A \circ \rho = A\rho A^\dagger$. The letters, X and Z , will usually denote the single qubit X and Z Pauli operators. Yet, where the Pauli matrices are out of context, we use these letters for other variables. The Pauli group itself, over n qubits, is denoted by $\mathcal{P}^{\otimes n}$, and when the number of qubits is either clear or doesn't matter,

we omit the superscript. For a Pauli $P \in \mathcal{P}$, we denote by $X(P)$ and $Z(P)$ the X - and the Z -components of P such that $P = X(P)Z(P)$. The gates 'Controlled Not' and 'Controlled Z' are denoted by $\mathbf{CX}$ and $\mathbf{CZ}$, and the gates Hadamard, S, and T are denoted by:

$$\mathbf{H} = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}, \mathbf{S} = \begin{bmatrix} 1 & 0 \\ 0 & i \end{bmatrix}, \mathbf{T} = \begin{bmatrix} 1 & 0 \\ 0 & e^{i\frac{\pi}{4}} \end{bmatrix}$$

For a quantum gate U , we denote by $|U\rangle$ the resource state that can be used to simulate U . Notice that, in general, the resource state is not unique. In this work, we address specific states that will be discussed later. The function $\mathbf{Pr}[\cdot]$ is the probability function. In the language of groups, for a group $\mathcal{G}$ and elements $g_1, \dots, g_l$, we denote by $\langle \{g_1, \dots, g_l\} \rangle$ the subgroup generated by $g_1, \dots, g_l$.

Throughout the paper, we will deal with CSS codes [CS96], [Ste96]. For a linear subspace A , we use $|A\rangle$ to denote the uniform superposition over the vectors in A . A CSS code will usually be denoted by the pairs $\mathcal{C}_X, \mathcal{C}_Z$ such that the Z -stabilizers match the codewords of $\mathcal{C}_Z^\perp$, and the logical codewords in the computational basis are the elements in $\mathcal{C}_X/\mathcal{C}_Z^\perp$. The logical zero codeword is $|\mathcal{C}_Z^\perp\rangle$. We denote the equivalence relation up to Z -stabilizer addition by $\sim_{\mathcal{C}_X}$. Namely, for $x_1, x_2 \in \mathbb{F}_2^n$, we have $x_1 \sim_{\mathcal{C}_X} x_2$ if $x_1 + x_2 \in \mathcal{C}_Z^\perp$. By the same way we define $\sim_{\mathcal{C}_Z}$.

The Computation Model. Informally, a quantum circuit is a placement of a collection of quantum gates, taken from a finite set, at different space-time locations. Each such location records when the gate is applied and on which qubits it acts. Formally:

Definition 3.1 (Quantum Circuit). *A space-time location is a pair $l = (t, S)$, where t is a time coordinate and $S = (s_1, \dots, s_\Delta)$ is an ordered tuple of qubit indices. For a finite collection of quantum gates*

$$\mathcal{G} \subset \left\{ g: L(\mathcal{H}_2^{\otimes \Delta}) \rightarrow L(\mathcal{H}_2^{\otimes \Delta}) \right\},$$

and integers $w, d \in \mathbb{N}$, a quantum circuit C of depth d and width w is a collection of tuples

$$C = \{u_i = (g_i, l_i)\}_i,$$

where each $g_i \in \mathcal{G}$ and each location $l_i = (l_{i,1}, l_{i,2})$ belongs to $[d] \times [w]^\Delta$. We call the tuples u_i the gates of C . They must satisfy the following compatibility condition: if two gates have the same time coordinate, then their supports are disjoint. Namely, for $u_i = (g_i, l_i), u_j = (g_j, l_j) \in C$, the equality $l_{i,1} = l_{j,1}$ implies that the sets of qubit indices appearing in $l_{i,2}$ and $l_{j,2}$ are disjoint. We denote by

$$\text{Loc}(C) := \{l_i : u_i = (g_i, l_i) \in C\}$$

the set of locations occupied by the gates of C . We define the t -layer of C to be its restriction to gates at time t :

$$C|_t := \{u_i = (g_i, l_i) : u_i \in C \text{ and } l_{i,1} = t\}.$$

Later, we will restrict $\mathcal{G}$ to contain **CX**, **H**, **S**, **T**, and measurements. To define the action of a quantum circuit C on a given mixed state, we first associate C with a directed acyclic graph.

Definition 3.2 (Associated Computation DAG). *Let $C = \{u_i\}_i$ be a quantum circuit with depth d and width w . We define the associated computation DAG of C , denoted by π_C , to be the directed acyclic graph $\pi_C = (\mathcal{U}, E)$ as follows:*

1. *The vertex set is the gate set of C , namely $\mathcal{U} = \{u_i\}_i$.*
2. *For $u_i = (g_i, l_i)$ and $u_j = (g_j, l_j)$, there is a directed edge $(u_i, u_j) \in E$ if and only if j is one of the minimizers of*

$$\{l_{k,1} : (g_k, l_k) \in \mathcal{U}, l_{k,1} > l_{i,1}, l_{k,2} \cap l_{i,2} \neq \emptyset\}.$$

Equivalently, u_j is one of the earliest later gates whose support intersects the support of u_i . There may be more than one such minimizer, so the out-degree may exceed one.

Suppose that $J = (j_1, j_2, \dots, j_k)$ is an ordering of the gate indices that agrees with a topological order induced by π_C . Then we define the associated partial circuits $C_J = \{C_{J,i}\}_{i=1}^k$ by

1. $C_{J,1} = \{u_{j_1}\}.$
2. *For $1 \leq i < k$, set $C_{J,i+1} = C_{J,i} \cup \{u_{j_{i+1}}\}.$*

We call $C_{J,i}$ the partial computation circuit up to step i .

Definition 3.3. *Let ρ be a density matrix over w qubits, and let $C = \{u_i\}_{i=1}$ be a quantum circuit with depth d and width w . For any density matrix σ , quantum gate $g \in \mathcal{G}$, and location l , the application of the gate-location tuple (g, l) on σ , denoted by $(g, l) \circ \sigma$, is the application of g on σ where wires are ordered such that the first qubit in the input to g is $l_{2,1}$, the second is $l_{2,2}$, and so on. The result of the application of C on ρ , denoted by $C \circ \rho$, is given by iterative applications of the gates $\{u_i\}_{i=1}$ on ρ , according to a topological order induced by π_C .*

We say that $\rho_1, \rho_2, \dots, \rho_k$ is a computation prefix if $\rho_1 = \rho$ and $\rho_{i+1} = u'_i \circ \rho_i$, where $\{u'_i\}_{i=1}$ is a prefix of a topological ordering of $\{u_i\}_{i=1}$. If $\{u'_i\}_{i=1}$ contains all the gates in $\{u_i\}_{i=1}$, then we call $\rho_1, \rho_2, \dots, \rho_k$ a running.

For a quantum circuit C , we would like to discuss its action in the presence of different outcomes of faults. The following definition encodes the running in the presence of a specific set of faults $\mathcal{E}$.

Definition 3.4 (C subjected to $\mathcal{E}$ every m computational layers). *Let $C = \{u_i\}_i$ be a quantum circuit with depth d and width w , let $m \in \mathbb{N}$, and set*

$$d_m := \left\lceil \frac{d}{m} \right\rceil.$$

Let $\mathcal{L} \subset [d_m] \times [w]$ be a set of locations, and let $\mathcal{E}$ be a quantum circuit, such $\text{Loc}(\mathcal{E}) = \mathcal{L}$. We name $\mathcal{E}$ subsets of faults, and we name the layers of $\mathcal{E}$ fault layers. The circuit C subjected to $\mathcal{E}$ every m computational layers, denoted by $C[\mathcal{E}]_m$, is obtained by setting the first layer to be the first layer of $\mathcal{E}$, then placing in order the layers of C and inserting one fault layer after every m consecutive C 's layers. Thus the second fault layer is inserted after the layers $1, \dots, m$ of C , the third is inserted after the layers $m+1, \dots, 2m$ of C , and so on. Concretely, each gate $u_i = (g_i, (t_i, S_i))$ of C is moved to the time coordinate

$$t_i + \left\lfloor \frac{t_i - 1}{m} \right\rfloor,$$

while each fault gate $(g, (t, S)) \in \mathcal{E}$ is placed at time coordinate

$$t(m+1).$$

When $m = 1$, the gates occupy the even time coordinates and the fault layers occupy the odd time coordinates. When the parameter is omitted, we mean $m = 1$. We denote the associated computation DAG of $C[\mathcal{E}]_m$ by $\pi_{C[\mathcal{E}]}$.

Remark 3.5. Formally, a fault circuit is just a particular quantum circuit whose gates are interpreted as faults rather than as computational operations. Thus the distinction between a quantum circuit and a fault circuit is semantic rather than structural: both are circuits in the sense of Definition 3.1, but they play different roles in the analysis.

Example 3.6. When $m = 3$, the 'computational' layers are sent to the times

$$2, 3, 4, 6, 7, 8, 10, 11, 12, \dots,$$

while the fault layers are sent to the times

$$1, 5, 9, 13, \dots$$

Definition 3.7 (Noise channel and local stochastic noise channel). A noise channel $\mathcal{N}_{\mathcal{E}}$ is a distribution over sets of fault locations. Let $p \in (0, 1)$, we say that the noise channel $\mathcal{N}_{\mathcal{E}}$ is local stochastic with rate p if for every finite set of locations $S \subset \mathbb{N} \times \mathbb{N}$,

$$\Pr_{\mathcal{E} \sim \mathcal{N}_{\mathcal{E}}} [S \subset \text{Loc}(\mathcal{E})] \leq p^{|S|}.$$

So, $\mathcal{N}_{\mathcal{E}}$ defines a probability space. The set of faults $\mathcal{E}$ is an atomic outcome, and for a given circuit C , the subjected circuit $C[\mathcal{E}]$ is a random variable. The following claim states that for local stochastic noise, the subsection of every 1-layer is equivalent to the subsection of every m -layer, up to a rescaling of the noise rate. Here, by equivalence, we mean that the actions of the subjected circuits are the same⁸.

We give the proof in the appendix for several reasons. First, the argument is not the core of this work. Second, versions of it appear in almost any fault tolerance work. Additionally, it relies on 'propagation of entire layers,' a term that conflicts with the 'propagation to a step' (Definition 5.25) needed for the intermediate theorem.

⁸For a given input (mixed) state ρ , they yield the same output

Claim 3.8. *Let $\mathcal{N}_{\mathcal{E}}$ be a local stochastic noise channel with rate p , and let C be a quantum circuit whose gates have arity at most Δ . Fix an even integer c . Then there exists a local stochastic noise channel $\mathcal{N}_{\mathcal{E}_q}$ with rate*

$$q := \left(2^{H(\Delta^{-2c})\Delta^{2c}} p \right)^{\Delta^{-c}},$$

such that for any given state ρ it holds that:

$$C[\mathcal{E}]_1 \circ \rho = C[\mathcal{E}_q]_{c/2} \circ \rho$$

In particular, up to a constant depending only on Δ and c , the rate q behaves like $p^{\Delta^{-c}}$.

Equipped with this claim, we are allowed to assume having constant depth cycles that don't observe faults.

4 The (d, d') -Dimensional Toric Code

The main error correction code that we will use is the Toric code invented by Kitaev. In the 2D case, the code is obtained by taking a Toric, tiling it with squares, setting qubits on the edges, and checks as follows: For every vertex, the parity of the qubits adjoining it is even, and for any square, the parity of the qubits lying on its sides, in the phase basis, is even. A high-dimensional version can be obtained by taking a high-dimensional Toric, encoded as a cubical complex, and then setting the X -checks, the qubits, and the Z -checks to be faces at sequenced dimensions. Throughout this paper, qubits will be set on d' -dimensional faces, and the X -checks and Z -checks on the $d' \mp 1$ dimensional faces. For a review, we refer to [Kit97], [Den+02].

For defining a decoder and proving its 'correctness,' we will have to overload notation. We derive the Toric complex from the skeleton of the Cayley graph of the cyclic group, using the generating subset as a collection of positive directions. The $\theta_{v, S'}$ encodes a face, one vertex v that serves as its 'root' and a subset of directions S' . Formally:

Definition 4.1 (d -Dimensional Toric Complex). *For integers d and n , denote the group resulting from the d -directed power of the cyclic group of order n by $\mathcal{G}_n^d = \mathbb{Z}_n^{\times d}$, and denote the standard generating set of $\mathcal{G}_n^d$ by $S = \{1_i : i \in [d]\}$. We define the d -dimensional Toric complex to be the hypergraph obtained from the Cayley graph $\text{Cay}(\mathcal{G}_n^d, S)$ by taking its vertices as the 0-faces and for any $i \in [d]$ we define the i -faces as follow: First, for a vertex $v \in \mathcal{G}_n^d$ and for a subset of generators $S' \subset S$ of size $|S'| = i$, the i -face rooted at v and spanned by S' , denoted by $\theta_{v, S'}$, is:*

$$\theta_{v, S'} := \bigcup_{I \subset S'} \left\{ \left(\prod_{g \in I} g \right) v \right\}$$

The i -faces of the complex are the collection of all the possible rooted spanned i -faces induced by $\text{Cay}(\mathcal{G}_n^d)$. We will denote them by $\mathcal{F}_i$. Additionally, let us define the multiplication of the generators subset $S' \subset S$ by a vertex $v \in V(\mathcal{G})$ as follows:

$$S'v = \left(\prod_{g \in S'} g' \right) v$$

Finally, we name n and d , the side length and the dimension of the complex.

Next, we formalize binary assignments on the Toric complex and introduce boundary operators. These boundary operators are the same as standard boundary operators in the chain complex formalism.

Definition 4.2 (Assignments on a Toric complex). *Let $\mathcal{T}$ be a d -dimensional Toric complex. For each i , the set of i -faces induces a vector space $\mathbb{F}_2^{|\mathcal{F}_i|}$. An element $z \in \mathbb{F}_2^{|\mathcal{F}_i|}$ is called an assignment on the i -faces of $\mathcal{T}$.*

For a vertex v and a subset of generators $S' \subset S$ of size $|S'| = i$, let $\theta_{\mathbf{v}, S'}$ denote the standard basis vector corresponding to the i -face $\theta_{v, S'}$. Thus the family

$$\{\theta_{\mathbf{v}, S'} : v \in \mathcal{F}_0, S' \subset S, |S'| = i\}$$

is the standard coordinate basis of $\mathbb{F}_2^{|\mathcal{F}_i|}$.

Let z be an assignment on the i -faces. Let $v_1, \dots, v_k \in \mathcal{F}_0$ and $S'_1, \dots, S'_k \subset S$ be such that $\theta_{v_1, S'_1}, \theta_{v_2, S'_2}, \dots, \theta_{v_k, S'_k}$ are precisely the i -faces in the support of z , namely:

$$z = \sum_{i=1}^k \theta_{\mathbf{v}_i, S'_i}$$

We refer to $\theta_{v_1, S'_1}, \theta_{v_2, S'_2}, \dots, \theta_{v_k, S'_k}$ as the collection of faces that form z , or briefly, the faces form z . For $i > 0$, we define the linear map $\partial^- : \mathbb{F}_2^{|\mathcal{F}_i|} \rightarrow \mathbb{F}_2^{|\mathcal{F}_{i-1}|}$ such that

$$\partial^- \theta_{\mathbf{v}, S'} := \sum_{g \in S'} (\theta_{\mathbf{v}, S'/g} + \theta_{\mathbf{g}\mathbf{v}, S'/g})$$

For $i < d$, we define the linear map $\partial^+ : \mathbb{F}_2^{|\mathcal{F}_i|} \rightarrow \mathbb{F}_2^{|\mathcal{F}_{i+1}|}$ such that

$$\partial^+ \theta_{\mathbf{v}, S'} := \sum_{g \in S/S'} (\theta_{\mathbf{v}, S' \cup g} + \theta_{\mathbf{g}^{-1}\mathbf{v}, S' \cup g})$$

We define the restriction of z at vertex v to be the linear map $|_v : \mathbb{F}_2^{|\mathcal{F}_i|} \rightarrow \mathbb{F}_2^{|\mathcal{F}_i|}$ such that:

$$\theta_{\mathbf{u}, S'}|_v := \begin{cases} \theta_{\mathbf{u}, S'} & v \in \theta_{\mathbf{u}, S'} \\ 0 & \text{else} \end{cases}$$

We refer to the faces rooted at v as the positive faces of v , and define the positive restriction at vertex v to be the linear map $|_{v+} : \mathbb{F}_2^{|\mathcal{F}_i|} \rightarrow \mathbb{F}_2^{|\mathcal{F}_i|}$ such that:

$$\theta_{\mathbf{u}, \mathbf{S}'}|_{v+} := \begin{cases} \theta_{\mathbf{u}, \mathbf{S}'} & u = v \\ 0 & \text{else} \end{cases}$$

Two i th dimensional faces will be said to be connected if their intersection is $i-1$ th dimensional face. So, an assignment z is connected if the graph that is induced by taking the faces that support z as vertices and defining the edges to be the adjacency relation between them is connected. We define the connected components of z to be its subsets such that each of them is a connected assignment. For a j -face f , we say that $f \in z$ if z has a support on $\theta_{\mathbf{v}, \mathbf{S}'}$ such that $f \in \theta_{\mathbf{v}, \mathbf{S}'}$. Consider an assignment $z \in \mathbb{F}_2^{\mathcal{F}_i}$ and a vertex $v \in \mathcal{G}_n^d$. We say that z is a rooted assignment at vertex v if there are subsets $S'_1, \dots, S'_k \subset S^{\times k}$ such that each of S'_j is at size i and z is decomposed as:

$$z = \sum_j^k \theta_{\mathbf{v}, \mathbf{S}'_j}$$

Put differently, z is the result of taking a collection of i -faces that are rooted at v .

We are ready to give a formal definition of the (d, d') -Toric code.

Definition 4.3 ((d, d') Toric Code). Let d, d' be integers such that $1 \leq d' \leq d$. For any integer n , we construct the n th instance of the quantum code family (d, d') Toric Code as follows: Let $G = (\mathcal{F}_0, \dots, \mathcal{F}_d)$ be the d -dimensional Toric complex obtained from $\mathcal{G}_n^d$ as defined in Definition 4.1. We associate the (q) bits with the $\mathcal{F}_{d'}$, the X -checks with the $\mathcal{F}_{d'-1}$, and the Z -checks with the $\mathcal{F}_{d'+1}$. An X -check, c_x , is satisfied if the parity of (q) bits set on the d' -face containing c_x is even. Similarly, a Z -check, c_z , is satisfied if the parity, in the Hadamard basis, of the (q) bits set on the d' -faces contained in c_z is even.

We argue that the (d, d') -Toric code is indeed a CSS, and prove that by showing that boundary operators $\partial^\pm : \mathbb{F}_2^{|\mathcal{F}_i|} \rightarrow \mathbb{F}_2^{|\mathcal{F}_{i\pm 1}|}$ induce a chain complex.

Claim 4.4 (CSS structure). The (d, d') -Toric code is a CSS code.

Proof. We will show that $\mathbb{F}_2^{|\mathcal{F}_i|} \xrightarrow{\partial^\pm} \mathbb{F}_2^{|\mathcal{F}_{i\pm 1}|}$ are a chain complex and its dual chain. First, we show that $(\partial^+)^T = \partial^-$. Consider vertices $v, u \in V$ and two generator subsets $A, B \subset S$. Consider the basis elements defined by them: $\theta_{\mathbf{v}, \mathbf{A}}$ and $\theta_{\mathbf{u}, \mathbf{B}}$. Suppose that:

$$(\partial^+ \theta_{\mathbf{v}, \mathbf{A}}) \theta_{\mathbf{u}, \mathbf{B}} = 1$$

Means that there exists $g \notin A$ such that $A \cup \{g\} = B$ and either $v = u$ or $g^{-1}v = u$. But that is equivalent to saying that there is $g \in B$ such that $B \setminus \{g\} = A$ and either $v = u$ or $v = g^{-1}u$. In other words:

$$\begin{aligned} (\partial^+ \theta_{\mathbf{v}, \mathbf{A}}) \theta_{\mathbf{u}, \mathbf{B}} &= \theta_{\mathbf{v}, \mathbf{A}} (\partial^- \theta_{\mathbf{u}, \mathbf{B}}) \\ \Rightarrow (\partial^+)^T &= \partial^- \end{aligned}$$

So it's left to show that the boundary of a boundary is zero. Consider a basis vector $\theta_{\mathbf{v}, \mathbf{S}'}$:

$$\begin{aligned} \partial^- (\partial^- \theta_{\mathbf{v}, \mathbf{S}'}) &= \partial^- \left(\sum_{g \in S'} \theta_{\mathbf{v}, \mathbf{S}'/g} + \sum_{g \in S'} \theta_{g\mathbf{v}, \mathbf{S}'/g} \right) \\ &= \sum_{\substack{g', g \in S' \\ g' \neq g}} (\theta_{\mathbf{v}, \mathbf{S}'/\{g, g'\}} + \theta_{g'\mathbf{v}, \mathbf{S}'/\{g, g'\}} + \theta_{g\mathbf{v}, \mathbf{S}'/\{g, g'\}} + \theta_{g'g\mathbf{v}, \mathbf{S}'/\{g, g'\}}) \\ &= 0 \end{aligned}$$

Where we used that the term in the closure is invariant to sweeping between g and g' , and therefore any term in the sum appears twice. Similarly:

$$\begin{aligned} \partial^+ (\partial^+ \theta_{\mathbf{v}, \mathbf{S}'}) &= \partial^+ \left(\sum_{g \in S/S'} \theta_{\mathbf{v}, \mathbf{S}' \cup g} + \sum_{g \in S/S'} \theta_{g^{-1}\mathbf{v}, \mathbf{S}' \cup g} \right) \\ &= \sum_{\substack{g', g \in S/S' \\ g' \neq g}} (\theta_{\mathbf{v}, \mathbf{S}' \cup \{g, g'\}} + \theta_{g'^{-1}\mathbf{v}, \mathbf{S}' \cup \{g, g'\}} + \theta_{g^{-1}\mathbf{v}, \mathbf{S}' \cup \{g, g'\}} + \theta_{g'^{-1}g^{-1}\mathbf{v}, \mathbf{S}' \cup \{g, g'\}}) \\ &= 0 \end{aligned}$$

The conditions $(\partial^\pm)^{\circ 2} = 0$ imply that $\mathbb{F}_2^{|\mathcal{F}_i|} \xrightarrow{\partial^\pm} \mathbb{F}_2^{|\mathcal{F}_{i\pm 1}|}$ are chain complexes. Therefore, for any d' , every $d' - 1$ face and every $d' + 1$ face overlap on an even number of d' -faces, so the corresponding X - and Z -checks commute. This is exactly the CSS condition. $\square$

The fact that $\partial^\pm : \mathbb{F}_2^{|\mathcal{F}_i|} \rightarrow \mathbb{F}_2^{|\mathcal{F}_{i\pm 1}|}$ induces a chain complex motivates referring to the collection of qubits/checks/vertices as chains of various dimensions. Next, we give a definition of the 'positive' vertices that a vertex v sees, which lie on a chain z .

Definition 4.5 (Positive Neighborhood). *For a vertex v and a chain z , we use the notation $\mathbf{Str}_{0+}(v, z)$ to denote the vertices in z that are connected to v through a stepping via generator, namely all the vertices that support a face in z of the form gv for some $g \in S$. We think of $\mathbf{Str}_{0+}(v, z)$ as the positive neighborhood of v in z .*

4.1 Sweep Rule

The sweep-rule is the core of our construction. A decoding cycle begins with each vertex v reading the syndrome on the X -checks rooted at v (the positive X -checks of v) and, based on what it sees, applying a Pauli correction supported on the qubits rooted at v . The sweep-rule defines the correction to be over the positive qubits such that it will turn off the positive syndrome. It's not hard to see that for the $(2, 2)$ -Toric complex, the Toom's and the sweep rules are the same. Compared to other correction cycles on topological codes, the main advantage of the sweep is that it is independent of time. The invention of the sweep-rule is attributed to Aleksander Kubica and John Preskill [KP19]. Here we extend their work to cubical complexes.

Definition 4.6 (Sweep Rule). *Let σ be a $d' - 1$ -dimensional chain. The sweep rule decoding procedure at a vertex $v \in \mathcal{G}_n^d$ is defined as follows: Measure the checks rooted at v and look for a rooted chain $\tilde{z}$ at v such that:*

1. *The vertices in $\tilde{z}$ are contained in the vertices of σ , namely: $\mathbf{Str}_{0+}(v, \tilde{z}) \subset \mathbf{Str}_{0+}(v, \sigma)$*
2. *The boundary of $\tilde{z}$ agrees with σ on the checks rooted at v , namely: $(\partial^- \tilde{z})|_{v+} = \sigma|_{v+}$.*

If such an assignment exists, flip $\tilde{z}$.

The first property is analogous to the second property of 'causal structure' defined in [KP19]⁹. There, they conjectured that vertices that see a syndrome only on their positive checks can suggest a correction that is 'supported' on the vertices of their unsatisfied checks. For lower dimensions, they mentioned that the property can be easily checked. Here, we provide a proof for any dimensions where the face dimension is greater than one. We use the language of cubical complexes; however, since the 'positive' local view at a vertex is identical for both cubical and simplicial complexes, generalizing the proof for simplicial complexes is a matter of adjusting notations. By that, we affirmatively answer their conjecture. Moreover, the proof will provide us with an algorithm to find $\tilde{z}$, which is significantly better than running all options.

Claim 4.7. *Consider a nontrivial $d' - 1$ -dimensional chain σ . If there is a d' -dimensional chain z that is rooted at v and whose boundary agrees with $\sigma|_{v+}$, then there is a d' -dimensional chain $\tilde{z}$ rooted at v that satisfies the sweep-rule conditions.*

Proof. We prove it by induction on d' . Without loss of generality, we can assume that σ has support only over the positive faces of v , namely $\sigma|_{v+} = \sigma$.

1. Base. Consider $d' = 2$ and let d be any integer equal to or greater than 3. We argue that the support size of σ has to be even. To see that, denote the faces in z by $\theta_{\mathbf{v}, \mathbf{S}'_i}$, where $\mathbf{S}'_i = \{g_{i,1}, g_{i,2}\}$. Using the fact that σ is the boundary of z at v , we have:

$$\sigma = (\partial^-(z))|_v = \sum_{\theta_{\mathbf{v}, \mathbf{S}'_i} \in z} \sum_{g \in \mathbf{S}'_i} \theta_{\mathbf{v}, \mathbf{S}'_i/g} = \sum_i \theta_{\mathbf{v}, \{g_{i,1}\}} + \theta_{\mathbf{v}, \{g_{i,2}\}}$$

So the number of terms in the sum is even, and since a term is zeroed out only if it appears an even number of times, we have that after excluding the vanishing terms, σ has an even number of faces. Since the number of faces is even, there is a perfect matching between them. It's easy to see that the 2-chain $\tilde{z}$, whose each face corresponds to the union of the two edges in a match, satisfies the sweep-rule conditions.

⁹In their language: $\diamond(\tilde{z}) = \diamond(\sigma|_v)$.

2. Assume the correctness of the claim for every face dimension lower than d' . Now for a face dimension d' and an arbitrary generator $g \in S$ such $gv \in \sigma$, divide σ into a sum of two chains, the first contain the faces of σ that supported on the vertex gv where the second contains the compliment. Denote them by $\sigma|_{gv}$ and $\sigma|_{/gv}$, so $\sigma = \sigma|_{gv} + \sigma|_{/gv}$.

Now let σ^* the chain obtained from $\sigma|_{gv}$ by removing from any face of it the vertex gv , thus σ^* is a $d' - 2$ -dimensional chain. Moreover, let z^* be the $d' - 1$ -dimensional chain complex obtained from $z|_{gv}$ by excluding gv from its faces. Observe that

$$(\partial^- z^*)|_v = \sum_{\theta_{\mathbf{v}, \mathbf{S}' \in z|_{gv}}} \sum_{g' \in S'/g} \theta_{\mathbf{v}, \mathbf{S}'/\{\mathbf{g}, \mathbf{g}'\}} = \sum_{\theta_{\mathbf{v}, \mathbf{S}' \in (\partial^- z)|_{gv}}} \theta_{\mathbf{v}, \mathbf{S}'/\{\mathbf{g}\}} = \sigma^*$$

Thus, by the induction assumption, there is a $d' - 1$ -dimensional chain $\tilde{z}^*$ such that:

$$(a) \text{Str}_{0+}(v, \tilde{z}^*) \subset \text{Str}_{0+}(v, \sigma^*)$$

$$(b) (\partial^- \tilde{z}^*)|_{v+} = \sigma^*|_{v+}$$

The first condition implies that $gv \notin \tilde{z}^*$. Hence, we can define the d' -dimensional chain $\tilde{z}$ by adding the vertex gv to each face of $\tilde{z}^*$. From the second it follows that:

$$\begin{aligned} (\partial^- \tilde{z})|_v &= (\partial^- (\sum_{\theta_{\mathbf{v}, \mathbf{S}' \in \tilde{z}^*}} \theta_{\mathbf{v}, \mathbf{S}' \cup \{\mathbf{g}\}}))|_v = \sum_{\theta_{\mathbf{v}, \mathbf{S}' \in \tilde{z}^*}} (\sum_{g' \in S'/g} \theta_{\mathbf{v}, \mathbf{S}'/\{\mathbf{g}'\} \cup \{\mathbf{g}\}}) + \theta_{\mathbf{v}, \mathbf{S}'} \\ &= \sum_{\theta_{\mathbf{v}, \mathbf{S}' \in (\partial^- \tilde{z}^*)|_v}} \theta_{\mathbf{v}, \mathbf{S}' \cup \{\mathbf{g}\}} + \tilde{z}^* = \sigma|_{gv} + \tilde{z}^* \end{aligned}$$

Therefore,

$$\begin{aligned} \sigma + (\partial^- \tilde{z})|_v &= \sigma|_{gv} + \sigma|_{/gv} + \sigma|_{gv} + \tilde{z}^* \\ &= \sigma|_{/gv} + \tilde{z}^* \end{aligned}$$

Thus we reduce the question to handle the $d' - 1$ -dimensional chain $\sigma + (\partial^- \tilde{z})|_v$. However, since $gv \notin \sigma + (\partial^- \tilde{z})|_v$ it follows that:

$$|\text{Str}_{0+}(v, \sigma + (\partial^- \tilde{z})|_v)| \leq |\text{Str}_{0+}(v, \sigma)| - 1$$

Thus, we can repeat recursively until the resulting $d' - 1$ chain is 'supported' on fewer than d' vertices. In that case, the chain must be trivial. □

The algorithm given in the proof of Claim 4.7 for finding $\tilde{z}$ is as follows: given σ , pick a $g \in S$ such that $gv \in \sigma$. Then, define σ^* by removing gv from each face of $\sigma|_{gv}$. Let $\tilde{z}^*$ be the result of the recursive call on σ^* , and return ' $\tilde{z}^* \cup \{gv\}$ ' plus the recursive call on $\sigma|_{/gv} + \tilde{z}^*$.

We stop the recursion whenever either $d' = 2$ or $|\mathbf{Str}_{0+}(v, \sigma)| \leq d'$. The work in each recursion step consists of decomposing σ into $\sigma|_{gv}$ and $\sigma|_{/gv}$ and summing $\binom{d}{d'}$ -length vectors. We summarize by introducing the following claim which gives the running time of the algorithm in the RAM computation model.

Claim 4.8. *Let $T(d, d', k) : \cdot \rightarrow \mathbb{R}$ be the upper bound function for the running time of the algorithm given in Claim 4.7, where d is the dimension of the skeleton, d' is the face dimension, and k is the number of vertices supporting the local view, namely $k = |\mathbf{Str}_{0+}(v, \sigma)|$. Then:*

$$T(k, d') = O(2^k \binom{d}{d'})$$

Proof. In each recursive call, the vertex support of σ matches a different subset of the vertices in $\mathbf{Str}_{0+}(v, \sigma)$; thus, the number of recursive calls is at most the number of subsets, namely 2^k . The decomposition of σ into $\sigma|_{gv}$ and $\sigma|_{/gv}$, and the summation $\sigma|_{/gv} + \tilde{z}^*$ are linear in the vector's length, namely linear in $\binom{d}{d'}$. So we bound the total work by the number of recursive steps multiplied by the maximal work at each recursive step. That becomes at most:

$$T(d, d', k) \leq O(2^k \binom{d}{d'})$$

□

Later (see Definition 4.34 and Definition 4.35), we will see that an isomorphism between the (d, d') -Toric code and its dual, which is the $(d, d - d')$ -Toric code, can be realized by transversal Hadamards followed by the following permutation over the qubits.

$$\phi(\theta_{\mathbf{v}, \mathbf{S}'}) \rightarrow \theta_{\mathbf{S}'\mathbf{v}, \mathbf{S}/\mathbf{S}'}$$

A full decoding cycle consists of four steps: first, applying the sweep rule regarding the X -syndrome; next, moving to the dual code; then applying the sweep rule again; and finally getting back to the code.

Definition 4.9 (Sweep cycle). *We define the sweep cycle to be the following procedure:*

1. *For each vertex, perform the sweep rule.*
2. *Apply the Hadamard transform $\mathbf{H}^{\otimes |\mathcal{F}_{d'}|}$ followed by the ϕ -transpose.*
3. *Again, apply the sweep rule at each vertex.*
4. *Reverse stage 2 by reordering the qubits according to the ϕ^{-1} -transpose, and then apply the Hadamard transform $\mathbf{H}^{\otimes |\mathcal{F}_{d'}|}$.*

4.2 Infinite Sides Torics

The Toric skeleton provides a notation for positive directions. However, to demonstrate that an 'error cluster' shrinks as a result of the decoding cycle, or at least does not expand too much during the operation of the logical operator, we will need a way to mark extremities (for example, the northernmost point). To achieve this, we will analyze the view over the finite Toric as a projection derived from a view over the infinite plane. For this purpose, given our initial (subjected) quantum circuit, we construct a representation in which any gate that originally acted on a quantum register, namely on the finite Toric, is mapped to an infinite set of quantum gates that act on the infinite number of qubits set on the infinite plane. This set respects a periodic condition such that any restriction of it to a finite rectangle at the size of the original Toric yields the original view, and therefore describes well the original running.

The following two definitions simplify decoding gates by replacing them with the Pauli 'correction' they suggest at the end of their cycle. The first defines what a decoding gate is, and the second defines the replacement.

Definition 4.10 (Decoding Gadget). *Let C be a quantum circuit, and let u be a gate in C with fan-in at most Δ . We say that u is a decoding gadget if it has the following form:*

1. *The wires of u can be decomposed into two subsets A and B such that the wires in A are reset wires, and there exists an encoded register R in C such that the wires in B are all set in R .*
2. *For a set of faults $\mathcal{E}$, there is a Pauli gate u' that acts only on B such that the action of u in $C[\mathcal{E}]$ equals u' . Put differently, the quantum circuit obtained from $C[\mathcal{E}]$ by replacing u with u' equals $C[\mathcal{E}]$.*

Definition 4.11 (Unitary Decoding Representation). *Let $C = \{u_i\}_{i=1}$ be a quantum circuit, and let $\mathcal{U} = \{u'_i\}_{i=1} \subset C$ be a subset of the gates in C that are decoding gadgets. For a set of faults $\mathcal{E}$, we define the unitary decoding representation of $C[\mathcal{E}]$ to be the quantum circuit obtained by replacing each of the gates in $\mathcal{U}$ with their corresponding Pauli gate according to $\mathcal{E}$.*

To have our constructions yield circuits that have ∞ -extensions Definition 4.17, the sweep-rule cycles have to be 'decoding gadgets' Definition 4.10.

Claim 4.12. *The local gates that perform the sweep rule are decoding gadgets.*

Proof. A local sweep update acts on a constant-size neighborhood around a vertex, uses only local syndrome information, and may be implemented with fresh ancillas that are reset after the update. Any dependence on faults inside the gadget can therefore be pushed onto the data wires as an effective Pauli operator supported on the encoded register, while the reset ancillas carry no residual information. This is exactly the form required in the definition of a decoding gadget. $\square$

Until the end of the subsection, we will assume that the subjected circuits are given in their unitary decoding representation.

The following definition extends the notion of location to differentiate between a coordinate that encodes a register and a coordinate of the location inside the register. This will be useful, as in the infinite plane representation, each register will be mapped to a unique infinite register. From a technical perspective, the map will stretch only inner locations to infinity.

Definition 4.13 (Toric register coordinate system). *Let $C = \{u_i\}_i$ be a quantum circuit of width w with registers $\mathcal{R} = R_1, \dots, R_k$, each encoded in the (d, d') -Toric code, induced by a d -dimensional Toric with a side length n . We define the 'Toric register coordinate system' to be the coordinate system that encodes spatial location with two coordinates as follows: In the first coordinate, the register number is recorded, and in the second, the qubit is encoded by the face associated with it. We call the first the register coordinate and the second the inner coordinate.*

We will need two types of gate translation. The first is the standard spatial translation, which means that any of the inner coordinates of the gate are advanced by the same vector. The second is 'rooted translation.' Here, we assume that the locations of the gate are given in a relative-to-root fashion, meaning that there is a root vertex, and all the locations are given by positive paths from the root. Then, we first translate the root and then advance the inner coordinate of any of the other locations to be on the path from the translated root. When considering a product gate, the translations act the same.

Definition 4.14 (Translation by vertex v). *Let u be a gate, and denote the inner coordinates of u 's locations by $\theta_{v_1, J_1}, \dots, \theta_{v_\Delta, J_\Delta}$. The translation of u by vertex v is defined to be the gate obtained from u by replacing the inner coordinates of its locations by: $\theta_{vv_1, J_1}, \dots, \theta_{vv_\Delta, J_\Delta}$, where the multiplication $v \cdot v_i$ is with respect to the infinity cover $\mathcal{G}^d$. Furthermore, for a d -dimensional Toric complex with side length n , we define a translation by an integer vector $a \in \mathbb{Z}^d$ as a translation by the vertex:*

$$(g_1^n)^{a_1}, (g_2^n)^{a_2}, \dots, (g_d^n)^{a_d}$$

*We denote the translation by an integer vector a as $a * u$.*

Definition 4.15 (Rooted gate at vertex v). *Let u be a quantum gate that acts on a register whose qubits are associated with a Toric complex at dimension d . We say that a gate u is rooted at vertex v if the inner coordinate of its location can be encoded by lists of pairs $(I_1, J_1), (I_2, J_2), \dots, (I_l, J_l)$ such that $I_i \in S^d, J_i \subset S$, and the i th qubit on which it acts is in the face $\theta_{I_i v, J_i}$.*

Definition 4.16 (Root translation in Register System). *For a gate u that is a rooted gate at vertex v , we define its translation by a vertex v' to be the gate obtained by copying u and replacing its inner location encoding the root by $v'v$.*

Similarly to the standard translation by an integer vector, we extend the action to a root translation by an integer vector $a \in \mathbb{Z}^d$, as the root translation by the vertex:

$$(g_1^n)^{a_1}, (g_2^n)^{a_2}, \dots, (g_d^n)^{a_d}$$

We denote the root translation by an integer vector a as $a *_r u$.

Now we are ready to give an explicit description of the mapping. Given a quantum circuit C , we name its presentation over the infinite plane the ∞ -extension of C .

Definition 4.17. Let n be an integer, and let $C = \{u_i\}_{i=1}$ be a quantum circuit with width w and quantum registers $\mathcal{R} = R_1, \dots, R_k$ such that each register R_i encodes a (d, d') -Toric code with side at length n . We define the ∞ -extension of C , denoted by C^∞ , as follows:

1. Each register R_i in $\mathcal{R}$ is mapped to a register R'_i , which its qubits are associated with the d' -dimensional faces of the infinite d -dimensional Toric.
2. Each gate $u_i \in C$ is mapped to the infinite collection of gates, where the gates are obtained by either the standard translation or the rooted translation by an integer vector, depending on whether u_i is a decoding gate as follows:

$$u_i \mapsto \begin{cases} \bigcup_{a \in \mathbb{Z}^d} \{a *_r u_i\} & \text{if } u_i \text{ is not a decoding gadget} \\ \bigcup_{a \in \mathbb{Z}^d} \{a_r *_r u_i\} & \text{else} \end{cases}$$

The restriction of the qubits in C^∞ 's registers to finite Torices with side length n will be denoted by $\cdot|_{[w]}$, and it includes all the faces whose root vertex has 'powers' (the powers of $g_1, \dots, g_d$) in the range $[0, n-1]$. For a set of faults $\mathcal{E} = \{f_i\}_{i=1}$, we denote by $C^\infty[\mathcal{E}^\infty]$ the ∞ -extension of $C[\mathcal{E}]$.

The infinite presentation of a circuit agrees with it in terms of action. Concretely, suppose that the input to the circuit C is the zero state; then the actions of C and its ∞ -extension are 'equivalent' in the sense that the restriction of C^∞ 's output to $[w]$ yields the same mixed state as the output of C . The following provides a formal definition of that 'equivalence', which we name ∞ -compatible.

Definition 4.18. For a quantum circuit C , of the form defined in Definition 4.17, and a set of faults $\mathcal{E}$, denote:

$$\begin{aligned} \rho &:= C[\mathcal{E}] |0\rangle\langle 0| C[\mathcal{E}]^\dagger \\ \rho_\infty &:= C^\infty[\mathcal{E}^\infty] |0\rangle\langle 0| (C^\infty[\mathcal{E}^\infty])^\dagger \end{aligned}$$

We say that a quantum circuit C is ∞ -compatible if, for any set of faults $\mathcal{E} = \{f_i\}_{i=1}$, it holds that:

$$\rho = \mathbf{Tr}_{/[w]}(\rho_\infty)$$

In the appendix, we provide a proof that any circuit as in Definition 4.17 is ∞ -compatible (Claim 7.4).

The computation DAG lifts naturally to the ∞ -extension: gates that arise as translates of the same local gate act on disjoint translated supports inside different periods, so they may be executed simultaneously. The same periodic lifting applies to

the underlying geometric objects, passing from the finite Toric complex to its infinite cover.

We close that subsection by introducing the *potential walls*. Analyzing the infinite presentation enables us to bound a subset of errors by 'triangles'. A potential wall is a linear function $\mathbb{R}^d \rightarrow \mathbb{R}$. For the subset of vertices (that might be the subset that supports some error), the maximal value of the potential wall on it is achieved at the point of tangency between the subsets and its bounding triangle. With the assistance of potential walls, we can discuss the endpoint of a bunch of qubits and the 'size' of the (hyper)triangle that bounds them. Soon, the qubits will be associated with the support of an error, and showing that an action decreases the maximums of the potential walls implies that the error shrinks. The next definition explains how the embedding of vertices of the infinite in $\mathbb{R}^d$ is handled.

Definition 4.19 (The standard embedding). *Consider the d -dimensional infinite grid $\mathcal{G}^d$. The standard embedding $\mathcal{G}^d \rightarrow \mathbb{R}^d$ is given by fixing a vertex in $\mathcal{G}^d$ to be sent to the origin and then mapping the generators of the grid to the standard basis. We extend notation as follows: for a vertex $v \in \mathcal{G}^d$, a generator $g \in S$, and a coordinate $i \in [d]$, such that under the embedding $g \rightarrow i$, we denote by $v_i = v_g$ the i 'th coordinate of the image of v .*

Now we are ready to define the *potential walls*:

Definition 4.20 (Potential Walls). *We define the potential walls to be the $d + 1$ functions $L_1, L_2, \dots, L_{d+1}$, from the vertices of $\mathcal{G}^d$ to the reals as follow:*

1. *For any coordinate $i \in [d]$, we define $L_i: \mathcal{G}^d \rightarrow \mathbb{R}$ as $L_i(v) := v_i$. In addition, for the generator $g \in S$ which is mapped to i , i.e $g \rightarrow i$, we identify $L_g = L_i$.*
2. *We define $L_{d+1}: \mathcal{G}^d \rightarrow \mathbb{R}$ by $L_{d+1}(v) := -\sum_{i=1}^d v_i$.*

Consider an assignment z over the faces. For a vertex $v \in z$, and a potential wall $l \in \{L_1, \dots, L_{d+1}\}$, we say that v is a local maximum of l in z if for any neighbor u of v , such that $u \in z$, we have $l(v) \geq l(u)$. When the assignment z is clear from the context, we might omit it and briefly say that v is a local maximum of l . In addition we define the function $\mathbf{Size}(\cdot): \mathcal{G}^d \rightarrow \mathbb{R}$, to act on a subset of points $A \subset \mathcal{G}^d$, as follows:

$$\mathbf{Size}(A) := \sum_{g \in S} \max_{v \in A} L_g(v) + \max_{v \in A} L_{d+1}(v)$$

and extend the function to act over assignment on faces z , by taking its action on the set of vertices that support z .

The following claim shows that size has a similar behavior to a metric:

Claim 4.21. *Let $A, B \subset \mathbb{R}^d$ such that there $A \cap B \neq \emptyset$ then:*

$$\mathbf{Size}(A \cup B) \leq \mathbf{Size}(A) + \mathbf{Size}(B)$$

Proof. Let p be a point in $A \cap B$. Then:

$$\begin{aligned}
\mathbf{Size}(A \cup B) &= \sum_{i \in [d+1]} \max_{v \in A \cup B} L_i(v) \\
&= \sum_{i \in [d+1]} \max\left\{\max_{v \in A} L_i(v), \max_{v \in B} L_i(v)\right\} \\
&= \sum_{i \in [d+1]} \max_{v \in A} L_i(v) + \max_{v \in B} L_i(v) - \min\left\{\max_{v \in A} L_i(v), \max_{v \in B} L_i(v)\right\} \\
&\leq \sum_{i \in [d+1]} \max_{v \in A} L_i(v) + \max_{v \in B} L_i(v) - L_i(p) \\
&\leq \mathbf{Size}(A) + \mathbf{Size}(B) - \sum_{i \in [d+1]} L_i(p) = \mathbf{Size}(A) + \mathbf{Size}(B)
\end{aligned}$$

Where in the last pass, we used that the size of a single point is zero. $\square$

4.3 Reduced Error

Errors can be thought of as assignments on the d' -dimensional faces that have non-vanishing boundaries. In particular, X -errors can be characterized as d' -dimensional chains $z \in \mathbb{F}_2^{\mathcal{F}^{d'}}$ such that $\partial^- z \neq 0$. In this section, we assume dealing with the ∞ -extension of a circuit (see Definition 4.17), so the cubical complex is the d -dimensional grid with infinite side length. The *reduced* form of an error is its representative, a translation by a stabilizer, which minimizes its size. Here we prove that the sizes of a reduced error and its boundary (derived from ∂^-) are the same. In particular, they share poles of the potential walls. That will come into use in analyzing the behavior of the sweep-rule, as it is defined to decide on correction upon the syndrome it sees.

For doing so we present two codes that describe local views of vertices. They are the *stabilizers code at vertex v* and the *small code at vertex v* . The first corresponds to stabilizers rooted at the vertex, while the second is the collection of local views on faces rooted at v that satisfy the checks rooted at v . Then, in Claim 4.24, we consider a vertex whose X -checks don't admit a syndrome, meaning that its local view is in its *small code*, and show that there is a codeword of its *stabilizers code* that agrees with its positive faces. Addition by that stabilizer yields a new chain excluding v . If v were a unique maximum of L_{d+1} , then the vertices on the yielded chain would have a lower value of L_{d+1} . Subsequently, since the exclusion decreases the size, it follows that the 'corner'—the vertex achieving the maximal L_{d+1} value—of a reduced error sees a syndrome.

Definition 4.22 (Stabilizers Code at Vertex v). *We define stabilizer code at vertex v , denoted by Stab_v , as the restriction of the (d, d') -Toric code to the d' -dimensional faces such that their vertices are at a distance of at most one positive step in each direction from v . That is equivalent to the union of all the d' -dimensional faces rooted at v with the d -dimensional faces of each of its sibling gv , excluding faces containing g*

as their generator. Namely:

$$\begin{aligned} \text{Stab}_v := \{ & z : \partial^- z = 0 \text{ and} \\ & z \in \text{Span}(\{\theta_{\mathbf{v}, \mathbf{S}'} : S' \subset S, |S'| = d'\} \\ & \bigcup_{g \in S} \{\theta_{\mathbf{g}\mathbf{v}, \mathbf{S}'/\mathbf{g}} : S' \subset S/g, |S'| = d'\})\} \end{aligned}$$

The checks of the code, are of the following three types:

1. Checks that are rooted at v , each identified by choosing $d - 1$ generators from S . Namely:

$$\{\theta_{\mathbf{v}, \mathbf{S}'} : S' \subset S, |S'| = d' - 1\}$$

For each check, the bits that connect to it, are given by restricting the coboundary of the check to bits domain of Stab_v :

$$(\partial^+ \theta_{\mathbf{v}, \mathbf{S}'})|_{\text{Stab}_v} = \left(\sum_{g \in S/S'} \theta_{\mathbf{v}, \mathbf{S}' \cup \mathbf{g}} + \sum_{g \in S'} \theta_{\mathbf{g}^{-1}\mathbf{v}, \mathbf{S}' \cup \mathbf{g}} \right)|_{\text{Stab}_v} = \sum_{g \in S/S'} \theta_{\mathbf{v}, \mathbf{S}' \cup \mathbf{g}}$$

2. Checks are rooted at the positive sibling of v . Consider such a sibling, let's say gv (namely the vertex that is given by stepping in the direction $g \in S$ from v). Notice that any face rooted at gv that contains ggv cannot be contained in faces rooted at v . Thus, the checks of gv correspond to choosing subsets from S/g . Namely:

$$\bigcup_{g \in S} \{\theta_{\mathbf{g}\mathbf{v}, \mathbf{S}'} : S' \subset S/g, |S'| = d' - 1\}$$

And each check is supported by:

$$\begin{aligned} (\partial^+ \theta_{\mathbf{g}\mathbf{v}, \mathbf{S}'})|_{\text{Stab}_v} &= \left(\sum_{g' \in S/S'} \theta_{\mathbf{g}\mathbf{v}, \mathbf{S}' \cup \mathbf{g}'} + \sum_{g' \in S'} \theta_{\mathbf{g}'^{-1}\mathbf{g}\mathbf{v}, \mathbf{S}' \cup \mathbf{g}'} \right)|_{\text{Stab}_v} \\ &= \left(\sum_{\substack{g' \in S/S' \\ g' \neq g}} \theta_{\mathbf{g}\mathbf{v}, \mathbf{S}' \cup \mathbf{g}'} \right) + \theta_{\mathbf{v}, \mathbf{S}' \cup \mathbf{g}} \end{aligned}$$

3. Checks that their roots are at a two-step distance from v , namely vertices of the form $g'gv$ where $g, g' \in S$, and $g \neq g'$. By the arguments presented in the second type, the checks are the $d' - 1$ faces which do not have the generators on the path leading from v to them, i.e., g, g' , in their generator set.

$$\bigcup_{g, g' \in S'} \{\theta_{\mathbf{g}'\mathbf{g}\mathbf{v}, \mathbf{S}'} : S' \subset S/\{g, g'\}, |S'| = d' - 1\}$$

And for each check, its support is:

$$\begin{aligned} (\partial^+ \theta_{\mathbf{g}'\mathbf{g}\mathbf{v}, \mathbf{S}'})|_{\text{Stab}_v} &= \left(\sum_{g'' \in S/S'} \theta_{\mathbf{g}\mathbf{g}'\mathbf{v}, \mathbf{S}' \cup \mathbf{g}''} + \sum_{g \in S'} \theta_{\mathbf{g}''^{-1}\mathbf{g}\mathbf{g}'\mathbf{v}, \mathbf{S}' \cup \mathbf{g}''} \right)|_{\text{Stab}_v} \\ &= \theta_{\mathbf{g}\mathbf{v}, \mathbf{S}' \cup \mathbf{g}'} + \theta_{\mathbf{g}'\mathbf{v}, \mathbf{S}' \cup \mathbf{g}} \end{aligned}$$

Definition 4.23 (Small Code at Vertex v). For a vertex v , we denote by $\mathcal{C}_v$ the small code at vertex v , which consists of all the d' -dimensional chains rooted at v that satisfy the checks rooted at v , namely:

$$\mathcal{C}_v := \left\{ z : (\partial^- z)|_{v+} = 0 \text{ and } z \in \text{Span}(\{\theta_{\mathbf{v}, \mathbf{S}'} : \mathbf{S}' \subset S, |\mathbf{S}'| = d'\}) \right\}$$

The checks of the code are exactly the first type checks of Stab_v .

Claim 4.24. Let v be a vertex. For any $d' \geq 2$, $d \geq d' + 1$, and any $z \in \mathcal{C}_v$, there is a stabilizer $w \in \text{Stab}_v$ such that $z = w|_v$; namely, z and w agree on the faces rooted at v . Moreover, the vertex support of $w|_v$ is contained in the vertex support of $z|_v$, namely:

$$\text{Str}_{0+}(v, w) \subset \text{Str}_{0+}(v, z)$$

Proof. By induction on the dimensions d, d' . We extend the notation of the stabilizer and the small codes at vertex v to $\mathcal{C}_v^J, \text{Stab}_v^J$, where J is a subset of S , to refer to the codes derived from the skeleton generated by J . Showing that the belonging $z \in \mathcal{C}_v^J$ implies the existence of $w \in \text{Stab}_v^J$ such that $w|_v = z|_v$ proves the claim.

1. Base. For the base cases, let d be any number greater than or equal to three, and let d' be two. First, observe that, in the case $d' = 2$, the code $\mathcal{C}_v^J$ is a Tanner code, where the bits are set on edges of the form $\{vg_i, vg_j\}$ and checks are associated with vertices of the form $\{vg_i\}$. Therefore, any codeword is a cycle. Let:

$$\{vg_1, vg_2\}, \{vg_2, vg_3\}, \dots, \{vg_k, vg_1\}$$

Be the cycle corresponding to z . Namely:

$$z = \theta_{\mathbf{v}, \{\mathbf{g}_1, \mathbf{g}_2\}} + \theta_{\mathbf{v}, \{\mathbf{g}_2, \mathbf{g}_3\}} + \dots + \theta_{\mathbf{v}, \{\mathbf{v}\mathbf{g}_k, \mathbf{v}\mathbf{g}_1\}}$$

And consider the stabilizer w defined as:

$$w = \partial^- (\theta_{\mathbf{v}, \{\mathbf{g}_1, \mathbf{g}_2, \mathbf{g}_3\}} + \theta_{\mathbf{v}, \{\mathbf{g}_1, \mathbf{g}_3, \mathbf{g}_4\}} + \dots + \theta_{\mathbf{v}, \{\mathbf{g}_1, \mathbf{g}_i, \mathbf{g}_{i+1}\}} + \dots + \theta_{\mathbf{v}, \{\mathbf{g}_1, \mathbf{g}_{k-1}, \mathbf{g}_k\}})$$

Clearly, $w \in \text{Stab}_v^J$, and in addition its restriction to v equals:

$$\begin{aligned} w|_v &= \left(\sum_{i=2}^{k-1} \partial^- \theta_{\mathbf{v}, \{\mathbf{g}_1, \mathbf{g}_i, \mathbf{g}_{i+1}\}} \right) |_v = \sum_{i=2}^{k-1} \theta_{\mathbf{v}, \{\mathbf{g}_i, \mathbf{g}_{i+1}\}} + \theta_{\mathbf{v}, \{\mathbf{g}_1, \mathbf{g}_i\}} + \theta_{\mathbf{v}, \{\mathbf{g}_1, \mathbf{g}_{i+1}\}} \\ &= \left(\sum_{i=3}^{k-1} \theta_{\mathbf{v}, \{\mathbf{g}_i, \mathbf{g}_{i+1}\}} \right) + \theta_{\mathbf{v}, \{\mathbf{g}_1, \mathbf{g}_2\}} + \theta_{\mathbf{v}, \{\mathbf{g}_k, \mathbf{g}_1\}} = z \end{aligned}$$

2. Assume the correction of the claim for dimensions lower than d, d' . Consider a generators subset $J \subset S$ and let $z \in \mathcal{C}_v^J$. In addition, consider a generator $g \in J$.

Decompose $z = z|_{gv} + z|_{/gv}$ where $z|_{gv}$ and $z|_{/gv}$ are the chains consisting of all the faces in z that have support on gv and its complement. Define z^* to be the $d' - 1$ -dimensional chain obtained by excluding gv from any face of $z|_{gv}$. Notice that $(\partial^- z^*)|_v$ has to be zero since it contains all the faces of $\partial^- z$ that contain gv (up to the removal). In other words, $z^* \in \mathcal{C}_v^{J/g}$. Thus, we can use the induction assumption and infer the existence of $w^* \in \text{Stab}_v^{J/g}$ such that $z^* = w^*|_v$.

Now, since $w^* \in \text{Stab}_v^{J/g}$, then $\partial^- w^* = 0$. And since the local complex is contractible, w^* can be decomposed into a sum of stabilizers rooted in v . Let $\theta_{\mathbf{v}, \mathbf{S}'_1}, \theta_{\mathbf{v}, \mathbf{S}'_2}, \theta_{\mathbf{v}, \mathbf{S}'_3}, \dots, \theta_{\mathbf{v}, \mathbf{S}'_i}$ be the d' -dimensional chains corresponding to the stabilizers composing w^* . Namely:

$$w^* = \sum_i \partial^- \theta_{\mathbf{v}, \mathbf{S}'_i}$$

Since those stabilizers are basis elements of $\text{Stab}_v^{J/g}$, we have that $g \notin S'_i$ for any i . Hence, we can define the extended d' -dimensional chain w to be the chain obtained by adding g to each of the stabilizers deriving w^* . Namely:

$$w = \sum_i \partial^- \theta_{\mathbf{v}, \mathbf{S}'_i \cup \mathbf{g}}$$

However, $\theta_{\mathbf{v}, \mathbf{S}'_1 \cup \mathbf{g}}, \dots, \theta_{\mathbf{v}, \mathbf{S}'_i \cup \mathbf{g}}$ are $d' + 1$ -dimensional chains corresponding to the stabilizers, which are basis elements of Stab_v^J . Therefore, $w \in \text{Stab}_v^J$. In addition, opening up the boundary operator yields:

$$\begin{aligned} w|_v &= \sum_i (\partial^- \theta_{\mathbf{v}, \mathbf{S}'_i \cup \mathbf{g}})|_v = \sum_i \left(\sum_{g' \in S'_i/g} \theta_{\mathbf{v}, \mathbf{S}'_i/g' \cup \mathbf{g}} + \theta_{\mathbf{v}, \mathbf{S}'_i} \right) \\ &= \sum_{\theta_{\mathbf{v}, \mathbf{S}'_i} \in w^*|_v} \theta_{\mathbf{v}, \mathbf{S}'_i \cup \mathbf{g}} + \sum_i \theta_{\mathbf{v}, \mathbf{S}'_i} = z|_{gv} + \sum_i \theta_{\mathbf{v}, \mathbf{S}'_i} \end{aligned}$$

Thus, adding $w|_v$ to z results in:

$$z + w|_v = z|_{/gv} + \sum_i \theta_{\mathbf{v}, \mathbf{S}'_i}$$

Which is a d' -dimensional chain, supported on vertices of the form $g'v$ for $g' \in J/g$, namely a d' -dimensional chain on the complex induced by the $d - 1$ -dimensional skeleton. In addition, since the boundary operator can only advance the root of a face forward, we have that $\partial^- w|_v = (\partial^- w)|_v$ (as w is rooted in v) and therefore $\partial^-(z + w|_v) = \partial^- z + (\partial^- w)|_v = 0$.

To summarize, we have both $z + w|_v \in \mathcal{C}_v^{J/g}$ and $z + w|_v$ sets on a $d - 1$ -dimensional skeleton. Therefore, we can repeat in a recursive manner, with each step resulting in a d' -dimensional chain in a skeleton with dimension decreased by one. We stop when $d = d'$; in that case, the resulting chain has to be trivial.

□

We would like to differentiate chains whose projection onto the finite Toric with side length n closes a non-trivial cycle from those which don't. For this, we fix a rough threshold via the size function. Connected chains which have size less than n will be called finite. Observe that a chain might have a projection that doesn't close a non-trivial loop on the Toric and still be marked, according to our definition, as infinite. So, the classification is rough, yet it will turn out to be sufficient.

Definition 4.25. *Let z be a d' -dimensional chain. Consider the connectivity notation defined in Definition 4.2. We will say that a connected chain z is finite if:*

$$\mathbf{Size}(z) \leq n - 1$$

Similarly, z is infinite if its size is greater than or equal to n .

Remark 4.26. *Observe that if $\mathbf{Size}(z) < n$, then the projection of z on the n -side length Toric does not close a non-trivial loop.*

The next claim states that the translation by stabilizer that minimizes its size of a finite connected chain admits a boundary with size equal to itself.

Claim 4.27. *Let $x \in \mathbb{F}_2^{\mathcal{F}_{d'}}$ be a d' -dimensional finite chain, consider the stabilizer $z \in \mathcal{C}_Z^\perp$ which adding it to x minimize the size, Namely:*

$$z = \arg \min_{z \in \mathcal{C}_Z^\perp} \mathbf{Size}(x + z)$$

Then:

$$\mathbf{Size}(x + z) = \mathbf{Size}(\partial^-(x + z))$$

Proof. First, observe that since the vertex domain of ∂^-x is contained in the vertex domain of x , it follows that for any x the inequality $\mathbf{Size}(x) \geq \mathbf{Size}(\partial^-x)$ holds. Combining this with the fact that $\partial^-z = 0$ for any $z \in \mathcal{C}_Z^\perp$, we have that $\mathbf{Size}(x + z) \geq \mathbf{Size}(\partial^-(x + z))$. Now suppose that $\mathbf{Size}(x) > \mathbf{Size}(\partial^-x)$; that implies the existence of $v \in x/\partial^-x$ such that v is a local maximum of one of the potential walls in x and $(\partial^-x)|_v = 0$. Let's split into cases, depending on the potential wall type.

1. If v is a maximum of L_{d+1} , then by Claim 4.24, there exists a stabilizer z rooted at v , such that $z|_v = z|_{v+} = x|_{v+} = x|_v$ (where we used that v is a local maximum of L_{d+1} in x to infer $x|_v = x|_{v+}$). Therefore, $(x + z)|_v = 0$, meaning adding z to x excludes v . In addition, since the rooted stabilizer z is supported only on the vertices in $\mathbf{Str}_{0+}(v, x)$, the addition cannot introduce a vertex with a higher L_g value than the maximal L_g value in $\mathbf{Str}_{0+}(v, x)$ into x . So, repeating the process until excluding all the vertices that get the maximal L_{d+1} -value in x , and summing the stabilizers, results in a stabilizer z for which $L_{d+1}(x + z) < L_{d+1}(x)$.

2. If, for some generator $g \in S$, v is a maximum of L_g , then without loss of generality, pick v among the maximum points of L_g with the highest value of $-L_{d+1}$ (maximize $\sum_{g \in S} L_g(v)$). Since stepping through the positive direction can only increase the value of $-L_{d+1}$, it follows that v is also a local max of $-L_{d+1}$. Now consider the skeleton system where the directions are flipped, and $\theta_{v,S'}$ is set to be the face containing the vertices that are reached by the selection of reverse steps from S' (i.e., $g^{-1}v \in \theta_{v,S'}$). Similarly, for the sake of the proof, let's present the following notation for the positive neighborhood (Definition 4.5) in the reversed system:

$$\mathbf{Str}_{0-}(v, x) := \{u \in x : \exists g \in S \text{ such } u = g^{-1}v\}$$

So, $\mathbf{Str}_{0-}$ plays the same role as $\mathbf{Str}_{0+}$ does in the original skeleton. In that system, we can apply Claim 4.24 (in the reversed direction) and get that there is a stabilizer z such that adding it to x excludes v from x . By the same argument as above, no vertex with a higher $L_{g'}$ -value (although L_{d+1} can increase, that's fine¹⁰) than what it already has in $\mathbf{Str}_{0-}(v, x)$ is introduced into x . Repeating the process excludes all the vertices with the maximal value of L_g in x , so eventually, we end up with $x + z$, such that $L_g(x + z) < L_g(x)$.

We keep excluding vertices until we get the desired equality. □



Figure 6: Consider the $(3, 2)$ -Toric code, with qubits and X -checks set on the 2- and 1-dimensional faces. In blue, we have a 2-dimensional chain where the vertex achieving the maximum of L_{d+1} , denoted by v , does not see any unsatisfied checks in its positive local view. The green faces represent a 2-dimensional chain, which is the **reduced** form of the blue. In particular, any local maximum of L_{d+1} (vertices u and w) sees a non-trivial syndrome in its positive local view. Namely, the maximal points of L_{d+1} in the chain are also the maximal points in the boundary ∂^- .

We close the section by defining the *reduced* form of a finite d' -dimensional chain.

¹⁰If it happens, we will apply the first case in the next exclusion.

Definition 4.28. Let $x \in \mathbb{F}_2^{\mathcal{F}_{d'}}$ be a d' -dimensional finite chain. We will say that x is reduced if for any vertex $v \in x$ that is a maximum point of a potential wall $L \in L_1, \dots, L_{d+1}$, v is also in $\partial^- x$. In particular, we have that $\mathbf{Size}(x) = \mathbf{Size}(\partial^- x)$ and that the reduced form of a given chain is obtained by the construction presented in the proof of Claim 4.27.

We emphasize that the reduced form is not unique. Throughout the paper, every time we choose a reduced form of a chain, we mean obtaining it from the given algorithm presented in Claim 4.27, so the reduced form is effectively unique.

4.4 Shrinking

We are about to prove that the sweep rule shrinks finite errors. The following claim relates the local maximum points of the potential walls $L_1, \dots, L_d$ in an assignment z to the structure of the syndrome they observe. Later, we will use that relation to infer that the decoding step of the sweep rule can't spread the error to a vertex¹¹ with a higher value than the current global maximum of $L_1, \dots, L_d$. In other words, the error does not spread in the positive directions.

Claim 4.29. Let z and ξ be d' -dimensional connected finite chains such that ξ results from the application of the sweep rule on z . For any generator $g \in S$, the maximal values of L_g in ξ and $\partial^- \xi$ are lower than the maximal values in z and $\partial^- z$. Namely:

$$\begin{aligned} \max \{L_g(u) : u \in \xi\} &\leq \max \{L_g(u) : u \in z\} \\ \max \{L_g(u) : u \in \partial^- \xi\} &\leq \max \{L_g(u) : u \in \partial^- z\} \end{aligned}$$

Proof. First, we argue that for a vertex v in $\partial^- z$, and a generator $g \in S$, such that v is a local maximum of L_g in either z or $\partial^- z$, it implies that $gv \notin \partial^- z$. The contradiction where v is a maximum of L_g in z follows immediately, since containment in $\partial^- z$ implies containment in z . Where v is a maximum in $\partial^- z$, suppose through contradiction that $gv \in \partial^- z$, then

$$L_g(gv) = (gv)_g = 1 + v_g = 1 + L_g(v)$$

In contradiction to v being a local maximum of L_g in $\partial^- z$.

For a vertex $v \in \partial^- z$ denote the correction it applies by $\varphi_v \in \mathbb{F}_2^{\mathcal{F}_{d'}}$, so

$$\xi = \sum_{v \in \partial^- z} \varphi_v + z \quad \text{and} \quad \partial^- \xi = \sum_{v \in \partial^- z} \partial^- \varphi_v + \partial^- z$$

So for any vertex $u \in \xi/z$, there is a vertex $v \in \partial^- z$ such that $u \in \varphi_v$. Moreover, since φ_v is a rooted chain, it follows that $L_g(u) \leq L_g(v) + 1$, where the equality holds when g belongs to the generator's subset of the face φ_v . Similarly for $u \in \partial^- \xi/\partial^- z$.

Now, combining with the sweep rule's first condition, which restricts the correction applied by a vertex to have a vertex support that is contained in the positive syndrome

¹¹To faces whose vertices have higher values of $L_1, \dots, L_d$ than the current global maximum values.

it sees, we have that for any vertex v which is a local maximum of L_g (in either z or $\partial^- z$), any $u \in \varphi_v$ satisfies $L_g(u) = L_g(v)$. In other words, ξ and $\partial^- \xi$ don't contain vertices of the form gv where v is a local maximum of L_g in z and $\partial^- z$, and the desired inequalities follow. $\square$

Next, we prove that for a finite X -error, the sweep-cycle advances the corner (the maximum of L_{d+1}) of its reduced form. We assume we are given a finite error in its reduced form, and therefore whose corner sees a syndrome, and show that the sweep-rule application excludes the corner. In other words, the sweep-rule shrinks X -errors.

Claim 4.30. *Let z be a connected finite reduced d' -dimensional chain, and let ξ be the reduced result of applying the sweep rule to it. Then:*

1. *For any generator $g \in S$: $\max \{L_g(v) : v \in z\} \geq \max \{L_g(v) : v \in \xi\}$*
2. *And for L_{d+1} : $\max \{L_{d+1}(v) : v \in z\} \geq \max \{L_{d+1}(v) : v \in \xi\} + 1$*

In particular, $\mathbf{Size}(\xi) \leq \mathbf{Size}(z) - 1$.

Proof. The first inequality was proved in Claim 4.29, so it's left to prove the second. Denote by φ the applied correction and by v the vertex on which L_{d+1} gets the maximal value in z , if there are more than one vertex then choose v arbitrarily. Since z is reduced, then by Claim 4.27 it holds that $v \in \partial^- z$. Moreover, since v is a local max of L_{d+1} in z the two follows:

1. Claim 4.7 guarantees that $\partial^- \varphi|_v = \partial^- z|_v$. Thus, adding $\varphi|_v$ to z excludes v from the boundary.
2. For any other vertex $u \in z$, there is no face $\theta_{\mathbf{u}, \mathbf{S}'} \in z$ such that $v \in \theta_{\mathbf{u}, \mathbf{S}'}$. Otherwise, it would contradict v being a local maximum of L_{d+1} in z . Thus, $v \notin \mathbf{Str}_{0+}(u, \partial^- z)$, so by Claim 4.7, $v \notin \varphi|_u$.

Combining the two, we get that $v \notin \partial^-(z + \varphi)$, so the reduced form of $z + \varphi$ does not contain v . Namely, $v \notin \xi$. Repeating the above for all the other vertices which get the maximal L_{d+1} -value in z yields that $L_{d+1}(\xi) \leq L_{d+1}(z) - 1$. $\square$

Write something here.

Claim 4.31. *Let z be a connected finite d' -dimensional chain, and let ξ be the result of applying the sweep rule to it. Denote by v the vertex which achieve the maximal L_i -value in $\partial^- \xi$, then there is a vertex $u \in \partial^- z$ such $L_i(u) \geq L_i(v)$ for $i \neq d + 1$ and $L_{d+1}(u) > L_{d+1}(v)$ if $i = d + 1$, and there is a d' -dimensional face f such $v, u \in f$.*

Proof. Without loss of generality, we can assume that z and ξ are reduced. We start by proving for $i = d + 1$. Now, it splits upon whether $v \in z + \xi$:

1. If $v \in z + \xi$, then v belongs to a face that flipped during the correction, denote it's root by u and by definition of the sweep-rule $(\partial^- z)|_{u+} \neq 0$. Therefore, $u \in \partial^- z$, and u, ξ share a face which u is it's root. Since u it's the root of the face we get $L_{d+1}(u) > L_{d+1}(v)$.
2. Else, if $v \notin z + \xi$, then v has to be in z . Combining with the fact that $v \in \xi$, it follows that v is not a local maximum of L_{d+1} in $\partial^- z$, otherwise the application of the sweep-rule would exclude v (Claim 4.30), so it wouldn't be in ξ . Pick u to be the neighbor of v in $\partial^- z$ such that $L_{d+1}(u) > L_{d+1}(v)$.

It remains to prove for $i \in [d]$, in this case if $v \in \partial^- z$ then take $u = v$. Otherwise, there must be a face f that was flipped during the sweep-rule containing v . Denote by g the corresponding generator for i and by u' the root of f . If $L_g(u') = L_g(v)$ then set $u = u'$. Otherwise, g is one of the generators that span f , the second condition of the sweep-rule (Definition 4.6) requires that for u' flipping a face in the g direction, it has to see gu' in $\partial^- z$, in that case, set $u = g'u$. $\square$

We finished analyzing the first part of the sweep-cycle that acts in the computational basis. In the next subsection, we introduce the transformation into the dual code. Since the dual code is a $(d, d - d')$ -Toric code, we also find that the sweep-cycle shrinks reduced phase errors.

4.4.1 Ideal decoder

Although the main role of the section is to show that a single sweep-cycle keeps error clusters under control during a noisy run (by admitting a modest shrinking), an immediate corollary of this is that the application of non-noisy multiple iterations of the sweep-cycle has the potential to erase errors, namely a decoder. Here, a classical measurement is given, we act on classical bits, and the sweep-cycle consists only of the first stage that works over the computational basis, without moving to the dual code. We formulate the corollary as follows:

Corollary 4.32. *Consider the (d, d') -Toric code, induced by a Toric with side length n . Given a connected finite reduced error $z \in \mathbb{F}_2^{\mathcal{F}_{d'}}$ on that Toric, with $\mathbf{Size}(z) \leq n - 1$. Then an application of $\mathbf{Size}(z)$ sweep cycles (partial cycles consisting only of the first stage) erases the error.*

When z is not connected, disjoint components might merge, and it is no longer true that the size of the chain decreases in each iteration. However, we will see later (Claim 5.31, Claim 5.32, and Lemma 6.2) that performing $\Theta(n^{d+1})$ sweep-cycles erases the error with high probability.

4.5 Logical Gates

The $\theta_{\mathbf{v}, \mathbf{S}'}$ notation gives a nice interpretation of the logical code words. Particularly, an X -type logical codeword is given as follows: Take a subset of generators $S' \subset S$ and consider the subgroup generated by it, $\langle S' \rangle$; let us denote the subgroup by K .

Then the logical codeword is the X -product over all $\theta_{\mathbf{u}, \mathbf{S}'}$ such that $u \in K$. Moreover, replacing K by a coset whose base subgroup is $\langle S' \rangle$, we get another presentation of the same logical codeword.

Claim 4.33. *For any generator subset S' of size d' , denote by $x_{S'}$ and $z_{S'}$ the following d' -dimensional chains:*

$$x_{S'} := \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}'} \quad z_{S'} := \sum_{u \in \langle S/S' \rangle} \theta_{\mathbf{u}, \mathbf{S}'}$$

Then $x_{S'}$ and $z_{S'}$ are binary representations of anticommuting logical codewords of the (d, d') -Toric code. In particular, $x_{S'} \in \mathcal{C}_X / \mathcal{C}_Z^\perp$, $z_{S'} \in \mathcal{C}_Z / \mathcal{C}_X^\perp$, and $x_{S'} z_{S'} = 1$. In addition, let $X_{S'}$ and $Z_{S'}$ be the Pauli matrices obtained by multiplying each non-trivial coordinate of them by X or Z , respectively.

Proof. By definition, $x_{S'}$ and $z_{S'}$ are binary assignments on the d' -dimensional faces. We start by showing that $x_{S'} \in \mathcal{C}_X$ and $z_{S'} \in \mathcal{C}_Z$; namely, that the boundaries $\partial^- x_{S'}$ and $\partial^+ z_{S'}$ vanish.

$$\begin{aligned} \partial^- \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}'} &= \sum_{u \in \langle S' \rangle} \sum_{g \in S'} \theta_{\mathbf{u}, \mathbf{S}'/g} + \theta_{g\mathbf{u}, \mathbf{S}'/g} \\ &= \sum_{g \in S'} \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}'/g} + \sum_{u \in \langle S' \rangle} \theta_{g\mathbf{u}, \mathbf{S}'/g} \\ &= \sum_{g \in S'} \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}'/g} + \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}'/g} = 0 \\ \partial^+ \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}/\mathbf{S}'} &= \sum_{u \in \langle S' \rangle} \sum_{g \in S'} \theta_{\mathbf{u}, \mathbf{S}/\mathbf{S}' \cup g} + \theta_{g^{-1}\mathbf{u}, \mathbf{S}/\mathbf{S}' \cup g} \\ &= \sum_{g \in S'} \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}/\mathbf{S}' \cup g} + \sum_{u \in \langle S' \rangle} \theta_{g^{-1}\mathbf{u}, \mathbf{S}/\mathbf{S}' \cup g} \\ &= \sum_{g \in S'} \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}/\mathbf{S}' \cup g} + \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}/\mathbf{S}' \cup g} = 0 \end{aligned}$$

Thus $x_{S'} \in \mathcal{C}_X$ and $z_{S'} \in \mathcal{C}_Z$, now for showing that they are not binary representations of stabilizers, it's enough to show they admit no vanishing product.

$$x_{S'} z_{S'} = \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}'} \sum_{u \in \langle S/S' \rangle} \theta_{\mathbf{u}, \mathbf{S}'} = \theta_{\mathbf{e}, \mathbf{S}'} \theta_{\mathbf{e}, \mathbf{S}'} = 1$$

Therefore, the Paulis $X_{S'}$ and $Z_{S'}$ are X - and Z -type anticommuting logical code words. $\square$

The (d, d') -Toric code admits a fold-transversal realization of an isomorphism gate into the $(d, d-d')$ -Toric code, for any dimensions d, d' . Particularly, the gate exchanges logical Pauli $X_{S'}$ (of the first code) with logical Pauli $Z_{S/S'}$ (of the second code) and vice versa. It is realized by an application of transversal Hadamards followed by permuting the qubits according to a transformation which we name the ϕ -transpose map. We define the transpose map as follows:




Figure 7: In blue and green are the binary presentations x_S and z_S of logical code words of the $(2, 1)$ -Toric code. The edge $\{(5, 4), (0, 4)\}$ is omitted for clarity (it should be a green edge that hovers over the green line).

Definition 4.34 (The Transpose Map). *Let d, d' be integers such that $d' < d$, and let $\mathcal{G}$ be a (d, d') -Toric complex generated by the generator set S . We define the transpose map $\phi : \mathbb{F}_2^{\mathcal{F}^{d'}} \rightarrow \mathbb{F}_2^{\mathcal{F}^{d-d'}}$ by acting on the basis elements as follows:*

$$\phi(\theta_{\mathbf{v}, \mathbf{S}'}) \rightarrow \theta_{\mathbf{S}'\mathbf{v}, \mathbf{S}/\mathbf{S}'}$$

Notice that the action of the transpose map over the qubits is a one-to-one map between the qubits of the (d, d') -Toric code and the qubits of the $(d, d - d')$ -Toric code, and therefore can be realized at constant depth using ancillas.

Definition 4.35 (Hadamard-Type Gate $\mathbf{H}_\phi$). *We define the logical Hadamard-type gate $\mathbf{H}_\phi : L(\mathcal{H}) \rightarrow L(\mathcal{H})$, to be a code isomorphism (d, d') -Toric $\rightarrow (d, d - d')$ -Toric, as follows, for any $S' \subset S$, of size d' :*

$$\mathbf{H}_\phi(X_{S'}) \rightarrow Z_{S/S'} \quad \mathbf{H}_\phi(Z_{S'}) \rightarrow X_{S/S'}$$

Claim 4.36. *The Hadamard-type gate $\mathbf{H}_\phi$ has a fold-transversal implementation by, first applying $\mathbf{H}^{\otimes |\mathcal{F}_{d'}|}$ over all the physical qubits, and then permuting them by the transpose map ϕ , namely $\hat{\mathbf{H}}_\phi = \phi \circ \mathbf{H}^{\otimes |\mathcal{F}_{d'}|}$.*

Proof. Clearly, for any subset of generators $S' \subset S$, the map ϕ sends $x_{S'}$ to $z_{S/S'}$:

$$\phi(x_{S'}) = \phi\left(\sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}'}\right) = \sum_{u \in \langle S' \rangle} \theta_{\mathbf{S}'\mathbf{u}, \mathbf{S}/\mathbf{S}'} = \sum_{u \in \langle S' \rangle} \theta_{\mathbf{u}, \mathbf{S}/\mathbf{S}'} = z_{S/S'}$$

Moreover, ϕ sends binary representations of Z -stabilizers of the (d, d') -Toric code to binary representations of X -stabilizers of the $(d, d - d')$ -Toric code, and vice versa. To see this, consider a $d' + 1$ dimensional face, and consider the image of the stabilizer it defines.

$$\begin{aligned}\phi(\partial^- \theta_{\mathbf{v}, \mathbf{S}'}) &= \phi\left(\sum_{g \in S'} \theta_{\mathbf{v}, \mathbf{S}'/g} + \theta_{g\mathbf{v}, \mathbf{S}'/g}\right) \\ &= \sum_{g \in S'} \theta_{g^{-1}\mathbf{S}'\mathbf{v}, (\mathbf{S}/\mathbf{S}') \cup g} + \theta_{g^{-1}\mathbf{S}'g\mathbf{v}, (\mathbf{S}/\mathbf{S}') \cup g} \\ &= \partial^+ \theta_{\mathbf{S}'\mathbf{v}, \mathbf{S}/\mathbf{S}'}\end{aligned}$$

And for the checks set on the $d' - 1$ -dimensional face we have:

$$\begin{aligned}\phi(\partial^+ \theta_{\mathbf{v}, \mathbf{S}/\mathbf{S}'}) &= \phi\left(\sum_{g \in S'} \theta_{g^{-1}\mathbf{v}, (\mathbf{S}/\mathbf{S}') \cup g} + \theta_{\mathbf{v}, (\mathbf{S}/\mathbf{S}') \cup g}\right) \\ &= \sum_{g \in S'} \theta_{(\mathbf{S}/\mathbf{S}')\mathbf{v}, \mathbf{S}'/g} + \theta_{g(\mathbf{S}/\mathbf{S}')\mathbf{v}, \mathbf{S}'/g} \\ &= \partial^- \theta_{(\mathbf{S}/\mathbf{S}')\mathbf{v}, \mathbf{S}'}\end{aligned}$$

Thus, $\phi \circ \mathbf{H}^{\otimes |\mathcal{F}_{d'}|}$ sends Z -stabilizers (at dimension d') to X -stabilizers (at dimension $d - d'$) and vice versa, thereby sending the satisfied checks to satisfied checks. Combining this with the above, $\phi \circ \mathbf{H}^{\otimes |\mathcal{F}_{d'}|}$ realizes the logical isomorphism $\mathbf{H}_\phi$ from the (d, d') -Toric code into the $(d, d - d')$ -Toric code. $\square$

Example 4.37. Consider the $(2, 1)$ -Toric code. Name the generators in S by $g_1 = (1, 0)$ and $g_2 = (0, 1)$. Consider the check corresponding to the Z -stabilizer $\partial^+ \theta_{(2,3), \{\}}.$

$$\partial^+ \theta_{(2,3), \{\}} = \{\theta_{(2,3), \{(1,0)\}}, \theta_{(2,3), \{(0,1)\}}, \theta_{(1,3), \{(1,0)\}}, \theta_{(2,2), \{(0,1)\}}\}$$

The ϕ -transpose map acts on it as follows:

$$\begin{aligned}\phi(\theta_{(2,3), \{(1,0)\}}) &= \theta_{(3,3), \{(0,1)\}} \\ \phi(\theta_{(2,3), \{(0,1)\}}) &= \theta_{(2,4), \{(1,0)\}} \\ \phi(\theta_{(1,3), \{(1,0)\}}) &= \theta_{(2,3), \{(0,1)\}} \\ \phi(\theta_{(2,2), \{(0,1)\}}) &= \theta_{(2,3), \{(1,0)\}}\end{aligned}$$

What gives edges on a rectangle that converges with the bits on the support of the X -stabilizer $\theta_{(4,3), \{(0,1), (1,0)\}}.$

Equipped with $\mathbf{H}_\phi$, we gain a mapping into the dual code that both keeps the Toric skeleton and sends the Z -checks to $d' - 1$ dimensional faces. This enables the use of the same local decoder to act on both $\mathcal{C}_X$ and its dual. We summarize it in the following immediate corollary:




Figure 8: An outline of Example 4.37. The ϕ -transpose sends Z -stabilizers to X -stabilizers and vice versa. Here we consider the $(2, 1)$ Toric code; namely, the Z -stabilizers (parity checks) are set on the vertices, and the X -stabilizers (phase checks) are set on the squares. In blue and green, we have the edges corresponding to checks $\partial^+ \theta_{(2,3),\{\}}^{\mathcal{F}}$ and $\partial^- \theta_{(2,3),\{(0,1),(1,0)\}}^{\mathcal{F}}$. The example shows that indeed ϕ maps between the two.

Corollary 4.38. *Consider an Pauli P , whose X - and Z -components are set on the reduced chains $x, z \in \mathbb{F}_2^{\mathcal{F}_{d'}}$. Denote by $\tilde{P}$ the Pauli operator result by an application of the sweep-rule on P and by $\tilde{x}, \tilde{z} \in \mathbb{F}_2^{\mathcal{F}_{d'}}$ the chains corresponding to the X - and Z -components of $\tilde{P}$. If both $x, \phi(z)$ are finite connected chain, Then:*

1. For any $i \in [d]$:

$$\max_{v \in x} L_i(v) \geq \max_{v \in \tilde{x}} L_i(v) \quad \max_{v \in \phi(z)} L_i(v) \geq \max_{v \in \phi(\tilde{z})} L_i(v)$$

2. And for $i = d + 1$:

$$\max_{v \in x} L_{d+1}(v) \geq \max_{v \in \tilde{x}} L_{d+1}(v) + 1 \quad \max_{v \in \phi(z)} L_{d+1}(v) \geq \max_{v \in \phi(\tilde{z})} L_{d+1}(v) + 1$$

In particular:

$$\text{Size}(\tilde{x}) \leq \text{Size}(x) - 1 \quad \text{Size}(\phi(\tilde{z})) \leq \text{Size}(\phi(z)) - 1$$

Throughout the construction, we encode each qubit in a code block; namely, we restrict ourselves to using only a single codeword of the (d, d') -Toric code. From now on, we will think of $S' \subset S$ as a fixed d' -size subset of generators that were picked from the head. And we will decode the logical $|1\rangle$ state by $x_{S'}$. To simplify terminology, we call such usage 'nice encoding'.

Definition 4.39 (The Nice Encoding). *Consider (d, d') -Toric code and fix a subset $S' \subset S$, at size d' . We define the nice encoding to be the subsystem code, which uses $x_{S'}$ to encode a single logical qubit. Namely, the logical codewords are:*

$$\begin{aligned} |\bar{0}\rangle &= |\mathcal{C}_Z^\perp\rangle \\ |\bar{1}\rangle &= |x_{S'} + \mathcal{C}_Z^\perp\rangle \end{aligned}$$

We are now ready to give a full description of the realization of the logical operators. The logical Paulis $\mathbf{X}$ and $\mathbf{Z}$ are realized by the application of $X_{S'}$ and $Z_{S'}$. For the gates $\mathbf{S}$ and $\mathbf{T}$, we use single-qubit gate teleportation as presented in Figure 9. In single-qubit gate teleportation, the resource $|\mathbf{S}\rangle$ is initialized as $\mathbf{S}|+\rangle$, similarly for $|\mathbf{T}\rangle$. The gate $\mathbf{1}_{S'}$ stands for the classical procedure that, given the measurement results, checks whether the logical bit $|1\rangle$ is turned on. We will return to this gate with a higher resolution depiction of it after the description of the Hadamard teleportation.



Figure 9: Simulating the $\mathbf{S}$ and $\mathbf{T}$ gates, using the $|\mathbf{S}\rangle$, $|\mathbf{T}\rangle$ magic states.

For the Hadamard gate, we use a modification of the standard two-qubit gate teleportation as shown in Figure 10. The realization differs from the standard realization as the measurement in the phase basis is done by first moving to the dual space via the application of physical $\mathbf{H}^{\otimes |\mathcal{F}_{d'}|}$, measuring, reorganizing the qubits through ϕ , and then checking for the $x_{S/S'}$ outcome (thus, it can be thought of as the application of $\mathbf{H}_\phi$ and then measurement). In addition, unlike the single-qubit gate teleportation, here $|\mathbf{H}\rangle$ is initialized as $I \otimes \mathbf{H}(|00\rangle + |11\rangle)$ (up to normalization).




Figure 10: Gate teleportation of the Hadamard gate, using the resource $|\mathbf{H}\rangle$. We emphasize that although the Hadamard (at least Hadamard-type) has an $O(1)$ -depth gate realization (at least when $d' = d/2$), we still need all the operations on the quantum register to be strictly transversal for the sake of the proof. Thus, we realize the Hadamard via gate teleportation.

The classical procedure of logical codeword checking, namely $\mathbf{1}_{x_{S'}} / \mathbf{1}_{x_{S/S'}}$, is divided into five stages. In the first, we find all the connected components of the ∂^- -boundary. Finding connected components can be done in $O(\log^2(n))$ depth (see [SV82] and [AS87] for logarithmic time connected components in the PRAM model and [Lev80] for the simulation of PRAM machines using bounded-fan circuits). Then, we broadcast the connected components. Conceptually, each connected component is sent to its own Toric; technically, we copy components of the syndromes into clean registers. The broadcast also requires at most $O(\log(n))$ depth since there can be at most $\# \text{checks} = O(n^d)$ connected components and therefore at most $O(n^d)$ copies. Then, we decode each of the copies in parallel. The separation of the components guarantees that during decoding, components will not merge with each other. For the decoding, we iterate sweep-rule cycles (defined in Definition 4.6) until erasing the syndrome. Later, we will see that n cycles succeed to erase the syndrome with high probability¹². We erase the syndrome by combining the results of all the decoding sessions together by Xoring their correction suggestions. Each of the Xor operations is between at most $O(n^d)$ qubits and therefore requires at most $O(\log(n))$ depth. At this point, we assume that the combined register stores the measurement without errors. We finish by getting back into the logical space, namely a multiplication of binary matrix, which can be done at logarithmic depth $O(\log(n))$ (at the cost of broadcasting into $O(n^{2d})$ -wide again). So the overall depth for performing a logical measurement is $O(n)$. The additional classical width cost is due to the broadcasting which introduce at most $O(n^{2d})$ additional classical bits.

Here. We might want to say something about the control-X operation (like that it requires another broadcast). In addition we might want to refer to the Barrier.

5 Toom's Contours

5.1 The Barrier Argument

We are going to move between the infinite and the finite back and forth, so we will need to extend the size function, and the reduced notion to chains on the finite-complex.

Definition 5.1. *extending notions.* **Return to here.**

Claim 5.2. *Let $x \in (\mathcal{F}_d^n; \mathbb{F}_2)$ be a d' -dimensional chain, such that the $d' - 1$ -dimensional chain ∂^-x is not connected. Suppose that for some constant number $\alpha \in (0, 1)$, each connected component of ∂^-x has size at most αn . Then there are d' -dimensional chains $y_1, y_2, \dots$ such that each ∂^-y_i is a connected component of ∂^-x and $\mathbf{Size}(y_i) \leq \alpha n$.*

Proof. Consider a d' -dimensional chain x such that ∂^-x has a decomposition $x = \tilde{x}_1 + \tilde{x}_2$, where $\tilde{x}_1, \tilde{x}_2$ are $d' - 1$ -dimensional chains, and $\tilde{x}_1$ is a connected component of ∂^-x . In particular $\tilde{x}_1$ is not connected to $\tilde{x}_2$. By disconnectivity, $\tilde{x}_1$ does not share

¹²In the proof of the Theorem 6.3 we show that $\Theta(n)$ -size connected component of the syndrome is not likely to occur.

a $d' - 2$ dimensional face with $\tilde{x}_2$. Thus, $\partial^- \tilde{x}_1$ and $\partial^- \tilde{x}_2$ are disjoint. On the other hand, since $\partial^{-\circ 2} x = 0$, then combining with disjointness, it follows that:

$$\partial^-(\tilde{x}_1 + \tilde{x}_2) = 0 \Rightarrow \partial^- \tilde{x}_1 = \partial^- \tilde{x}_2 = 0$$

Namely, $\tilde{x}_1 \in \ker \partial^- : (\mathcal{F}_{d'-1}^n; \mathbb{F}_2) \rightarrow (\mathcal{F}_{d'-2}^n; \mathbb{F}_2)$. Suppose $\tilde{x}_1 \notin \text{Im } \partial^- : (\mathcal{F}_{d'}^n; \mathbb{F}_2) \rightarrow (\mathcal{F}_{d'-1}^n; \mathbb{F}_2)$, then we have that $\tilde{x}_1$ is a non-trivial codeword of the $(d, d' - 1)$ -Toric code. Since $d' \geq 2$, the size of $\tilde{x}$ is at least n , which is a contradiction. So we get that $\tilde{x}_1 \in \text{Im } \partial^- : (\mathcal{F}_{d'}^n; \mathbb{F}_2) \rightarrow (\mathcal{F}_{d'-1}^n; \mathbb{F}_2)$. In particular, there is a d' -dimensional chain x_1 such that $\partial^- x_1 = \tilde{x}_1$. Choose x_1 such it doesn't supported on a non-trivial codeword, namely $\tilde{x}_1$ is finite. The reduced form of x_1 has size at most $\text{Size}(\tilde{x}_1) \leq \alpha n$, and satisfies the desired condition for the connected component $\tilde{x}_1$. To complete, we repeat recursively on the connected components of $\tilde{x}_2$. $\square$

Definition 5.3. Let $x \in (\mathcal{F}_{d'}^n; \mathbb{F}_2)$ be a d' -dimensional chain. x will be said to be a fillable chain if each of the connected components of its ∂^- -boundary, namely $\partial^- x$, has size at most n . By Claim 5.2, there are d' -dimensional reduced chains $x_1, \dots$, such that each $\partial^- x_i$ is a connected component of $\partial^- x$. Thus:

$$\partial^-(x + \sum_i x_i) = 0 \Rightarrow x + \sum_i x_i \in \ker \partial^- : (\mathcal{F}_{d'}^n; \mathbb{F}_2) \rightarrow (\mathcal{F}_{d'-1}^n; \mathbb{F}_2)$$

Namely, $x + \sum_i x_i$ is a codeword in $\mathcal{C}_X / \mathcal{C}_Z^\perp$. We define the filling operator, denoted by $[x]$, which for a given fillable chain x , returns the homology class (logical operator) of $x + \sum_i x_i$.

Claim 5.4. Let $x_1, \dots, x_t \in (\mathcal{F}_{d'}^n; \mathbb{F}_2)$ be d' -dimensional chains, such that x_1 is a trivial chain and for any $i \in [t - 1]$, the Hamming distance between x_i and x_{i+1} is one. There is a constant $c \in (0, 1)$, depending only on d, d' , such that if n is sufficiently large, and for any $x_i \in x_1, \dots, x_t$, the size of each of the connected components of $\partial^- x_i$ is bounded by cn , then:

$$[0] = [x_1] = [x_2] = \dots = [x_t]$$

Proof. Assume through contradiction that there is i such that $0 = [x_i] \neq [x_{i+1}]$. Denote by $x_i^1, x_i^2, \dots$ and $x_{i+1}^1, x_{i+1}^2, \dots$ the reduced d' -dimensional chains such that $\partial^- x_i^j$ is a connected component of $\partial^- x_i$. Hence:

$$[x_{i+1}] = [x_{i+1}] + [x_i] \sim_{\mathcal{C}_{Z^\perp}} x_i + x_{i+1} + \sum_j x_i^j + \sum_j x_{i+1}^{j+1}$$

Denote by w the chain:

$$w := x_i + x_{i+1} + \sum_j x_i^j + \sum_j x_{i+1}^{j+1}$$

Since the Hamming distance between x_i and x_{i+1} is at most one, and since each d' -dimensional face contributes a constant support over $d' - 1$ -dimensional chains, the

∂^- -boundaries of x_i and x_{i+1} are at a Hamming distance that is at most constant. Therefore, there are only constant connected components of ∂^-x_i which are not connected components in ∂^-x_{i+1} and vice versa. Denote by $y_1, \dots, y_m$ the chains in $x_i^1, \dots$ and $x_{i+1}^1, \dots$ such that they belong to only one of the collections. Where m is bounded by a function of d, d' , namely $m \leq M(d, d')$. So:

$$\sum_j x_i^j + \sum_j x_{i+1}^{j+1} = \sum_j^m y_j$$

Moreover, the vertices of $y_1, \dots, y_m$ and $x_{i+1} + x_i$ have a nontrivial intersection. This is because the face $x_{i+1} + x_i$ is responsible for the splits and merges of the x_i 's ∂^- -boundaries. Therefore:

$$\mathbf{Size}(w) \leq \mathbf{Size}(x_i + x_{i+1}) + \sum_j^m \mathbf{Size}(y_j) \leq d + M(d, d')cn$$

Thus, for sufficiently large n , if $c < \frac{1}{M(d, d') + 1}$, we have that $\mathbf{Size}(w) < n$, a contradiction for w being a non-trivial codeword. $\square$

Corollary 5.5 (Barrier argument). *Barrier argument.*

5.2 Toom's Contours

We are about to zoom out from a single register encoded in the Toric code to describe a fully fault-tolerant computation. For a given quantum circuit, assumed to consist of gates **X**, **Z**, **CX**, **H**, **S**, **T** and measurements, we define its Toric encoded circuit as its version where each qubit is encoded by the Toric code. Formally:

Definition 5.6 (Toric Encoded Circuit). *Let C be a quantum circuit. We say that C is a Toric encoded circuit if the following holds:*

1. *Any layer of the circuit consists of classical and quantum registers, such that there is a Toric code $\mathcal{C}$ such that the quantum registers are encoded by the nice encoding of $\mathcal{C}$. At time $t = 0$, the quantum registers are initialized to be either $|\bar{0}\rangle$, encoded $|\mathbf{H}\rangle$, encoded $|\mathbf{S}\rangle$, or encoded $|\mathbf{T}\rangle$. The classical registers are initialized to be zeros.*
2. *The layers of C on the quantum registers at even times are sweep-cycles, and the layers at odd times are realizations of either logical Paulis, **CX**, or measurements in the computational and parity bases. On the classical registers, including the measured register, we allow any two-bit gates.*

For Toric encoded circuits, we define their *base graph*, which can be thought of as a recording of a running from the Toric's point of view. The *base graph* is obtained by taking the Toric's vertices at different times as the vertices and connecting any two sequential in time if they are support faces — qubits — which are either the same or are entangled by a quantum gate.

Definition 5.7 (Base graph G_o). Let C be a Toric encoded circuit at depth μ . Notice that measurements take quantum registers into classical, Thus any quantum register has a maximal depth, which we denote by $\{\mu_i\}$. Let's denote by $\{\mathcal{G}_i\}$ the skeleton of the ∞ -extension of the Toric complexes that induce the codes. We define the base graph of C , denoted by $G_o = (V_o, E_o)$, as follows:

1. The vertices V_o are tuples, where the first coordinates correspond to a vertex in the disjoint union over the vertices in $\{\mathcal{G}_i\}$, and the second coordinate is associated with a time. Namely:

$$V_o := \bigcup_i (V(\mathcal{G}_i) \times [\mu_i])$$

2. The edges E_o are derived from the original edges in $\mathcal{G}_1, \dots, \mathcal{G}_k$, with the addition that any two vertices in preceding times t and $t + 1$ that support qubits (faces) such that C^∞ acts non-trivially on those qubits at time t are also connected by an edge. Namely:

$$E_o := \{ \{(v, t), (u, t')\} : \text{if} \\ \text{either there exists } i \text{ such } \{v, u\} \in \mathcal{G}_i \text{ and } t' = t + 1 \\ \text{or there exists } u_i = (g_i, l_i) \in C^\infty \text{ such } v, u \in l_{i,2} \text{ and } t = l_{i,1}, t' = l_{i,1} + 1 \}$$

Where we define a vertex v to belong to location l if there is a qubit-face f in l such that v belongs to its associated face.

As a concept, if the base graph is a recording of a run from the Torices' point of view, then a *Toom's contour* is a recording of the deviation's evolution in the eyes of the Torices' vertices. We use the notion of Pauli-path to describe a deviation:

Definition 5.8. Let P be a Pauli. We denote by P^∞ the view of P in the ∞ -extension. We denote by $x(P^\infty)$ the d' -dimensional chain in the infinite complex that supports the X -component of P^∞ , and by $z(P^\infty)$ the $d - d'$ -dimensional chain, such that the inversion of the ϕ -transpose sends it to the d' -dimensional chain that supports the Z -component of P^∞ , namely $\phi^{-1}(z(P^\infty))$ is the chain that supports the Z -component.

Consider removing. We extend the notion to Paulis in finite spaces (for example, Paulis on the circuit's register) by saying that a Pauli $P \in \mathcal{P}^{(\cdot)}$ is reduced if its view in the ∞ -extension is reduced.

Definition 5.9. Consider a quantum circuit subjected to noise $C[\mathcal{E}]$, with width w and depth μ , whose registers are all encoded in the (d, d') -Toric code with side length n such that $C[\mathcal{E}]$ has an ∞ -extension. We say that the set of Pauli $\{(P_i)\}_1^{\mu+1}$, where for each i , $P_i \in \mathcal{P}^{[w]}$, is a Pauli-path for $C[\mathcal{E}]$. If P_1 is the identity and for any $t \in [\mu]$, the Pauli P_{t+1} is defined such it satisfies the follow:

$$P_{t+1} \in C|_t \circ P_t$$

Consider removing. If there is a Pauli-path $P_1, \dots, P_{\mu+1}$ and a time $t \in [\mu + 1]$ such P_t is infinite, then we say that time t admits an infinite error.

For a Pauli-path $P_1, \dots, P_{\mu+1}$, the Toom's contour is the restriction of the base graph to vertices that witness a deviation, namely, those that belong to qubits supporting $P_1, P_2, \dots, P_{\mu+1}$.

Definition 5.10 (Toom's Contour Graph of $C(\mathcal{E})$, (Standard form)). *For a Toric encoded quantum circuit C , set of faults $\mathcal{E}$ and a Pauli-path $P_1, \dots$, let us denote the base graph of $C[\mathcal{E}]$ by G_o . We define the Toom's contour graph $\mathcal{H} = (V_{\mathcal{H}}, E_{\mathcal{H}})$ as follows:*

1. The vertices $V_{\mathcal{H}} = V_{\mathcal{H},X} \cup V_{\mathcal{H},Z}$ correspond to vertices of faces supporting the Pauli-path. They consist of the triple $(v, t, \cdot)$ where (v, t) is a vertex of the base graph, and the third is a symbolic coordinate that is either $\{X\}$ or $\{Z\}$ corresponding to the component of the Pauli that v supports at time t .

$$V_{\mathcal{H},X} := \{v_t \times \{X\} : v_t = (v, t) \in V_o$$

$$\text{and there exists a face } f \in \partial^- x(P_t) \text{ such } v \in f\}$$

$$V_{\mathcal{H},Z} := \{v_t \times \{Z\} : v_t = (v, t) \in V_o$$

$$\text{and there exists a face } f \in \partial^- z(P_t) \text{ such } v \in f\}$$

2. The edges $E_{\mathcal{H}}$ is partitioned into two types $A_{\mathcal{H}}$ and $F_{\mathcal{H}}$, where for any logical gate g which applied in time t in $C[\mathcal{E}]^\infty$ we add edge to $A_{\mathcal{H}}$ as follow:

- (a) If g is either a logical Pauli, a measurement in either the computational or the phase basis, a sweep-cycle, or an identity, then for any face f in the register on which g acts, we add:

$$\{\{\tilde{u} = (u, t, \cdot), \tilde{v} = (v, t+1, \cdot)\} : u, v \in f \text{ and } \tilde{v}, \tilde{u} \in V_{\mathcal{H},X}$$

$$\text{or } u, v \in \phi(f) \text{ and } \tilde{v}, \tilde{u} \in V_{\mathcal{H},Z}\}$$

- (b) If g is a logical **CX** then for any physical **CX** in its realization, between faces f_1 into f_2 at time t we add:

$$\{\{\tilde{u} = (u, t, \cdot), \tilde{v} = (v, t+1, \cdot)\} : u \in f_1 \cup f_2, v \in f_1 \cup f_2 \text{ and } \tilde{v}, \tilde{u} \in V_{\mathcal{H},X}$$

$$\text{or } u \in \phi(f_1) \cup \phi(f_2), v \in \phi(f_1) \cup \phi(f_2) \text{ and } \tilde{v}, \tilde{u} \in V_{\mathcal{H},Z}\}$$

- (c) If g is a fault Pauli, then for any face f in the register, we add:

$$\{\{\tilde{u} = (u, t, \cdot), \tilde{v} = (v, t+1, \cdot)\} : u, v \in f, \tilde{v}, \tilde{u} \in V_{\mathcal{H},X}, \text{ and } f \notin X(g)$$

$$\text{or } u, v \in \phi(f), \tilde{v}, \tilde{u} \in V_{\mathcal{H},Z}, \text{ and } f \notin Z(g)\}$$

To construct $F_{\mathcal{H}}$, we connect any pair of vertices (v, t) , (u, t) if there exists a vertex $(w, t+1)$ such both of them are connected to it through $A_{\mathcal{H}}$. Namely:

$$F_{\mathcal{H}} = \{e = \{u, v\} : \text{exists } w \in V_{\mathcal{H}} \text{ such } \{v, w\}, \{u, w\} \in A_{\mathcal{H}}\}$$

We call $A_{\mathcal{H}}$ and $F_{\mathcal{H}}$ the arrows and forks edges respectively.

Remark 5.11. *One might expect to have a differentiation between measurement in the computational and the phases basis. Namely, a rule of connecting arrows of the form:*

If g is a measurement in the phases base, then for any face f in the register on which g acts, we add:

$$\begin{aligned} & \{ \{ \tilde{u} = (u, t, \cdot), \tilde{v} = (v, t+1, \cdot) \} : u \in f, v \in \phi(f) \text{ and } \tilde{u} \in V_{\mathcal{H},X}, \tilde{v} \in V_{\mathcal{H},Z}, \\ & \text{or } u \in \phi(f), v \in f \text{ and } \tilde{u} \in V_{\mathcal{H},Z}, \tilde{v} \in V_{\mathcal{H},X} \} \end{aligned}$$

Such intuition is right in the sense that this is the right way to keep track of the faults. Yet, later we will state Claim 5.32 whose corollary is that instead of analyzing the running of the Toric encoded circuit, it is sufficient to analyze the running of a 'simpler' circuit which doesn't contain measurement in the phase base. In particular, the claim considers the circuit obtained from a given Toric encoded circuit by replacing each of the classical decoding procedures with noisy sweep-rule cycles. Then it states that to derive an upper bound on the failure probability of the original circuit, it's enough to bound the probability that the modified circuit does not admit an infinite 'time-connected cluster' (slice, Definition 5.16).

Definition 5.12. *With the same notations as in Definition 5.9, let τ be a step in $C[\mathcal{E}]$. Recall that each gate in $C[\mathcal{E}]$ acts on at most Δ qubits. Let $\Delta' \in [\Delta]$. In addition, denote by g and t the gate and its time which is applied at step τ . We say that the set of Pauli $P_1, \dots, P_t, P_\star$ is a Pauli-path up to step (τ, Δ') , if the $P_1, \dots, P_t$ is a Pauli-path for the circuit obtained from $C[\mathcal{E}]$ by taking only its first $t-1$ layers. And $P_\star$ satisfies the following: let $g'_1, \dots, g'_t = g$ be the gates applied at time t at step up to, and including, τ , placed by order. Let $P'_0, P'_\Delta \in \mathcal{P}^{(\cdot)}$ be Paulis such:*

$$P'_0 \in (g'_{t-1}g'_{t-2} \dots g'_1) \circ P_t \quad P'_\Delta \in g \circ P'_0$$

Now, since P'_Δ is at a Hamming distance of at most Δ from P'_0 , there is a set of Paulis:

$$P'_0, P'_1, P'_2, \dots, P'_{\Delta'}, \dots, P'_\Delta$$

Then set $P_\star = P'_{\Delta'}$.

Rewrite. For the analyses, we would like to extend Toom's contour to record in their final layer a partial time-step, namely a computation up to a certain step Definition 3.2. For this, we substitute in the definition above the Pauli-path with a Pauli-path up to a time **Prefix?**. Moreover, since we would like to refer to the connected chain at the final layer as the 'final error clusters', we add at the end a layer whose purpose is to connect in time (by arrows) the connected components of Toom's contour in the prior layer, which are connected by checks (their union is identified with a connected chain).

Definition 5.13 (Analysis Toom's Contour Up to Step (τ, Δ')). For a Toric encoded quantum circuit C , a set of faults $\mathcal{E}$, a step τ according to the standard order, and an integer $\Delta' \in [\Delta]$, we define the analysis Toom's contour up to step (τ, Δ') , denoted by $\mathcal{H}_{\tau, \Delta'} = (V_{\mathcal{H}}, E_{\mathcal{H}})$, to be the graph obtained by taking a Pauli-path up to step (τ, Δ') $P_1, \dots, P_t, P_*$, and gluing to the Toom's contour induced by the prefix $P_1, \dots, P_t$ a layer defined by adding vertices and edges as follows:

1. We add to the vertices in $V_{\mathcal{H}, \{X\}}$ the vertices that support P_* , namely the subset:

$$\begin{aligned} & \{v_{t+1} : v_{t+1} = (v, t+1, \{X\}) \text{ if there exists a face } f \in \partial^- x(P_*) \\ & \text{such } v \in f \in \partial^- x(P_*)\} \end{aligned}$$

Similarly we add to $V_{\mathcal{H}, Z}$ the subset:

$$\begin{aligned} & \{v_{t+1} : v_{t+1} = (v, t+1, \{Z\}) \text{ if there exists a face } f \in \partial^- z(P_*) \\ & \text{such } v \in f\} \end{aligned}$$

2. For the edges, we add the following arrows:

$$\begin{aligned} & \{\{\tilde{u} = (u, t+1-1, \cdot), \tilde{v} = (v, t+1, \cdot)\} \text{ if either there exists} \\ & \text{a } d' - 2 \text{ dimensional face } f \text{ such } \tilde{u}, \tilde{v} \in V_{\mathcal{H}, X} \text{ and } u, v \in f \\ & \text{or there exists a } d - d' - 2 \text{ dimensional face } f \text{ such } \tilde{u}, \tilde{v} \in V_{\mathcal{H}, Z} \text{ and } v, u \in f\} \end{aligned}$$

3. Finally, we add forks for vertices at time $t+1$ to preserve the property that any two vertices with the same time coordinate, which are connected to a third (with a $+1$ higher time coordinate) via arrows, are also connected by a fork.

5.3 Space-Time Graph

Not a few of the results about Tooms contours are derived from their 'directed' structure, regardless of their connection to Torices. Thus, we split the discussion to talk about *space-time graphs*, which are graphs whose edges can be partitioned into arrows and forks.

Here, we present the results and proofs first given by Berman and Simon [BS88] and later by Gacs [Gac21] about *space-time graphs*. Although the proofs we give are almost identical, we chose to present them rather than referring to the sources. The reason for this is that in the language appearing in [BS88] and [Gac21], vertices are bits¹³, so conceptually their Toom's contours track different creatures. Thus, jumping into their text while keeping our image in mind might cause confusion, which we would like to avoid.

Definition 5.14. Consider a graph $G = (V, A, F)$, where V is a set of vertices, and A and F are sets of undirected edges. Let $T: V \rightarrow \mathbb{N}$. We say that (G, T) is a *space-time graph* if and only if:

¹³It is also true that the gates are different, yet that difference alone does not constitute a conceptual change.

1. $\{x, y\} \in A \Rightarrow T(y) = T(x) \pm 1$
2. For $x, y \in V$ $\{x, y\} \in F$ if and only if there exists $z \in V$ such both $\{x, z\}, \{y, z\} \in A$, and $T(x) = T(y) = T(z) - 1$.

Additionally, for any vertex $v \in V$ of the space-time graph, denote by $\text{pre}(v)$ the set of vertices defined as follow:

$$\text{pre}(v) := \{u \in V : \{u, v\} \in A \text{ and } T(v) = T(u) + 1\}$$

We extend the notion for subsets of vertices $X \subset V$ by $\text{pre}(X) = \bigcup_{v \in X} \text{pre}(v)$. We will call T the time coordinate, A arrows, and F forks. Observe that from the second requirement, it follows that forks can connect only vertices that have the same time value (set by T).

Remark 5.15. The Toom's contour graph defined in Definition 5.10 is, by definition of it's construction, a time-graph.

The history of a vertex consists of all the vertices that are connected to it by arrows. Here we allow passing through arrows in both directions, so it might be possible that the history of vertex u contains vertices with the same time coordinate as u . That motivates to define *slices*. A *slice* is a (maximal) subset of vertices with the same time coordinate such they share the same history. In [BS88], slices take the role of error droplets.

Definition 5.16 (History, Slice). For a time graph (G, T) and a vertex $u \in V$, let the history of u , denoted by H_u , be the restriction of G to vertices that both precede u in time and connect to u via path of arrows. A slice is a maximal subset of vertices $K \subset V$ such that for any two vertices $u, v \in K$, $H_u = H_v$ is satisfied. As a consequence, the time coordinate in K has to be constant. Thus, we can extend the time coordinate and the history to slices and write for a slice K , $T(K)$ and H_K .

The following claim ensures us an 'independent' argument about slices. Here we state that two slices with the same time coordinate have to hold disjoint histories. Another way to think about this is that if the first and the second slices are 'induced by n_1 and n_2 faults,' then the slice obtained by merging them 'is induced by $n_1 + n_2$ faults.'

Claim 5.17. Let K_1, K_2 be two slices such that $T(K_1) = T(K_2)$, then either $K_1 = K_2$ or H_{K_1}, H_{K_2} are disjoint.

Proof. Assume a non-empty intersection, so there exists $z \in H_{K_1} \cap H_{K_2}$ such that for any $x \in K_1$ and $y \in K_2$, there exist arrow-paths $z \rightarrow x$ and $z \rightarrow y$. Thus, for any $w \in H_{K_1}$, one can concatenate a path $w \rightarrow x \rightarrow z \rightarrow y$ and conclude that $y \in H_{K_2}$. $\square$

Next, we would like to study how slices look a moment before they merge. For this, we define the *graph of a slice*. We construct this graph by taking a slice and considering all the slices that precede it by a single time unit. We take those slices as vertices and connect any two of them that have vertices sharing a fork.

Definition 5.18 (The Graph of Slice). *Let K be a slice. The graph of K , denoted by $G_K = (V_K, E_K)$, is the graph whose vertices are slices $R \subset H_K$ with $T(R) = T(K) - 1$. Two vertices are connected by an edge if and only if there is a fork whose ends belong to the slices associated with them. Namely, $S, R \in V_K$ are connected in G_K if and only if there is $\{s, r\} = f \in F$ such that $s \in S$ and $r \in R$.*

The graph of slices is connected. That fact means that a spanning tree of a (subgraph of) a slice's history can be found by combining the spanning trees of its processors. The following gives the proof:

Claim 5.19. *G_K is connected.*

Proof. Let $R, S \in V_K$ and consider $r, s \in V$ such that $r \in R, s \in S$. Recall that $R, S \subset H_K$ and therefore $r, s \in H_K$. Thus, by definition of the 'History' there exists a path, passing through arrows only, from r to s in H_K . Consider a path that is minimal in the number of vertices with a time coordinate equal to $T(K)$; denote that number by k . We prove by induction on k that R and S are connected in G_K .

1. Base. $k = 0$, In this case, any vertex on the path between r and s has an equal or lower time coordinate than r and is also connected to r via a path of arrows, namely belongs to H_r . In particular, s is also a vertex on the path, and therefore $s \in H_r \Rightarrow H_r = H_s$, so $R = S$.
2. Induction step. Since $k > 0$, there exists a vertex y on the path from r to s such that $T(y) = T(K)$. Let x be the vertex which precedes y on the path, and z be the vertex that follows y . Since the values of T at the two ends of an arrow differ by 1, and since y gets the maximal T value in H_K it follows that:

- (a) $\{x, z\} \subset \text{pre}(y) \Rightarrow \{x, z\} \in F$.
- (b) The slices of x and z , say K_x and K_z has time value $T(K_x) = T(K_z) = T(K) - 1$, Namely they are both members of V_K .

Since the path connecting r, x in H_K has $k - 1$ vertices with $T(\cdot) = T(K)$, then by the induction hypothesis, R and K_x are connected in G_K , and by similar argument so are K_z and S . On the other hand, $\{x, z\} \Rightarrow \{K_x, K_z\} \in E_K$. Thus, R and S are connected in G_K .

□

Now, we would like to plug in the fact that the vertices originate from the infinite plane. To do so, we equip slices with poles. The poles are the vertices in the slice, which are the maximal points of potential walls.

Definition 5.20 (Poles Selection). *Let (G, T) be a space-time, whose vertices are associated with points in $\mathbb{R}^d$. Let A be either a slice, arrow, or a fork in (G, T) .*

We define a poles selection for A to be a selection of $d + 1$ vertices, $v_1, \dots, v_{d+1}$, not necessary different, from A . The size of A is defined to be:

$$\mathbf{Size}(A) := \sum_i^{d+1} \max_{v \in A} L_i(v)$$

Thus, for any selection of poles $v_1, \dots, v_{d+1}$ we have that:

$$\mathbf{Size}(A) \geq \max_{v \in A} L_i(v_i)$$

Consider splitting the definitions. For a vertex u , we define the i th predecessor of a u , denoted by $\text{pre}_i(u)$, to be its predecessor which achieves the maximal value of L_i , namely:

$$\text{pre}_i(u) := \arg \max_{v \in \text{pre}(u)} L_i(v)$$

The following claim relates the size of a given slice to the sizes of its predecessor slices.

Claim 5.21. Consider a time graph (G, T) whose vertices are associated with points in $\mathbb{R}^d$. Let K be a slice in G , such that $K \neq H_K$. Let $v_1, \dots, v_{d+1}$ be a pole selection for K . Then for some integer k , there exist $k + 1$ slices $K_0, K_1, \dots, K_k$ and k forks $F_1, \dots, F_k$ with pole selections for them that satisfy:

1. For any $i \in [d + 1]$, the i th predecessor $\text{pre}_i(v_i)$ belongs to the selected poles of $K_0, \dots, K_k$.
2. Consider the graph whose vertices are associated with $K_0, \dots, K_{k+1}$, and there is an edge between K_i and K_j if and only if there is a fork in $F_1, \dots, F_k$ that intersects both K_i 's and K_j 's selected poles. Then the graph formed from $K_0, \dots, K_k$ is connected.
3. For any $j \in [0, k]$, denote by $v_j^1, \dots, v_j^{d+1}$ the selected poles for K_j , then:

$$\sum_{j=0}^k \sum_{i=1}^{d+1} L_i(v_j^i) + \sum_{i=1}^k \mathbf{Size}(F_i) \geq \sum_{i=1}^{d+1} L_i(\text{pre}_i(v_i))$$

Proof. First, notice that K doesn't contain leaves in (V, A) , since a leaf is connected only to itself via arrows, and therefore its 'slice' equals its 'History'. Thus, $\text{pre}_i(v_i)$ is not empty for any $i \in [d + 1]$. Next, consider the graph $G_K = (V_K, E_K)$ from Definition 5.18. Since $\{v_i, \text{pre}_i(v_i)\}$ is an arrow, $\text{pre}_i(v_i)$ belongs to a slice in V_K . For $i = 1, \dots, d + 1$, denote by R_i the slice in V_K that contains $\text{pre}_i(v_i)$.

By Claim 5.19 G_K is connected and therefore there exist a minimal collection $K_0, \dots, K_k$ of slices from V_K which contains $R_1, R_2, \dots, R_{d+1}$, and which is connected in G_K . We pick a set of forks $\mathbf{F} = \{F_1, \dots, F_k\}$ which realizes a spanning tree for this collection of slices.

Let us denote by $J_0, \dots, J_{2k}$ the enumerated union $K_0, \dots, K_k, F_1, \dots, F_k$. Observe that for any choice of $(d+1)(2k+1)$ elements $z_0^1, z_0^2, z_0^3, \dots, z_0^{d+1}, \dots, z_{2k}^{d+1}$ in $\mathbb{R}^d$, such that $z_j^i \in J_j$ it holds that:

$$\sum_{j=0}^k \sum_{i=1}^{d+1} L_i(z_j^i) + \sum_{i=1}^k \text{Size}(F_i) \geq \sum_{j,i} L_i(z_j^i) \quad (1)$$

Later, we are going to take the points $\{z_j^i\}$ from the intersection of J_j 's. For convenience, let us assume that for the index subset $X \subset [2k] \cup \{0\}$, the intersection $\bigcap_{j \in X} J_j$ returns a single point. If there is more than one, then decide on the point in a way that all the chosen points simultaneously satisfy that for any index subsets Y, X , such that $X \cap Y$ and they both introduce a non-trivial intersection $\bigcap_{j \in Z} J_j : Z \in \{X, Y\}$, then:

$$\bigcap_{j \in X} J_j = \bigcap_{j \in Y} J_j$$

Now consider a non-trivial maximal intersection $\bigcap_{j \in X} J_j$. The intersection is maximal in the sense that for any $j' \notin X$, it holds that $\bigcap_{j \in X \cup \{j'\}} J_j = \emptyset$. We claim that for any i and for any such X , there is $t \in X$ such that J_t is the successor of J_j on the path to R_i for any $j \in X/\{t\}$. Assume not. Let $x, y \in X$, and denote the paths from J_x, J_y to the R_i by $J_x, J_{x_2}, \dots, R_i$, and $J_y, J_{y_2}, \dots, R_i$, such that $J_y \neq J_{x_2}$. If $y = x_2$, $y_2 = x$, or $x_2 = y_2 \in X$, then pick other x, y (if there is no such, we get an immediate contradiction). In that case, $J_x, J_y, J_{y_2}, \dots, R_i$ is a path from J_x to R_i which is different from $J_x, J_{x_2}, \dots, R_i$. Therefore, there are at least two paths from J_x to the R_i , namely, there is a cycle in contradiction to the fact that the collection $K_0, \dots, K_k, F_1, \dots, F_k$ is a tree.

We will choose the z_j^i by the following process. Pick a z_j^i that has not been set yet. If $J_j = R_i$, then set $z_j^i \leftarrow \text{pre}_i(v_i)$ - we call this a type-1 assignment. Otherwise, consider the path from J_j to R_i (by connectivity, it has to be such a one, and by minimality, there is exactly one). Let J_t be the successor on the path, then pick $z_j^i \in J_j \cap J_t$ - similarly, we call this assignment a type-2 assignment. Since any type-2 assignment is associated with a non-trivial intersection, and any non-trivial intersection induces a type-2 assignment for all i 's (since for all i 's, one of the J_j on its support is a successor of the rest on their path to the R_i), we have that:

$$\begin{aligned} \sum_{j,i} L_i(z_j^i) &= \sum_{j,i,z_j^i \in \text{type-1}} L_i(z_j^i) + \sum_{j,i,z_j^i \in \text{type-2}} L_i(z_j^i) \\ &= \sum_i^{d+1} L_i(\text{pre}_i(v_i)) + \sum_{X: \bigcap_{j \in X} J_j \neq \emptyset} (|X| - 1) \sum_i^{d+1} L_i \left(\bigcap_{j \in X} J_j \right) \\ &= \sum_i^{d+1} L_i(\text{pre}_i(v_i)) \end{aligned} \quad (2)$$

To complete, we select the i th pole of K_j to be z_j^i . Then, for $i = 1, \dots, d+1$, the i th predecessor $\text{pre}_i(v_i)$ is the i th pole of R_i . This assures (1). Since for any intersection

$K_i \cap F_j$ a type-2 assignment is picked from the intersection to K_i , we have that the type-2 selected poles of $K_0, K_1, \dots, K_k$ are covered by the vertices of $F_1, F_2, \dots, F_k$. This assures (ii). Plugging Equation (2) into Equation (1) assures (3). $\square$

Until now, the claims stated were true for any kind of space-time graph (where their vertices can be associated with coordinates in $\mathbb{R}^d$). The next claims assume that the given time graph is a Toom's contour induced by the noisy running of a Toric encoded circuit.

Claim 5.22. *For Toom's contours, the maximal size of a fork is at most $2d$.*

Proof. Consider the rules of forks connecting in Definition 5.10. For forks obtained by rules (a), (b), or as a result of a fault, it's clear that their vertices at their ends are on intersecting faces (when projecting the vertices on $\mathcal{G}$). The maximal fork, in those cases, has the form of:

$$F = \{v, (\prod_{g \in S'_1} g)(\prod_{g \in S'_2} g)v\}$$

Where S'_1 and S'_2 are subsets of S of size either d' or $d - d'$, and therefore the size of such forks is at most $\mathbf{Size}(F) = 2 \max\{d', d - d'\} \leq 2d$. $\square$

While in the construction presented in [BS88], all the layers were the same (consisting of the same gates up to wiring), here different layers might differ both in realization and role. Thus, we would like to distinguish between them. We will define the *predecessor decomposition* of a slice to be the collection of slices and forks which satisfy the condition in Claim 5.21. The order of a slice is a number in $\{0, 1, 2, 3\}$ that matches the time modulo four of the layer that results in the slice. Therefore, it encodes the layer type.

Definition 5.23. *Let K be a slice, and let $v_1, \dots, v_{d+1}$ be a selection of poles for it. Then, let $K_0, \dots, K_k$ be slices in H_K , and let $F_1, \dots, F_k$ be forks, such that $K_0, \dots, K_k, F_1, \dots, F_k$ satisfy the conditions in Claim 5.21 with respect to the poles selection. We call $K_0, \dots, K_k, F_1, \dots, F_k$ the predecessor decomposition of K with respect to the pole selection $v_1, \dots, v_{d+1}$. When the pole selection is either clear from the context or can be chosen arbitrarily, we will omit it.*

Furthermore, we say that K is a result of a sweep-cycle, faults, or a logical operation if the layer at time $T(K) - 1$ consists of sweep-cycle gates, fault gates (whose source is $\mathcal{E}$), or either measurement or logical Pauli. The order of the slice is its time plus 3 modulo 4, namely $\text{order}(K) = T(K) + 3 \pmod{4}$. So the orders of slices that result from a sweep-cycle, faults, logical operation, and again faults are 0, 1, 2, 3 respectively.

The distinction between slices resulting from different layers allows us to adjust Claim 5.21 by conditioning on the slice type. Yet, before doing so, we will have to extend the infinity characterization to slices.

Do something with the infinite slices.

Definition 5.24 (Inifinite Slice). *Let K be a slice in a Toom's contour. We will say that K is infinite slice if:*

$$\mathbf{Size}(K) \geq n$$

That motivates us to define the notion of a step that is the first to admit an infinite slice/error.

Definition 5.25 (Time t Admits Inifinte Error/Slice). *Let $C[\mathcal{E}]$ be a Toric encoded circuit with depth μ . We say that a time $t \in [\mu + 1]$ admits an infinite error if there exists a Pauli path, $P_1, P_2, \dots$, such that in its prefix up to time t , there is at least one Pauli in the path whose view in the ∞ -extension is infinite (Definition 4.25). Consider the analysis of Toom's contour that is induced by the prefix. We say that a time t admits an infinite slice if there is a Pauli path such that $\mathcal{H}_t$ has an infinite slice (Definition 5.24). Moreover, we say that t is the first time to admit an infinite error/slice if t is the minimal time for which there is a Pauli path such that $\mathcal{H}_t$ admits an infinite error/slice. In exactly the same way, we define the notation of time t admitting an error/slice at a given size.*

Corollary 5.26. *Consider a slice K in it, and a selection of poles $v_1, \dots, v_{d+1}$, such that $K \neq H_K$. Since $K \neq H_K$, it follows that K has a predecessor decomposition with respect to $v_1, \dots, v_{d+1}$, denoted by $K_0, \dots, K_k, F_1, \dots, F_k$. For each slice K_j , denote its selected poles by $z_j^1, \dots, z_j^{d+1}$. The inequality in Claim 5.21 becomes:*

1. *If K is a result of a sweep-cycle, then:*

$$\sum_{j=0}^k \sum_i^{d+1} L_i(z_j^i) + \sum_{i=1}^k \mathbf{Size}(F_i) \geq \sum_i^{d+1} L_i(v) + 1$$

2. *If K is a result of either faults or a logical operation, then:*

$$\sum_{j=0}^k \sum_i^{d+1} L_i(z_j^i) + \sum_{i=1}^k \mathbf{Size}(F_i) \geq \sum_i^{d+1} L_i(v)$$

Proof. First, since the arrows and forks connection rules do not allow vertices of $V_{\mathcal{H},X}$ to connect to vertices of $V_{\mathcal{H},Z}$, we can assume that $K \subset V_{\mathcal{H},X}$. Let us split upon what K results from:

1. For the sweep-cycle part, we use Corollary 4.38 and Claim 4.31. The first states that within an application of the sweep-cycle over a chain, the maximal L_{d+1} value of the ∂^- -boundary of the resulting chain decreases by at least one, while the values of each L_g do not increase. The second assures that for any local maximum of L_i in the resulting ∂^- -boundary, there is a vertex in the ∂^- -boundary of the original such that they (1) share a d' -dimensional face and (2) the L_i -value of the vertex in the result is at most the L_i -value of the vertex in the original, where if $i = d + 1$, the relation becomes strictly less. Thus, prior to forming the spanning tree of the predecessor slices, for every v_i in $v_1, \dots, v_{d+1}$ we select a vertex $\text{pre}_i(v_i)$ that has either at most the same L_i -value as v_i if $i \neq d + 1$ or an L_{d+1} -value that is bigger by at least 1. That gives the inequality.

2. If K is a result of faults, then it cannot contain a vertex that supports an X -fault. That is because no arrow connection is made in that case, and hence, if it were to contain a vertex that supports a X -fault, that vertex would be the only one in K , namely $K = H_K$. So, we can imagine that all the vertices in K support the identity operator (or Z -faults, which are treated as identity for vertices in $V_{\mathcal{H},X}$) and handle that case as a result of a logical operator.
3. Where K is the result of a logical operator, we use that the operators are transversal¹⁴, so we have that the support K is contained in the support of $\cup_0^k K_i$, and therefore for any $i \in [d+1]$:

$$\max\{L_i(v) : v \in K\} \leq \max\{L_i(v) : v \in \bigcup_{i=0}^k K_i\}$$

□

The following claim extends the corollary to slices at the final layer of an *analysis Toom's contour up to step* (τ, Δ') . **Understand if the admitting is required here.**

Claim 5.27. *For a step τ and integer $\Delta' \in [\Delta]$, let $\mathcal{H}_{\tau, \Delta'}$ be Toom's contour analysis up to step (τ, Δ') . Denote by t the time at which the gate at step τ is applied. Then, for a slice K and selection for its poles $v_1, \dots, v_{d+1}$ such that $T(K) = t + 1$ and $K \neq H_K$, the inequality in Claim 5.21 becomes:*

$$\sum_{j=0}^k \sum_i^{d+1} L_i(z_j^i) + \sum_{i=1}^k \mathbf{Size}(F_i) \geq \sum_i^{d+1} L_i(v) - 2d$$

Proof. We allow vertices of $V_{\mathcal{H},X}$ at the final time layer to be connected by an arrow if their counterpart on the skeleton shares a $d' - 2$ dimensional face. Similarly, two vertices of $V_{\mathcal{H},Z}$, such that one of them is at the top layer, are connected if their images in the map ϕ^{-1} share a $d - d' - 2$ dimensional face. Therefore, for any $i \in [d]$, the value of L_i for v_i might be greater by at most 1 than the value of $\text{pre}_i(v_i)$, and the L_{d+1} for v_{d+1} might be greater by at most $\max\{d' - 2, d - d' - 2\}$ than the value of the $d + 1$ pole of K . Hence, overall, we get that:

$$\begin{aligned} \sum_{i=1}^{d+1} L_i(\text{pre}_i(v_i)) &\geq \sum_{i=1}^d (L_i(v_i) - 1) + L_{d+1}(v_{d+1}) - \max\{d' - 2, d - d' - 2\} \\ &\geq \mathbf{Size}(K) - 2d \end{aligned}$$

□

¹⁴In fact, a pure logical operator doesn't change the slices, yet we analyze by the containment argument to capture also the case of resulting of faults.

5.4 Minimal Explantation Argument

We are about to prove the main lemma stating that the number of faults causing an occurrence of slice is linear in both the slice's size and the minimal number of edges required to span a Toom's contour that induces it. First, we define the explanation tree of a slice to be the subtree of its history that connects its poles. The relation between the leaves of the subtree and its arrows, forks, and the size of its origin slice is embedded as a requirement in the definition. Then, we prove that, assuming no time admits an infinite error, any slice has an explanation tree.

Definition 5.28. *Consider a Tooms contour graph induced by Toric encoded circuit, with sweep-cycle parameter $\xi = 1$, and let α, β be constants and $\mathcal{X}_0, \mathcal{X}_1, \mathcal{X}_2, \mathcal{X}_3 \in \mathbb{R}$ be quadruple, defining such:*

$$\alpha := 40(d+1)^2, \quad \beta := 5(d+1)$$

$$\mathcal{X}_r := 1 - r \frac{d+1}{\alpha} = 1 - \frac{r}{40(d+1)}$$

In addition, we would like to refer to $\mathcal{X}_0, \mathcal{X}_1, \mathcal{X}_2, \mathcal{X}_3$ in a cyclic manner, so we define:

$$\mathcal{X}_{-1} := \mathcal{X}_3$$

Given sets W of points, A of arrows, and F of forks, we define U as the set of vertices of the graph induced by the edges of $A \cup F$. The sets W , A , and F form an explanation of a slice K , with poles selection $v_1, \dots, v_{d+1}$, and order $\text{order}(K) = r$, if and only if:

1. $W \cup \{v_1, \dots, v_{d+1}\} \subset U \subset H_K$ and the graph $(U, A \cup F)$ is connected.
2. All elements of W are leaves.
3. For $m = |W|$ and $s = \sum_i^{d+1} L_i(v_i)$ we have $|A| \leq \alpha(m - \mathcal{X}_r) - \beta s$ and $|F| \leq m - 1$.

Furthermore, for a step τ , integer $\Delta' \in [\Delta]$, and a slice K in the analysis of Toom's contour up to step (τ, Δ') , with pole selection $v_1, \dots, v_{d+1}$, we extend the definition of explanation as above, but require the following weaker inequality instead: For $m = |W|$ and $s = \text{Size}(K)$, we have $|A| \leq \alpha(m + 1) - \beta s$ and $|F| < m - 1$.

Claim 5.29. *Let $P_1, \dots, P_{\mu+1}$ be the Pauli-path that induces Toom's contour. For every poles selection $v_1, \dots, v_{d+1}$ of K , there is an explanation. Moreover, any slice in the analysis of Toom's contour up to step (τ, Δ') , and any poles selection for it, there is an explanation.*

Proof. Let $h = T(K) - \min_{u \in H_K} T(u)$. For slices in a standard Toom's contour, we prove by induction on h .

1. Base: $h = 0$, meaning T is constant on H_K , so no arrows connect points of H_K . Thus, H_K , as well as K , consists only of a single leaf, namely K is a result of a fault, and $r = \text{order}(\mathcal{X}_r) \in \{1, 3\}$. In that case, there is a single vertex $u \in V_{\mathcal{H}}$ such that $U = W = \{u\}$, so $m = 1$, $s = 0$, and $|A|, |F| = 0$. It is easy to check that those values satisfy the requirement.

2. Induction Step: $h \neq 0$ implies $K \neq H_K$. By Claim 5.21, we know that for some k there exist slices $K_0, \dots, K_k$, pole selections for them, and forks $F_1, \dots, F_k$ that satisfy Claim 5.21 with respect to the pole selection. In particular, denote by $z_j^1, \dots, z_j^{d+1}$ the pole selection for K_j . Then, for $r \neq 0$ it holds that:

$$\sum_{j=0}^k \sum_{i=1}^{d+1} L_i(z_j^i) + \sum_{i=1}^k \mathbf{Size}(F_i) \geq s$$

Whereas, for $r = 0$:

$$\sum_{j=0}^k \sum_{i=1}^{d+1} L_i(z_j^i) + \sum_{i=1}^k \mathbf{Size}(F_i) \geq s + \xi = s + 1$$

Consider $i \in \{0, \dots, k\}$, and denote $s_j = \sum_i^{d+1} L_i(z_j^i)$. By the inductive hypothesis and since $T(K_i) = T(K) - 1$, we know that for K_i and its selected poles $z_i^1, \dots, z_i^{d+1}$, there is an explanation formed by a set of leaves $\mathbf{W}_i$, a set of arrows $\mathbf{A}_i$, and a set of forks $\mathbf{F}_i$. Let $m_i = |\mathbf{W}_i|$. Thus, the sets $\mathbf{A}_i$ and $\mathbf{F}_i$ have at most $\alpha(m_i - \mathcal{X}_{r-1}) - \beta s_i$ arrows and $m_i - 1$ forks. We define:

- (a) $\mathbf{W} = \cup_{i=0}^k \mathbf{W}_i$, Notice that $\mathbf{W}_i \subset K_i$ and that $T(K_i) = T(K_j)$ for all i, j , thus by Claim 5.17 $\mathbf{W}$ is a union of disjoint set, Hence W contains $m = \sum_{i=0}^k m_i$ leaves.
- (b) $\mathbf{F} = \{F_1, \dots, F_k\} \cup_{i=0}^k \mathbf{F}_i$ with at most $k + \sum_{i=0}^k m_i - 1 = m - 1$ elements.
- (c) $\mathbf{A} = \{\{v_i, \text{pre}_i(v_i)\}\} \cup_{i=0}^k \mathbf{A}_i$. For bounding the number of arrows in $\mathbf{A}$ we condition upon whether $r = 0$, For convenient, let us denote by Y the indicator, which indicate if $r = 0$.

$$\begin{aligned} & d + 1 + \sum_{i=0}^k \alpha(m_i - \mathcal{X}_{r-1}) - \beta s_i \\ & \leq d + 1 + \alpha m - \mathcal{X}_{r-1} \alpha(k + 1) - \beta(s + Y\xi - k \cdot \mathbf{Size}(\text{fork})) \\ & = \alpha(m - \mathcal{X}_r) - \beta s + (\alpha(\mathcal{X}_r - \mathcal{X}_{r-1}) + d + 1 - \mathcal{X}_{r-1} k \alpha - \beta Y \xi + \beta k \cdot \mathbf{Size}(\text{fork})) \end{aligned}$$

So it's left to show that the term in the closure is negative for all integers k :

$$\alpha(\mathcal{X}_r - \mathcal{X}_{r-1}) + d + 1 - \mathcal{X}_{r-1} k \alpha - \beta Y \xi + \beta k \cdot \mathbf{Size}(\text{fork})$$

We do that by showing that for the choice of the parameter values, $\alpha, \beta, \mathcal{X}_0, \dots, \mathcal{X}_3$, the slope is negative, and then the value of the term at k equals zero is negative. For the slope to be negative, we require:

$$-\alpha \mathcal{X}_{r-1} + \beta \mathbf{Size}(\text{fork}) \leq 0 \Rightarrow \mathbf{Size}(\text{fork}) \leq \frac{\alpha}{\beta} \mathcal{X}_{r-1}$$

For negativity at the origin, $k = 0$, we require for $r \neq 0$:

$$\begin{aligned}\alpha(\mathcal{X}_r - \mathcal{X}_{r-1}) + d + 1 &\leq 0 \\ \Rightarrow \mathcal{X}_r &\leq \mathcal{X}_{r-1} - \frac{d+1}{\alpha}\end{aligned}$$

And for $r = 0$:

$$\begin{aligned}\alpha(\mathcal{X}_0 - \mathcal{X}_3) + d + 1 - \beta\xi &\leq 0 \\ \Rightarrow \mathcal{X}_0 &\leq \mathcal{X}_3 - \frac{d+1}{\alpha} + \frac{\beta}{\alpha}\xi\end{aligned}$$

So in total:

$$\begin{aligned}\mathcal{X}_1 &\leq \mathcal{X}_0 - \frac{d+1}{\alpha} \\ \mathcal{X}_2 &\leq \mathcal{X}_0 - 2\frac{d+1}{\alpha} \\ \mathcal{X}_3 &\leq \mathcal{X}_0 - 3\frac{d+1}{\alpha} \\ \mathcal{X}_0 &\leq \mathcal{X}_0 - 4\frac{d+1}{\alpha} + \frac{\beta}{\alpha}\xi \\ \mathbf{Size}(\text{fork}) &\leq \frac{\alpha}{\beta}\mathcal{X}_r\end{aligned}$$

The first three inequalities hold by the definitions of the $\mathcal{X}_r$, the fourth is satisfied since $\frac{\beta}{\alpha}\xi = 5\frac{d+1}{\alpha}$, the last is obtained due to Claim 5.22, that states that the maximal size of a fork is at most $2d$, and that $\frac{\alpha}{\beta}\mathcal{X}_r$ is greater than $4d$.

Finally, by Claim 5.21, the graphs induced by $\mathbf{A}_i \cup \mathbf{F}_i$ are connected to each other through $F_1, \dots, F_k$, thus the graph induced by $\mathbf{A} \cup \mathbf{F}$ is also connected. Therefore, $\mathbf{W}, \mathbf{A}, \mathbf{F}$ form an explanation of K .

It's left to prove for slices in the analysis of Toom's contours. Notice that the 'forming rules' of the analysis of Toom's contour are different from the standard Toom's contour only at the last layer. Therefore, the induction proof captures all the slices not in the last layer. For any slice at the last layer, we repeat the induction step while using Claim 5.27 to bound the expansion in size. So the bounding over $\mathbf{A}$'s size becomes:

$$\begin{aligned}|\mathbf{A}| &\leq \alpha(m+1) - \beta s \\ &\quad + (-\alpha(1 + \mathcal{X}_{r-1}) + d + 1 - \mathcal{X}_{r-1}k\alpha + 2d + \beta k \cdot \mathbf{Size}(\text{fork})) \\ &\leq \alpha(m+1) - \beta s \\ &\quad + (-\alpha + d + 1 - \mathcal{X}_{r-1}k\alpha + 2d + \beta k \cdot \mathbf{Size}(\text{fork}))\end{aligned}$$

Where we used the fact that $\mathcal{X}_{r-1} \in (0, 1)$ for $r \in \{0, 1, 2, 3\}$. Clearly, the bound is satisfied since $\alpha \geq 3d$ and $\mathcal{X}_{r-1}\alpha \geq \beta \cdot \mathbf{Size}(\text{fork})$. $\square$

We summarize by using the linear relation between the leaves of a Toom's contour and the faults that induce it. That gives an upper bound on the probability of a slice occurring when the noise is local stochastic.

Corollary 5.30. *Let K be a slice with size s , and explanation with $|A|$ arrows, and $|F|$ forks. Then it is induced by at least the following number of faults:*

$$\frac{1}{2^{\max\{d', d-d'\}}} \frac{1}{2} (|F| + 1 + \alpha^{-1}(|A| + \beta s - 1))$$

Furthermore, the ratio of leaves to edges in an explanation is greater than $\frac{1}{2\alpha} = \frac{1}{80(d+1)^2}$.

Proof. We bound Claim 5.29 by the number of leaves in the explanation and use the fact that any d' -face has $2^{d'}$ vertices, so each X -fault and Z -fault contributes at most $2^{d'}$ and $2^{d-d'}$ leaves to the explanation. $\square$

5.5 The Replacement Argument

Rewrite. Throughout our final construction, logical measurements will be done via first measuring the qubits of the target register, clean the error with non-noisy computation, and eventually inferring the results, by turning to the logical space. In this section we introduce the cleaning procedure, which is the excitation of sweep-cycles for $2dn^{2d}$ times - the number of vertices times the side's length. In the case of logical measurement, the cleaning is done over classical bits, and therefore only the relevant stage of the sweep-cycle is done (either in the computational or the phase base).

The following claim states that for ideal running (not subject to noise), the cleaning succeeds in erasing the syndrome within at most $\Theta(n^{2d})$ sweep-cycles.

Claim 5.31. *Let $|\psi\rangle$ be a logical codeword of the (d, d') -Toric code with side length n . Let $P \in \mathcal{P}^{(\cdot)}$. Assuming that any two vertices u, w which are two different translations of the same vertex on the finite Toric (a vertex in $\mathcal{G}_n^d$), it holds that they are not connected by arrows. Then at the end the syndrome is erased.*

Proof. We will prove that the decoder erases the ∂^- -boundary corresponding to the X -component of P , and the proof for the Z -component is identical. Let P^∞ the view of P in the ∞ -extension and let $v_1, v_2, \dots$ be the vertices supporting the faces of $\partial^-x(P^\infty)$. During the ideal computation, no new faults are introduced, and therefore $v_1, v_2, \dots$ are the only leaves in the induced Toom's contour. Without loss of generality, assume that the first m vertices $v_1, \dots, v_m$ align with the vertices of the finite Toric (namely non-translated vertices), denote them by V' .

They set a natural relation over its slices as follows: two slices K_1 and K_2 , not necessarily with the same time coordinate, will be said to be *history-related*, $K_1 \sim K_2$, if there are leaves u, w such that both u and w are translations of a vertex $v_i \in V'$ and u, w belong to the histories of K_1, K_2 respectively. Thus, for a time t and a leaf $v \in V'$, we can define the notation $K_{v,t}$ to be the slice at time t that is reached from leaf v by arrows. We emphasize that $\sim$ is not an equivalence relation, since it doesn't satisfy transitivity. Yet, let us match slices to 'classes' when the map is many-to-many.

Here, a class is a grouping of translations of slices which are all history-related. For a vertex $v \in V'$, the delegate of the slices reached from vertex v and its translations, namely the translations of slices $\{K_{v,t} : t\}$, will be denoted by $\mathbf{K}_v$. Proving that at the end of time $2dn^{2d}$, for any leaf v , the slice $K_{v,t}$ is empty, or equivalently is of size zero, and therefore its translations, implies the correctness of the claim.

For proving this, let us present the following terminology. For leaves v, u, w , and a time coordinate t , consider the following:

1. Suppose that $K_{v,t}$ and $K_{u,t+1}$ are slices at times t and $t+1$, such $K_{v,t} \sim K_{u,t+1}$. If $\mathbf{Size}(K_{u,t+1}) > \mathbf{Size}(K_{v,t})$, we will say that the class $\mathbf{K}_v$ grows at time t . Similarly, if $\mathbf{Size}(K_{u,t+1}) < \mathbf{Size}(K_{v,t})$, then we will say that class $\mathbf{K}_v$ shrank at time t .
2. Suppose that $K_{v,t}$ and $K_{u,t}$ are slices at time t , and $K_{w,t+1}$ is a slice at time $t+1$, such that $K_{v,t} \not\sim K_{u,t}$, $K_{w,t+1} \sim K_{v,t}$, and $K_{w,t+1} \sim K_{u,t}$. Then we will say that $\mathbf{K}_v$ and $\mathbf{K}_u$ merged at time t .
3. The size of a class $\mathbf{K}_v$ at time t is $\mathbf{Size}(K_{v,t})$.

Since no new faults are introduced during the computation, we have that in each step, the only option for a class to grow is by merging. Notice that, by definition, classes cannot 'split'. Namely, if $K_{v,t} \sim \mathbf{K}_u$, then for any time $t' \geq t$, it holds that $K_{v,t'} \sim \mathbf{K}_u$. Specifically, a correction by the sweep rule that disconnects two connected components of the syndrome that were connected, prior to correction, belonged to the same slice and therefore to the same class, would still belong to the same class after the correction.

Therefore, there can be at most $m - 1$ merges. And since the number of classes is bounded by the number of vertices on the Toric, we have that $m \leq n^d$, namely at most $n^d - 1$ merges can occur. Hence, during the $2dn^{2d}$ sweep-cycles, there has to be a continuous sequence of $2dn^d$ cycles for which no merge occurs.

Moreover, we argue that the maximum size of a slice is bounded by $2dn^d$. For this, for a time t we define $V'_t \subset V'$, to be the maximal subset such for any $v, u \in V'_t$ it holds $K_{v,t} \not\sim K_{u,t}$. Then, we define the following potential function:

$$\Phi_t = \sum_{v \in V'_t} (\mathbf{Size}(K_{v,t}) + 1)$$

Observe that $\Phi_1 = m \leq n^d$. Now, for a vertex $v \in V'_{t+1}$, denote the predecessor of $K_{v,t+1}$ by $K_{u'_0,t}, K_{u'_1,t}, \dots, K_{u'_{k_v},t}$ and their translation into the finite Toric by $K_{u_0,t}, K_{u_1,t}, \dots, K_{u_{k_v},t}$, so $u_i \in V'_t$. Using Corollary 5.26 and Claim 5.22, where the

poles for $K_{v,t+1}$ are selected to be the extremis of $L_1, \dots, L_{d+1}$, we have:

$$\begin{aligned}
\Phi_{t+1} &= \sum_{v \in V'_{t+1}} (\text{Size}(K_{v,t+1}) + 1) \leq \sum_{v \in V'_{t+1}} \left(\left(\sum_{i=0}^{k_v} \text{Size}(K_{u'_i,t}) \right) + 2dk_v + 1 \right) \\
&= \sum_{v \in V'_{t+1}} \left(\left(\sum_{i=0}^{k_v} \text{Size}(K_{u_i,t}) \right) + (2d-1)k_v + 1 + k_v \right) \\
&= \sum_{v \in V'_{t+1}} \left(\left(\sum_{i=0}^{k_v} \text{Size}(K_{u_i,t}) + 1 \right) + (2d-1)k_v \right)
\end{aligned}$$

Where, in the last transition, we used the fact that the size function is invariant to translations. Now, any vertex of V'_t has to appear at most once. To see this, assume towards contradiction the existence of two predecessor slices $K_{u'_i,t}$, $K_{w'_j,t}$ which are translations of the same slice. Denote by $K_{u,t+1}$ and $K_{w,t+1}$ the slices whose $K_{u'_i,t}$, $K_{w'_j,t}$ are their predecessors. Let's split to cases, whether $u = w$.

- If $u \neq w$, then u'_i, w'_j belong to the histories of different slices $K_{u,t+1}, K_{w,t+1}$, and they are a translation of the same vertex in $V'_t \subset V'$. Hence, $K_{u,t+1} \sim K_{w,t+1}$, in contradiction to $u, w \in V'_{t+1}$.
- Otherwise, if $u = w$, then we found two vertices u'_i and w'_j , which are translations of the same origin vertex of the finite Toric, such that they belong to the same history $H_{K_{w,t+1}}$, namely connected by an arrow, in contradiction to the given of.

Now, observe that each fork connecting two slices is associated with a merge; therefore, $\sum_t \sum_{v \in V'_t} k_v$ is lower than the number of merges, and the bound on the potential function becomes:

$$\begin{aligned}
\Phi_{t+1} &\leq \sum_{v \in V'_t} \text{Size}(K_{v,t}) + 1 + \sum_{v \in V'_{t+1}} (2d-1)k_v \\
&\leq \Phi_t + (2d-1)k = \Phi_1 + (2d-1) \cdot \# \text{ merges} \leq 2dn^d
\end{aligned}$$

Thus, since $\text{Size}(K_{v,t}) \leq \Phi_t$, we have that during the computation, no slice reaches size $2dn^d$. During the continuous sequence of $2dn^d$ cycles, each sweep cycle decreases the size of any slice by at least 1. This is followed by Corollary 5.26, combining that no merging occurred, namely there is only one predecessor slice. Hence, at the end of that sequence of $2dn^d$ cycles, all the slices' sizes shrink to zero; namely, the connected components are erased after those cycles. $\square$

To prove that no logical error is admitted (namely, through the computation, no connected component of the boundary is sufficiently large to pass the barrier), and that the assumption that no two translated vertices are arrow-connected holds, we use a replacement argument, stating that the probability of an event occurring during a computation where part of its layers are non-noisy is lower than the probability of having the same event in a computation where all its layers are noisy. More precisely,

any bound that assumes only that the noise channel is local stochastic captures both cases.

Claim 5.32 (Replacement argument). *Consider a running of Toric encoded circuit with depth l . Denote by A the event in which, after an additional $2dn^{2d}$ **ideal** (non-noisy) iterations of the sweep-cycle, the error is not erased, and denote by B the event that, during the application of an additional $2dn^{2d}$ **noisy** iterations, either the Toom's contour admits an infinite error, or two translated vertices are become arrows connected. Let $f : (n, d, d', l, p) \rightarrow \mathbb{R}$ be an upper bound on $\Pr[B]$ that assumes only that the noise is distributed according to a local stochastic noise channel with noise rate p . Then, we have:*

$$\Pr[A] \leq f(n, d, d', l, p)$$

Proof. A corollary of Claim 5.31 is that if during the $2dn^{2d}$ ideal sweep-cycles, there is no Pauli-path up to step that admits a sufficiently large syndrome connected component, and no two translated vertices become arrow-connected, then the error is erased. Consider the set of atomic events corresponding to the faults in the first l noisy layers, such that during the additional $2dn^{2d}$ ideal iterations, there is a step in which a connected component of the syndrome reaches the barrier, or that its corresponding Pauli-path up to step induces an 'analysis Toom's counter,' in which two translated vertices are connected by arrows. All those atomic events are also contained in B , since local stochastic noise is our only assumption. The bound over the B scenario has to hold for a noise channel that assigns zero probability for a fault occurring during the additional $2dn^{2d}$ iterations. For this specific channel, $\Pr[A] \leq \Pr[B]$ and therefore $\Pr[A] \leq f(n, d, d', l, p)$. $\square$

6 Proof of the Theorem

We are ready to prove the main theorem (Theorem 6.3). We start by giving a formal definition of a simulation.

Definition 6.1 (Simulation and Logical Errors). *Consider a quantum code $\mathcal{C}$ of length n and dimension k . A quantum circuit $D: L(\mathcal{H}_2^{\otimes n}) \rightarrow L(\mathcal{H}_2^{\otimes n})$ will be said to be a referee for $\mathcal{C}$ if, when given a codeword of $\mathcal{C}$, it results in the first k output qubits, the logical state of the given. We say that a Pauli operator $P \in \mathcal{P}^{\otimes n}$ is a logical error with respect to D if $DP \neq D$.*

For quantum circuits C and $\tilde{C}$, a noise channel $N_{\mathcal{E}}$, and $p \in (0, 1)$, we say that $\tilde{C}$ simulates C with probability at least $1 - p$ if

1. *There is a mapping between the qubits of C and the registers of $\tilde{C}$, along with a referee such that, at the end of running $\tilde{C}$ (on the zero state), decoding the registers back to the logical space results in the same quantum states as those produced by running C .*

2. Let $\mathcal{E} \sim N_{\mathcal{E}}$ be a set of faults. The probability induced by $N_{\mathcal{E}}$ that at the end of the running of $\tilde{C}[\mathcal{E}]$, one of its quantum registers introduces a logical error (with respect to the referee) is less than p .

The next lemma is the core of the theorem; it analyzes the probability of a logical error occurring after the initialization stage, namely when assuming the given input is encoded.

Lemma 6.2. *For any d, d' , there are constants $a, b \in \mathbb{R}^+$ and $p_0 \in (0, 1)$ that depend only on d, d' and satisfy the following. Consider a quantum circuit C with depth μ and width w , and let $N_{\mathcal{E}}$ be a local stochastic noise channel with a noise rate $p < p_0$. There is a Toric encoded circuit $\tilde{C}$ such that each of its quantum registers is encoded by an n -side length, (d, d') -Toric code, and it simulates C with probability:*

$$q = 1 - O(\text{poly}(w, \mu, n)) (ap^b)^{\frac{n}{2d}}$$

To prove the lemma, we will need the following two claims. The first bounds, for a given regular graph, the number of its subgraphs having at most k edges. This claim will be used to infer the number of possible explanation trees for a given slice. The second states that each tree has a subtree with fewer than half of the edges and a leaves-to-edge ratio that is at least half of the ratio of the original tree. This claim will be used to infer that any explanation has a subtree where, under the projection to the finite Toric, its faults are mapped uniquely, meaning no fault is counted twice.

Claim (Claim 7.6). *Let $G = (V, E)$ be a $\Delta + 1$ -regular graph, and let k be an integer less than $|E|$. Fix a vertex $r \in V$. There is a constant number σ , depending only on Δ , such that the number of subgraphs rooted at r with fewer than k edges is at most σ^k .*

Claim (Claim 7.9). *Let $\gamma, \beta \in \mathbb{R}_{>0}$ and n be a positive integer. For any tree T with $\tau \geq \beta n$ edges and m leaves such that $\tau \leq \gamma m$, there is a subtree T' of T with a number of edges τ' less than $\frac{\beta}{2}n$ and a number of leaves greater than $\frac{\beta}{4\gamma}n$.*

The proofs of the claims are given in the appendix. We continue with the proof of the lemma.

Proof. We construct $\tilde{C}$ as follows: for each input wire of C , we initialize the logical zero $|0\rangle$ of the (d, d') -Toric code. For each of the gates **H** and **S**, we initialize the encoded $|\mathbf{H}\rangle$ and $|\mathbf{S}\rangle$ states, and for each **T**, we initialize copies of both $|\mathbf{T}\rangle$ and $|\mathbf{S}\rangle$ states. The even layers of $\tilde{C}$ are sweep cycles, while on the odd layers, we set the realization of the logical gates. For **CX**, we set the transversal realization, and for **H**, **S**, and **T**, we use gate teleportation. Thus the width of $\tilde{C}$ expands to at most $(w + 2\mu)n^{2d}$.

We realize the gate teleportation as described in Definition 4.39, so the gate teleportation admits $O(n + \log n)$ and $O(n^{2d})$ depth and width overhead. Thus, the depth of $\tilde{C}$ expands to $\Theta(\mu n)^{15}$. We will denote the procedure of inferring the logical measurement in the gate teleportation Definition 4.39¹⁶ by D , and we also

¹⁵The constant inside the Θ also accounts for the fact that the realization of a logical gate might contain multiple gate teleportations.

¹⁶The codeword checking procedure - $\mathbf{1}_{x_{S'}} / \mathbf{1}_{x_{S/S'}}$.

use D as a realization of measurement in the computational basis and as our referee (Definition 6.1). Denote by S the number of locations, so $S = O(w\mu^2 n^{2d+1})$.

For the sake of proving correctness, we can assume that at the end we measure all registers, each in either the computational or the phase basis. Any proof that gives correctness without distinguishing in which basis we choose for the registers captures the referee test at the end of the run. Then, saying that the referee, D , does not fail means that the logical measurement indeed measures their given encoded states. In turn, the corrections followed in the gate teleportation procedures are the right ones, and therefore, the logical gates that were applied throughout the simulation match the gates in the original circuit. Lastly, we pass the referee test. In other words, the original circuit was successfully simulated according to Definition 6.1.

Here. Thus, showing that D does not fail, and that at the end of the running, the resulting state is Corollary 5.5

$$\tilde{C}[\mathcal{E}]_{\rightarrow \tau_{\mathcal{E}}} = \tilde{C}_{\rightarrow \tau} \sum_i p_i P_i$$

Claim 5.31 states that D successfully decodes, assuming that, first, no infinite error occurs and, second, no vertex is connected to a non-trivial translation of itself by arrows. And Claim 5.32 states that any upper bound on the occurrence probability of at least one of either infinite error or a vertex connecting to its translation, in the running where the decoding procedure is noisy¹⁷, and which assumes only that the noise channel is local stochastic, is sufficient to bound the probability that one of them occurs during the original running.

Since connectivity by faces in the analysis of Toom's contour implies connectivity by arrows, these two conditions are followed by requiring the slightly stronger restriction that no time t is the first time such that the analysis of Toom's contour at time t , namely $\mathcal{H}_t$, has an infinite slice.

Thus, it remains to prove that with probability at least q , no time t is the first time such $\mathcal{H}_t$ has an infinite slice of in the setting in which the decoding procedures are also subjected to noise. From now on, we consider that setting.

Denote by B_t the event where t is the first time such that $\mathcal{H}_t$ has an infinite slice, and by A_t the event where all explanations of slices (and any of their pole selections) in $\mathcal{H}_t$ have fewer than $\frac{n}{2d}$ edges. Using the conditional probability formula, we have:

$$\Pr_{\sim \mathcal{N}_{\mathcal{E}}} [B_t] = \Pr_{\sim \mathcal{N}_{\mathcal{E}}} [B_t \cap A_t] + \Pr_{\sim \mathcal{N}_{\mathcal{E}}} [B_t \cap \bar{A}_t]$$

We begin by bounding the contribution of the right term. Intersection with B_t guarantees that any slice, and a pole selection for it, have an explanation. According to Corollary 5.30, there exists a constant c_1 , depending only on d , such that the ratio of leaves to edges in the explanation is at least c_1 . By Claim 7.9, it follows that any such explanation has more than $\frac{n}{2d}$ edges containing a subtree with fewer than $\frac{n}{4d}$ edges, with a leaves-to-edges ratio greater than $\frac{1}{2}c_1$. Using the union bound, we bound

¹⁷Observe that in this setting, the classical procedure that infers a logical measurement makes no sense

the probability that there is a vertex in the base graph that is a root of such a subtree. Since the construction of the ∞ -extension sets periodic conditions, it's sufficient to ask if there is such a vertex in the restriction to the finite n -side length Toric.

$$\Pr_{\sim \mathcal{N}_\varepsilon} [\bar{A}_\tau \cap B_\tau] \leq \bigcup_{\substack{v \in V_o, \tau' \leq \tau \\ L_g(v) \in [0, n]}} \Pr_{\sim \mathcal{N}_\varepsilon} [v \text{ is a root of subtree in } \mathcal{H}_{\tau'} \\ \text{with less than } \frac{n}{4d} \text{ edges and more than } \frac{1}{2} c_1 \frac{n}{4d} \text{ leaves}]$$

Now, since any two connected vertices are at a space distance of at most

$$\max\{\mathbf{Size}(\text{arrow}), \mathbf{Size}(\text{fork})\} \leq 2d$$

It follows that all the vertices in the subtree are at a space distance of at most $\frac{1}{2}n$ from each other. Thus, no two of them have the same space coordinate modulo n . In particular, projecting the infinite plane back to the finite Toric sends each of the leaves of the subtree to a unique target, namely no fault is counted twice. Therefore the number of faults that induce the explanation is at least:

$$\frac{1}{2^{\max\{d', d-d'\}}} \frac{1}{2} \frac{1}{2} c_1 \frac{n}{4d} \geq \frac{1}{2^{d+4}} c_1 \frac{n}{d}$$

Plugging into the union bounds and combining Claim 7.6 yields:

$$\Pr_{\sim \mathcal{N}_\varepsilon} [\bar{A}_t \cap B_t] \leq tO(S) \sigma^{\frac{n}{4d}} p^{\frac{c_1 n}{2^{d+4}d}} \leq O(S^2) \left(\sigma p^{\frac{1}{2^{d+2}} c_1} \right)^{\frac{n}{4d}}$$

It's left to bound the probability of $B_t \cap A_t$. Since prior to, and including time t , no time admits an infinite error, it follows that any slice in $\mathcal{H}_t$, and any pole selection of it, has an explanation (Claim 5.29). Confining to the event A_t , any explanation implies that any explanation has fewer than $\frac{n}{2d}$ edges and therefore, by the same argument as above, the projection of each explanation tree into the finite Toric sends its leaves to unique targets.

Any infinite slice, namely a slice of size at least n , has poles selection $v_1 \dots, v_{d+1}$ such:

$$\sum_{i=1}^{d+1} L_i(v_i) \geq n \tag{3}$$

Consider a slice with size at least n and a pole selection for it satisfying Equation (3). Claim 5.29 and Corollary 5.30 state that there are constants $c_2, c_3 > 0$ such that any explanation for the slice and the selection, with x edges, has at least $c_2 n + c_3 x$ leaves. Therefore, we obtain:

$$\Pr_{\sim \mathcal{N}_\varepsilon} [B_t \cap A_t] \leq \bigcup_{\substack{v \in V_o, x \geq 1 \\ L_g(v) \in [0, n]}} \Pr_{\sim \mathcal{N}_\varepsilon} [v \text{ is a root of explanation with } x \text{ edges}$$

and at least $c_2 n + c_3 x$ leaves,
which are projected to different targets]

So, for sufficiently small p such that $\sigma p^{c_2} < 1$, the probability is bounded by:

$$\sum_{x=1}^{\infty} O(S) p^{c_2 n} (\sigma p^{c_3})^x \leq O(S) p^{c_2 n}$$

Combing the bounds we have, that there is constant c_4 depending only on d, d' such:

$$\Pr_{\sim \mathcal{N}_{\mathcal{E}}} [B_{\tau}] \leq O(S^2) p^{c_4 n}$$

Now, by the union bound over the time steps, with a probability greater than:

$$1 - O(S^3) p^{c_4 n}$$

No time admits an infinite slice at its corresponding analysis Toom's contour, so the decoder successfully restores the logical state; in other words, no logical error occurs. $\square$

The lemma gives the theorem almost immediately. To complete the proof, we need to plug in the initialization phase, which initializes encoded zeros and resource states subjected to local stochastic noise. We emphasize that the resulting states are subjected to local stochastic noise (in addition to the computation). We neither prove nor describe that phase here; instead, we refer to [Gro19], [Got14], [FGL18], and [NP25].

Theorem 6.3. *There exists a constant $p_0 \in (0, 1)$ such that for any local stochastic noise channel $\mathcal{E}$ with rate $p < p_0$, and a quantum circuit C with depth μ and width w , there exists a Toric encoded circuit $\tilde{C}$, with depth $\Theta(\mu n^8 + \text{Polylog}(n))$ and width $\Theta((w + \mu)n^8 \text{Polylog}(n))$ that simulates C with probability:*

$$1 - O(\text{poly}(w, \mu, n)) p^{\Omega(n)}$$

and in addition:

1. *After a short initialization stage (at depth $\Theta(\text{poly}(\log n))$), each of the qubits is encoded in the n -length $4D$ Toric code.*
2. *An error correction cycle takes $O(1)$ depth, and the gates in each correction cycle don't depend on time. In particular, if C is the identity circuit, then after the initialization stage, the layers in $\tilde{C}$ are the same.*

References

- [Neu56] J. von Neumann. “Probabilistic Logics and the Synthesis of Reliable Organisms From Unreliable Components”. In: *Automata Studies*. Ed. by C. E. Shannon and J. McCarthy. Princeton: Princeton University Press, 1956, pp. 43–98. ISBN: 9781400882618. DOI: [doi:10.1515/9781400882618-003](https://doi.org/10.1515/9781400882618-003). URL: <https://doi.org/10.1515/9781400882618-003>.

- [Lev80] G. Lev. “Size Bounds and Parallel Algorithms for Networks”. PhD thesis. University of Edinburgh, 1980. URL: <https://era.ed.ac.uk/handle/1842/17015>.
- [SV82] Yossi Shiloach and Uzi Vishkin. “An $O(\log n)$ parallel connectivity algorithm”. In: *Journal of Algorithms* 3.1 (1982), pp. 57–67. ISSN: 0196-6774. DOI: [https://doi.org/10.1016/0196-6774\(82\)90008-6](https://doi.org/10.1016/0196-6774(82)90008-6). URL: <https://www.sciencedirect.com/science/article/pii/0196677482900086>.
- [AS87] Awerbuch and Shiloach. “New Connectivity and MSF Algorithms for Shuffle-Exchange Network and PRAM”. In: *IEEE Transactions on Computers* C-36.10 (1987), pp. 1258–1263. DOI: [10.1109/TC.1987.1676869](https://doi.org/10.1109/TC.1987.1676869).
- [BS88] Piotr Berman and J’anos Simon. “Investigations of fault-tolerant networks of computers”. In: *Proceedings of the Twentieth Annual ACM Symposium on Theory of Computing*. STOC ’88. Chicago, Illinois, USA: Association for Computing Machinery, 1988, pp. 66–77. ISBN: 0897912640. DOI: [10.1145/62212.62219](https://doi.org/10.1145/62212.62219). URL: <https://doi.org/10.1145/62212.62219>.
- [CS96] A. R. Calderbank and Peter W. Shor. “Good quantum error-correcting codes exist”. In: *Physical Review A* 54.2 (Aug. 1996), pp. 1098–1105. DOI: [10.1103/physreva.54.1098](https://doi.org/10.1103/physreva.54.1098). URL: <https://doi.org/10.1103/PhysRevA.54.1098>.
- [Ste96] Andrew Steane. In: *Proceedings of the Royal Society of London. Series A: Mathematical, Physical and Engineering Sciences* 452.1954 (Nov. 1996), pp. 2551–2577. ISSN: 1471-2946. DOI: [10.1098/rspa.1996.0136](https://doi.org/10.1098/rspa.1996.0136). URL: <http://dx.doi.org/10.1098/rspa.1996.0136>.
- [Kit97] A Yu Kitaev. “Quantum computations: algorithms and error correction”. In: *Russian Mathematical Surveys* 52 (1997), pp. 1191–1249. URL: <https://api.semanticscholar.org/CorpusID:250816585>.
- [AB99] Dorit Aharonov and Michael Ben-Or. *Fault-Tolerant Quantum Computation With Constant Error Rate*. 1999. arXiv: [quant-ph/9906129](https://arxiv.org/abs/quant-ph/9906129) [quant-ph].
- [ZLC00] Xinlan Zhou, Debbie W. Leung, and Isaac L. Chuang. “Methodology for quantum logic gate construction”. In: *Physical Review A* 62.5 (Oct. 2000). ISSN: 1094-1622. DOI: [10.1103/physreva.62.052316](https://doi.org/10.1103/physreva.62.052316). URL: <http://dx.doi.org/10.1103/PhysRevA.62.052316>.
- [Den+02] Eric Dennis et al. “Topological quantum memory”. In: *Journal of Mathematical Physics* 43.9 (Sept. 2002), pp. 4452–4505. DOI: [10.1063/1.1499754](https://doi.org/10.1063/1.1499754). URL: <https://doi.org/10.1063/1.1499754>.
- [Ali+08] R. Alicki et al. *On thermal stability of topological qubit in Kitaev’s 4D model*. 2008. arXiv: [0811.0033](https://arxiv.org/abs/0811.0033) [quant-ph]. URL: <https://arxiv.org/abs/0811.0033>.
- [Got14] Daniel Gottesman. *Fault-Tolerant Quantum Computation with Constant Overhead*. 2014. arXiv: [1310.2984](https://arxiv.org/abs/1310.2984) [quant-ph].

- [Bom15] Héctor Bombín. “Gauge color codes: optimal transversal gates and gauge fixing in topological stabilizer codes”. In: *New Journal of Physics* 17.8 (Aug. 2015), p. 083002. DOI: [10.1088/1367-2630/17/8/083002](https://doi.org/10.1088/1367-2630/17/8/083002). URL: <https://doi.org/10.1088/1367-2630/17/8/083002>.
- [FGL18] Omar Fawzi, Antoine Grospellier, and Anthony Leverrier. “Constant Overhead Quantum Fault-Tolerance with Quantum Expander Codes”. In: *2018 IEEE 59th Annual Symposium on Foundations of Computer Science (FOCS)*. IEEE, Oct. 2018, pp. 743–754. DOI: [10.1109/focs.2018.00076](https://doi.org/10.1109/focs.2018.00076). URL: <http://dx.doi.org/10.1109/FOCS.2018.00076>.
- [Gro19] Antoine Grospellier. “Constant time decoding of quantum expander codes and application to fault-tolerant quantum computation”. Theses. Sorbonne Université, Nov. 2019. URL: <https://theses.hal.science/tel-03364419>.
- [KP19] Aleksander Kubica and John Preskill. “Cellular-Automaton Decoders with Provable Thresholds for Topological Codes”. In: *Physical Review Letters* 123.2 (July 2019). ISSN: 1079-7114. DOI: [10.1103/physrevlett.123.020501](https://doi.org/10.1103/physrevlett.123.020501). URL: <http://dx.doi.org/10.1103/PhysRevLett.123.020501>.
- [Gac21] Peter Gacs. *A new version of Toom’s proof*. 2021. arXiv: [2105.05968](https://arxiv.org/abs/2105.05968) [cs.FL]. URL: <https://arxiv.org/abs/2105.05968>.
- [Ber+24] Noah Berthussen et al. “Experiments with the four-dimensional surface code on a quantum charge-coupled device quantum computer”. In: *Physical Review A* 110.6 (Dec. 2024). ISSN: 2469-9934. DOI: [10.1103/physreva.110.062413](https://doi.org/10.1103/physreva.110.062413). URL: <http://dx.doi.org/10.1103/PhysRevA.110.062413>.
- [BB24] Nikolas P. Breuckmann and Simon Burton. “Fold-Transversal Clifford Gates for Quantum Codes”. In: *Quantum* 8 (2024), p. 1372. ISSN: 2521-327X. DOI: [10.22331/q-2024-06-13-1372](https://doi.org/10.22331/q-2024-06-13-1372). URL: <http://dx.doi.org/10.22331/q-2024-06-13-1372>.
- [GV24] Alexandre Guernut and Christophe Vuillot. *Fault-Tolerant Constant-Depth Clifford Gates on Toric Codes*. 2024. arXiv: [2411.18287](https://arxiv.org/abs/2411.18287) [quant-ph]. URL: <https://arxiv.org/abs/2411.18287>.
- [Rei+24] Ben W. Reichardt et al. *Demonstration of quantum computation and error correction with a tesseract code*. 2024. arXiv: [2409.04628](https://arxiv.org/abs/2409.04628) [quant-ph]. URL: <https://arxiv.org/abs/2409.04628>.
- [Aas+25] David Aasen et al. *A Topologically Fault-Tolerant Quantum Computer with Four Dimensional Geometric Codes*. 2025. arXiv: [2506.15130](https://arxiv.org/abs/2506.15130) [quant-ph]. URL: <https://arxiv.org/abs/2506.15130>.
- [NP25] Quynh T. Nguyen and Christopher A. Pattison. *Quantum fault tolerance with constant-space and logarithmic-time overheads*. 2025. arXiv: [2411.03632](https://arxiv.org/abs/2411.03632) [quant-ph]. URL: <https://arxiv.org/abs/2411.03632>.

- [BDL26] Shankar Balasubramanian, Margarita Davydova, and Ethan Lake. “A local automaton for the 2D toric code”. In: *Quantum* 10 (2026), p. 2125. ISSN: 2521-327X. DOI: [10.22331/q-2026-06-02-2125](https://doi.org/10.22331/q-2026-06-02-2125). URL: <http://dx.doi.org/10.22331/q-2026-06-02-2125>.
- [DST26] Gesa Dünnweber, Georgios Styliaris, and Rahul Trivedi. *Quantum Memory and Autonomous Computation in Two Dimensions*. 2026. arXiv: [2601.20818](https://arxiv.org/abs/2601.20818) [quant-ph]. URL: <https://arxiv.org/abs/2601.20818>.
- [SST26] Jan M. Swart, Réka Szabó, and Cristina Toninelli. *Peierls bounds from Toom contours*. 2026. arXiv: [2202.10999](https://arxiv.org/abs/2202.10999) [math.PR]. URL: <https://arxiv.org/abs/2202.10999>.

7 Appendix

7.1 Time Coarse-Graining

In this section, we prove that when subjecting a circuit to local stochastic noise, changing the frequency at which we introduce the fault layers translates to rescaling the noise rate. For this purpose, we are abusing the notation “propagation” by reusing it to describe a different context.

Definition 7.1 (Layer propagation). *Let $C = \{u_i\}_{i=1}$ be a quantum circuit. For each time coordinate t , write $U_t := C|_t$. The family of non-empty t -cuts $\{U_t\}_t$ is called the layer decomposition of C .*

Let $t < t'$ be such that U_t and $U_{t'}$ are non-empty and $U_s = \emptyset$ for every $t < s < t'$. The forward propagation of the layer U_t across the next layer $U_{t'}$ is the operation that replaces the pair $(U_t, U_{t'})$ by a pair $(U_{t'}, V_{t \rightarrow t'})$ satisfying

$$\prod_{u \in U_t} u \prod_{v \in U_{t'}} v = \prod_{v \in U_{t'}} v \prod_{w \in V_{t \rightarrow t'}} w.$$

The layer $V_{t \rightarrow t'}$ is called the propagation expression of U_t across $U_{t'}$.

More generally, if $t < t'$, the forward propagation of U_t to time t' is obtained by iterating the above operation through the successive non-empty cuts between times t and t' . The resulting propagated layer is denoted by $V_{t \rightarrow t'}$ and is again called the propagation expression of U_t to time t' . The resulting circuit has layer decomposition

$$\dots, U_{t-1}, U_t, U_{t+1}, \dots, U_{t'}, U_{t'+1}, \dots \mapsto \dots, U_{t-1}, U_{t+1}, \dots, U_{t'}, V_{t \rightarrow t'}, U_{t'+1}, \dots$$

After propagation we update the time coordinates so that they agree with the new layer in which each gate is applied.

Definition 7.2 ($C \sim_P C'$). *For quantum circuits C and C' , we write $C \sim_P C'$ if one can be obtained from the other by a finite sequence of layer propagations.*

By construction, $\sim_P$ preserves the implemented channel. Indeed, if $C \sim_P C'$ and ρ is any mixed state on the input register, then

$$C \circ \rho = C' \circ \rho.$$

We state Claim 3.8 in the language of layer propagation.

Claim 7.3. *Let $\mathcal{N}_{\mathcal{E}}$ be a local stochastic noise channel with rate p , and let C be a quantum circuit whose gates have arity at most Δ . Fix an even integer c . Then there exists a local stochastic noise channel $\mathcal{N}_{\mathcal{E}_q}$ with rate*

$$q := \left(2^{H(\Delta^{-2c})\Delta^{2c}} p \right)^{\Delta^{-c}},$$

such that $C[\mathcal{E}]_1 \sim_P C[\mathcal{E}_q]_{c/2}$. In particular, up to a constant depending only on Δ and c , the rate q behaves like $p^{\Delta^{-c}}$.

Proof of Claim 3.8. Partition the subjected circuit $C[\mathcal{E}]_1$ into consecutive blocks of c layers. Inside each block, propagate all the fault layers to the end of the block and then grouping the resulted expressions together. Doing this block by block yields a new circuit C' in which each block of $c/2$ computational layers is followed by a single fault layer. Therefore $C' = C[\mathcal{E}']_{c/2}$ for a suitable random fault set $\mathcal{E}'$, and by construction: $C[\mathcal{E}]_1 \sim_P C[\mathcal{E}']_{c/2}$.

It remains to be proven that $\mathcal{E}'$ is distributed according to a local stochastic channel. Fix one propagation expression $V^{(i)}$ and a subset $S \subset \text{Loc}(V^{(i)})$. Denote by $U_1^i, \dots, U_c^i$ the layers for which $V^{(i)}$ is their (grouped) expression. Denote the light cone of S in $U_1^i, \dots, U_c^i$ by $\Lambda(S)$. Since the fan is bounded by Δ , we have that $|\Lambda(S)| \leq \Delta^c |S|$, and any $S' \subset \Lambda(S)$ whose propagation covers S has size at least $|S'| \geq \Delta^{-c} |S|$, and in particular contains a subset of size exactly $\Delta^{-c} |S|$.

Now, having $S \subset \text{Loc}(V^{(i)})$, implies that there is a minimal S' (the minimality is regarding inclusion) such that $S' \subset \text{Loc}(\mathcal{E})$. Denote by $\mathcal{W}(S)$ the set of such minimal S' 's, and observe that its size is at most:

$$|\mathcal{W}(S)| \leq \binom{\Delta^c |S|}{\Delta^{-c} |S|}$$

Hence, using the fact that $\mathcal{E}$ is an outcome of local stochastic noise channel with a noise rate p we have:

$$\begin{aligned} \Pr[S \subset \text{Loc}(V^{(i)})] &\leq \Pr\left[\bigcup_{S' \in \mathcal{W}(S)} \{S' \subset \text{Loc}(\mathcal{E})\}\right] \\ &\leq \sum_{S' \in \mathcal{W}(S)} \Pr[S' \subset \text{Loc}(\mathcal{E})] \\ &\leq \binom{\Delta^c |S|}{\Delta^{-c} |S|} p^{\Delta^{-c} |S|} \leq \left(2^{H(\Delta^{-2c})\Delta^{2c}} p \right)^{|S|/\Delta^c} = q^{|S|}. \end{aligned}$$

□

7.2 Checks Commutation

We give an alternative proof of Claim 4.4 by straightforwardly multiplying checks. Consider subsets of generators $A, B \subset S$ with sizes $d' - 1$ and $d' + 1$, respectively. Thus, $\partial^+ \theta_{\mathbf{v}, \mathbf{A}}$ and $\partial^- \theta_{\mathbf{u}, \mathbf{B}}$ are X - and Z -checks. Let us examine their multiplication:

$$(\partial^+ \theta_{\mathbf{v}, \mathbf{A}})(\partial^- \theta_{\mathbf{u}, \mathbf{B}}) = \left(\sum_{g \in S/A} \theta_{\mathbf{v}, \mathbf{A} \cup \mathbf{g}} + \theta_{\mathbf{g}^{-1} \mathbf{v}, \mathbf{A} \cup \mathbf{g}} \right) \left(\sum_{g \in B} \theta_{\mathbf{u}, \mathbf{B}/\mathbf{g}} + \theta_{\mathbf{g} \mathbf{u}, \mathbf{B}/\mathbf{g}} \right)$$

Split into cases based on whether v equals u .

1. If $v = u$, then for any $g, g' \in S$, it holds that $gv \neq u$, $g'^{-1}u$, and $v \neq g'^{-1}u$. Thus, only terms of the form

$$\theta_{\mathbf{v}, \mathbf{A} \cup \mathbf{g}} \theta_{\mathbf{u}, \mathbf{B}/\mathbf{g}'}$$

can contribute to the product, yet for any g, g' for which this term is not zeroing out, it follows that $g \in B$ and $g' \notin A$, and therefore the term

$$\theta_{\mathbf{v}, \mathbf{A} \cup \mathbf{g}'} \theta_{\mathbf{u}, \mathbf{B}/\mathbf{g}}$$

Is not zero as well. Therefore, the number of non-zero terms is even, namely the sum equals zero.

2. If there is g such that $gv = u$, then first notice that for any $g', g'' \in S/g \cup \{e\}$, we have $g'v \neq u$. So only terms of the form

$$\text{either } \theta_{\mathbf{g} \mathbf{v}, \mathbf{A} \cup \mathbf{g}} \theta_{\mathbf{u}, \mathbf{B}/\mathbf{g}'} \text{ or } \theta_{\mathbf{v}, \mathbf{A} \cup \mathbf{g}'} \theta_{\mathbf{g}^{-1} \mathbf{u}, \mathbf{B}/\mathbf{g}}$$

can contribute. Similarly to the above, if

$$\theta_{\mathbf{g} \mathbf{v}, \mathbf{A} \cup \mathbf{g}} \theta_{\mathbf{u}, \mathbf{B}/\mathbf{g}'} \neq 0$$

then $g \in B$ and $g' \notin A$. So

$$\theta_{\mathbf{v}, \mathbf{A} \cup \mathbf{g}'} \theta_{\mathbf{g}^{-1} \mathbf{u}, \mathbf{B}/\mathbf{g}}$$

is also not zeroing out, and again the number of non-zeros is even.

7.3 Infinite Compatibility

Here, we provide a proof that circuits for which an ∞ -extension is defined are $\times$ -compatible. Formally:

Claim 7.4. *Let C be quantum circuit of the form defined in Definition 4.17. Then C is ∞ -compatible.*

Proof. Obviously, translation by the zero vector (or by the unit vertex) keeps the gate in its place, or in other words, in the restriction of C^∞ to $[w]$ - the original Toric - any non-decoding gate is restored. We will need the following claim, which states that this is also true for a decoding gate.

Claim 7.5. *Let $u = (g, l)$ be a decoding gate. The restriction of the image of u in C^∞ to $[w]$ is u (when applied to the infinite domain). Namely:*

$$\left(\prod_{a \in \mathbb{Z}^d} a *_r u \right) |_{[w]} = u$$

Proof. Denote by $\hat{a}^j \in \mathbb{F}_2^d$ the vector which assigns in its i th coordinate 1 if the root that supports the face of u_j has a g_i power greater than n . We argue that the j th face of $-\hat{a}^j *_r u$ is in $[w]$ and that the j th face of any other root translation of u is not in $[w]$. That implies that the restriction

$$\left(\prod_{a \in \mathbb{Z}^d} a *_r u \right) |_{[w]}$$

contains each of u 's faces exactly once, and combined with the fact that u is a Pauli product, it follows that they are the same operator.

By the definition of root translation, we get that $(\hat{a}^j *_r u)_j$ is in $[w]$. Let $a' \in \mathbb{Z}^d$ be such that $a' \neq \hat{a}^j$, and denote by i the coordinate on which they differ. Then, we get that the g_i power of the root of $(a' *_r u)_j$ is at a distance n from the g_i power of the root of $(\hat{a}^j *_r u)_j$; therefore, their roots can't both be in the range $[0, n-1]$, which implies that $(a' *_r u)_j \notin [w]$. □

From here we continue by induction on the number of the gates in C :

1. Base. For an empty circuit, ρ_∞ is just the infinite string of zeros; thus, the restriction of it to $[w]$ locations, namely $\mathbf{Tr}_{/[w]}(\rho_\infty)$, yields a single matrix element corresponds to the w -length string of zeros, which is exactly ρ .
2. Step. Assume the claim holds for every circuit with fewer than τ gates, and consider a circuit C with τ gates. Denote by C' the circuit obtained from C by erasing its last gate; then C' has $\tau - 1$ gates. Denote by ρ' and ρ'_∞ the output states of $C'[\mathcal{E}]$ and $(C')^\infty[\mathcal{E}^\infty]$, respectively. By the induction assumption we have:

$$\rho' = \mathbf{Tr}_{/[w]}(\rho'_\infty)$$

Denote the last gate in C by u , and by $I = \{v_j\}$ the collection of gates in C^∞ to which u is mapped (notice that the collection is countable). Let us split into cases, depending on whether u is or is not a decoding gate.

- (a) Suppose that u is not a decoding gate. Let g and l be the gate element and the location of u . By the definition of ∞ -extension, I equals:

$$I = \bigcup_{a \in \mathbb{Z}^d} \{a * u\}$$

So, there is only one gate in I , corresponding to a being the zero vector, that has support on qubits with spatial locations in the range $[w]$. Denote that gate by v_0 .

$$\begin{aligned} \mathbf{Tr}_{/[w]}(\rho_\infty) &= \mathbf{Tr}_{/[w]} \left(\prod_{j=-\infty}^{\infty} v_j \rho'_\infty \prod_{j=-\infty}^{\infty} v_j^\dagger \right) \\ &= v_0 \mathbf{Tr}_{/[w]} \left(\prod_{j=-\infty, j \neq 0}^{\infty} v_j \rho'_\infty \prod_{j=-\infty, j \neq 0}^{\infty} v_j^\dagger \right) v_0^\dagger \\ &= v_0 \mathbf{Tr}_{/[w]} \left(\prod_{j=-\infty, j \neq 0}^{\infty} v_j^\dagger \prod_{j=-\infty, j \neq 0}^{\infty} v_j \rho'_\infty \right) v_0^\dagger \\ &= v_0 \mathbf{Tr}_{/[w]}(\rho'_\infty) v_0^\dagger \\ &= v_0 \rho'_0 v_0^\dagger \\ &= \rho \end{aligned}$$

- (b) In the case that u is a decoding gate, each of the v_j is a Pauli product. Thus, we can decompose it into the product of two terms: the first has support on $[w]$, and the second is its complement. Exactly as in the above case, the second v_j component can be treated as a disjoint gate that has no support on qubits in $[w]$ and therefore does not contribute to the trace. For the rest, we get:

$$\mathbf{Tr}_{/[w]}(\rho_\infty) = \left(\prod_{-\infty}^{\infty} v_j \right) |_{[w]} \rho' \left(\left(\prod_{-\infty}^{\infty} v_j \right) |_{[w]} \right)^\dagger$$

Then, by Claim 7.5, we have that:

$$\left(\prod_{-\infty}^{\infty} v_j \right) |_{[w]} = u$$

therefore:

$$\mathbf{Tr}_{/[w]}(\rho_\infty) = u \rho' u^\dagger = \rho$$

□

7.4 Counting Subtrees.

Claim 7.6. *Let $G = (V, E)$ be a $\Delta + 1$ -regular graph, and let k be an integer less than $|E|$. Fix a vertex $r \in V$. There is a constant number α , depending only on Δ , such that the number of subgraphs rooted at r with fewer than k edges is at most α^k .*

Proof. Denote by T_Δ and T_2 the infinity Δ and binary tree. There is an injection between the subtrees of G , rooted at v , with at most k edges, to the subtrees of T_Δ rooted at v , with at most k edges. Moreover, there is an injection from the subtrees of T_Δ rooted at v , with at most k edges to the subtrees of T_2 , rooted at v , with at most $k \log \Delta$ edges. The injection is obtained by replacing each of the vertex with a (complete) binary tree over $\Delta - 1$ leaves, each leaf is associated with a outgoing edge of the original vertex.

Thus, it's enough to bound the number of of subtrees in T_2 rooted at v , with at most $k \log \Delta$ edges, that number is less than the number of binary tree with $k \log \Delta + 1$ leaves. For exact leaves number we get the $k \log \Delta$ Catalan number, which its limit is:

$$\sim \frac{4^{k \log \Delta}}{(k \log \Delta)^{\frac{3}{2}} \sqrt{\pi}} \leq \frac{\Delta^{2k}}{(k \log \Delta)^{\frac{3}{2}} \sqrt{\pi}}$$

Therefore we can bound by:

$$\sum_{j=1}^k \frac{\Delta^{2j}}{(j \log \Delta)^{\frac{3}{2}} \sqrt{\pi}} \leq \frac{\Delta^{2k}}{\sqrt{\pi} \log^{\frac{3}{2}} \Delta} \sum_{j=1}^k \frac{1}{j^{\frac{3}{2}}} \leq \zeta\left(\frac{3}{2}\right) \cdot \frac{\Delta^{2k}}{\sqrt{\pi} \log^{\frac{3}{2}} \Delta}$$

□

7.5 Tree Separator Lemma

For a tree $T = (V, E)$ and a vertex $v \in V$, we call the subtree obtained by taking the descendants of v the subtree rooted at v .

Claim 7.7. *Let $T = (V, E)$ be a tree with $m > 3$ leaves and τ edges, then it has a subtree with $m' \in (\frac{1}{3}m, \frac{2}{3}m)$ leaves and at most $\frac{m'}{m}\tau$ edges.*

Proof. Order the vertices according to their heights. Denote by $v \in V$ the lowest vertex for which the subtree rooted at v has at least $\frac{1}{3}m$ leaves. Now, denote by T_1 the subtree obtained through the following iterative process over v 's children:

1. Initialize T_1 as the subtree rooted at v .
2. For a child u of v :
 - (a) Check if erasing from T_1 the subtree rooted at u doesn't decrease the number of leaves below $\frac{1}{3}m$. If so, then we update T_1 to be the result of the erasure.

Clearly, the algorithm is defined such that T_1 has to have at least $\frac{m}{3}$ leaves. Furthermore, since any child of v is a root of a subtree with fewer than $\frac{m}{3}$ leaves (otherwise we get a contradiction to the minimality of v 's height), we have that if T_1 would have more than $\frac{2}{3}m$ leaves, then the subtraction of any of its children would yield a tree with at least $\frac{2}{3}m - \frac{1}{3}m = \frac{1}{3}m$ leaves. Namely, v has in T_1 a child u such that the subtraction of the subtree rooted at u from T_1 results in a subtree with more than $\frac{1}{3}m$ leaves. But if indeed there were such a child, then the algorithm would erase it, so there is not.

Now let $T_2 = T/T_1 \cup v$, namely the tree induced by the subtraction of T_1 's edges from T . Observe that since the summation of leaves in T_1 and T_2 is the leaves in T , and T_1 has no more than $\frac{2}{3}m$ leaves, and no less than $\frac{1}{3}m$ leaves, it holds that the number of leaves in T_2 is also in the range $(\frac{1}{3}m, \frac{2}{3}m)$. We claim that at least one of the subtrees has fewer than $\frac{m'}{m}\tau$ edges.

To see this, denote by m_1, m_2 and by τ_1, τ_2 the leaves and the edges in T_1, T_2 respectively. Now:

$$\frac{\tau}{m} = \frac{\tau_1 + \tau_2}{m_1 + m_2} = \frac{m_1}{m_1 + m_2} \frac{\tau_1}{m_1} + \frac{m_2}{m_1 + m_2} \frac{\tau_2}{m_2}$$

Namely, $\frac{\tau}{m}$ is a weighted average of $\frac{\tau_1}{m_1}$ and $\frac{\tau_2}{m_2}$, and therefore it is greater than their minimum. Peak the subtree with the minimal quotient. $\square$

By repeating recursively on the above construction, we get the following corollary:

Corollary 7.8. *Let $T = (V, E)$ be a tree with m leaves and τ edges. Then, for any $\alpha > 0$, it has a subtree with $m' \in (\frac{1}{3}\alpha m, \frac{2}{3}\alpha m)$ leaves and at most $\frac{m'}{m}\tau$ edges.*

Claim 7.9. *Let $\gamma, \beta \in \mathbb{R}_{>0}$ and n be a positive integer. For any tree T with $\tau \geq \beta n$ edges and m leaves such that $\tau \leq \gamma m$, there is a subtree T' of T with a number of edges τ' less than $\frac{\beta}{2}n$ and a number of leaves greater than $\frac{\beta}{4\gamma}n$.*

Proof. Since the construction of claim 7.7 preserves the ratio of edges to leaves, we have that for a T' obtained using the construction, the number of edges satisfies $\tau' \leq \gamma m'$. So, it's enough to look for the maximal α such that:

$$\gamma m' \leq \gamma \frac{2}{3}\alpha m \leq \frac{\beta}{2}n \Rightarrow \alpha = \frac{3\beta}{4\gamma} \frac{n}{m}$$

Which gives $m' \geq \frac{1}{3}\alpha m = \frac{\beta}{4\gamma}n$. $\square$